

# Exercise Session 1

Performance, Amdahl's Law, Pipeline

Advanced Computer Architectures

Politecnico di Milano

March 11th, 2024

Davide Conficconi <davide.conficconi@polimi.it>



**POLITECNICO**  
MILANO 1863

POLITECNICO MILANO 1863

**NECST**  
laboratory



# Who am I

Assistant Professor @ Politecnico di Milano

Research in the  
System Architecture Area



Focus: (Co-)Design Domain-Specific Computer  
Architectures and Systems

Performance

Energy Efficiency

Reconfigurable Systems  
(e.g., FPGAs)





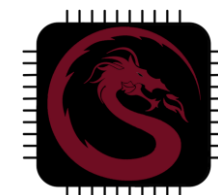
# Who am I

Assistant Professor @ Politecnico di Milano

Research in the  
System Architecture Area



Focus: (Co-)Design Domain-Specific Computer  
Architectures and Systems

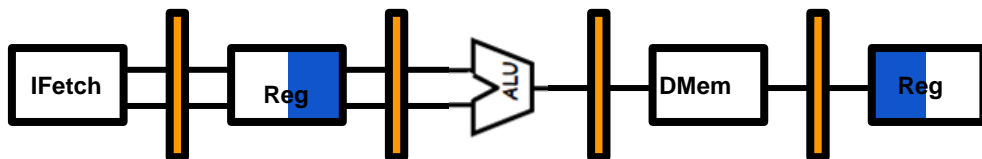


Performance

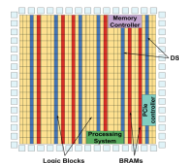
Energy Efficiency

Reconfigurable Systems  
(e.g., FPGAs)

Adjunct Prof. for Bachelor CS101 and Advanced Computer Architecture



Lecturer of FPGA101 Passion In Action



**POLITECNICO**  
MILANO 1863



# Who am I

Assistant Professor @ Politecnico di Milano

Research in the  
System Architecture Area



Focus: (Co-)Design Domain-Specific Computer  
Architectures and Systems

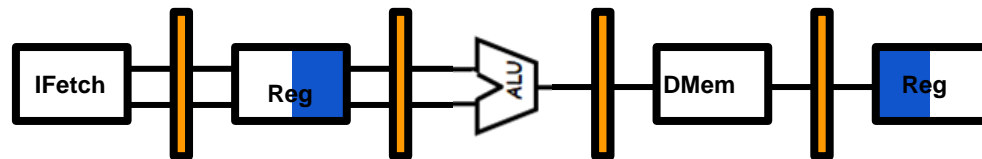


Performance

Energy Efficiency

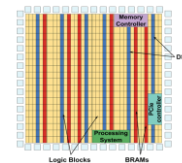
Reconfigurable Systems  
(e.g., FPGAs)

Adjunct Prof. for Bachelor CS101 and Advanced Computer Architecture



**POLITECNICO**  
MILANO 1863

Lecturer of FPGA101 Passion In Action



Intern at research teams of IBM (21/22), Xilinx (18/19) Oracle (18)

IBM **Research** | Zurich

 **XILINX**

**ORACLE**

# Important things: Material

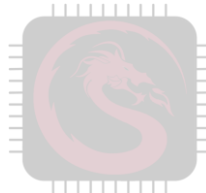


Assistant Professor @ Politecnico di Milano

Research in the  
System Architecture Area



Focus: (Co-)Design Domain-Specific Computer  
Architectures and Systems

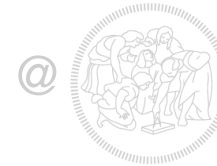
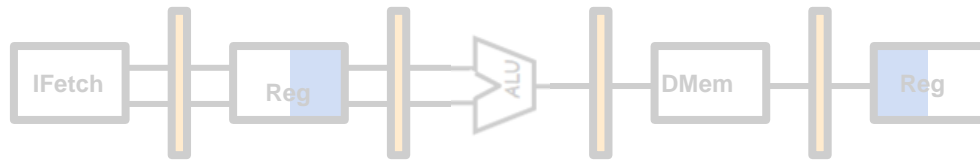


Performance

Energy Efficiency

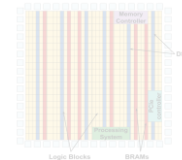
Reconfigurable Systems  
(e.g., FPGAs)

Adjunct Prof. for Bachelor CS101 and Advanced Computer Architecture



**POLITECNICO**  
MILANO 1863

Lecturer of FPGA101 Passion In Action



Intern at research teams of IBM (21/22), Xilinx (18/19) Oracle (18)

**<https://webeep.polimi.it/course/view.php?id=10616>**

IBM Research | Zurich

XILINX

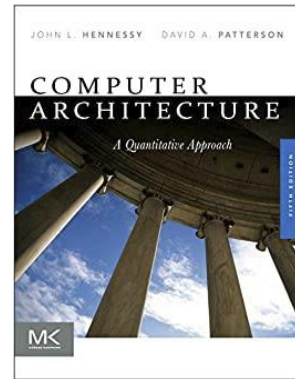
ORACLE

# Material

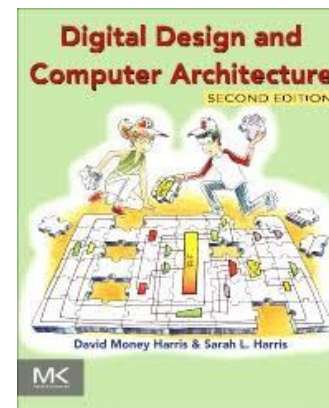
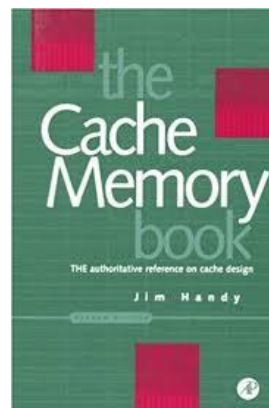
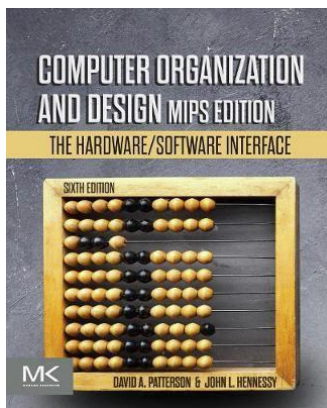
## EVERYTHING OPTIONAL

<https://webeep.polimi.it/course/view.php?id=10616>

Textbook: Hennessy and Patterson, Computer Architecture: A Quantitative Approach



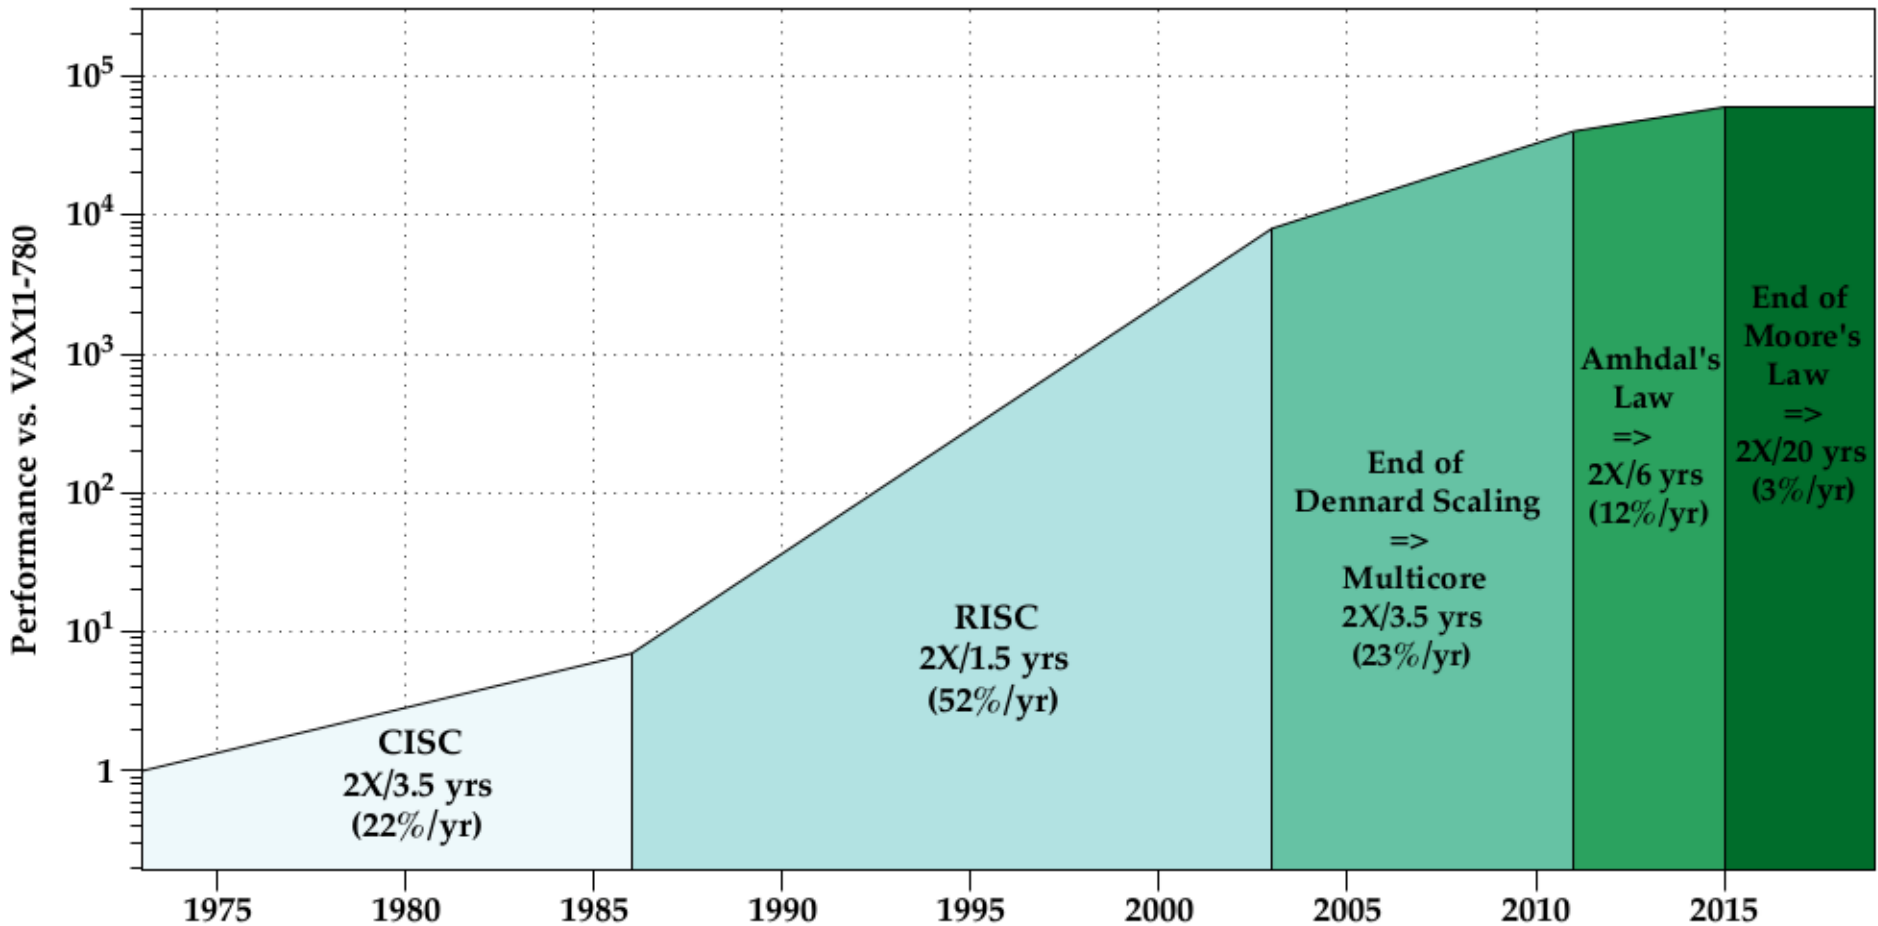
Other Interesting Reference



# Who Are You?



# Motivations

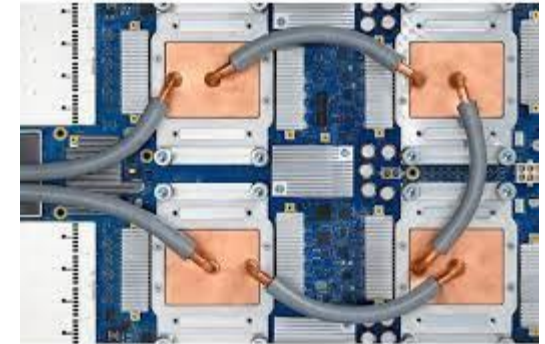
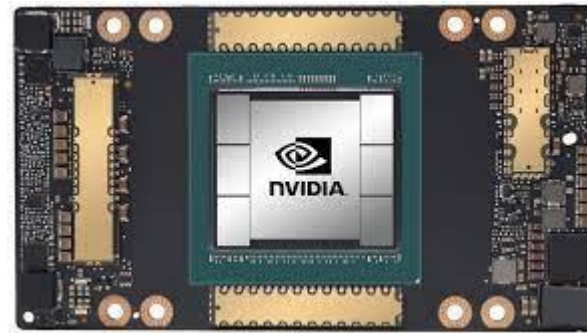


Adapted from E. Del Sozzo. On how to effectively target fpgas from domain specific tools. 2019.  
Data from: J. L. Hennessy and D. A. Patterson. Computer architecture: a quantitative approach 6th edition. Elsevier, 2018.



# The End Goal is... Pick the Best Architecture

(it is a matter of trade-offs)



# News from the outer world: Reviving the Spartan Technology

<https://www.nextplatform.com/2024/03/08/amd-flexing-spartan-fpga-muscles-in-clouds-and-at-edges/>



<https://www.nextplatform.com/2022/07/08/now-comes-the-hard-part-amd-software/>



<https://www.nextplatform.com/2024/02/29/the-once-and-future-fpga-maker-altera/>

# News from the outer world: Reviving the Spartan Technology

<https://www.nextplatform.com/2024/03/08/amd-flexing-spartan-fpga-muscles-in-clouds-and-at-edges/>



<https://www.nextplatform.com/2024/03/08/amd-flexing-spartan-fpga-muscles-in-clouds-and-at-edges/>  
[now-comes-the-hard-part-amd](#)

## AMD SPARTAN™ ULTRASCALE+™ FPGA: OPTIMIZED FOR THE EDGE

### Proven 16 nm Technology

#### Low Power

- 16 nm FinFET & hardened DDR, PCIe®

#### Block RAM and UltraRAM

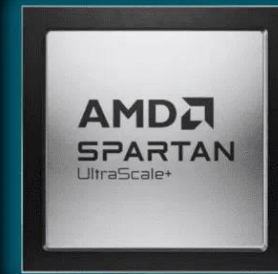
- Up to 26 Mb of on-chip memory

#### Transceivers and PCIe®

- Up to 8 GTx with 16.3 Gb/s
- PCIe® Gen4 x8

#### Digital Signal Processing

- Up to 384 DSP48E2 blocks
- Floating/fixed point



### Future-Ready Capabilities

#### Flexible I/O Interfaces

- Up to 572 I/Os, 3.3V support
- 3.2G MIPI D-PHY

#### State-of-the-Art Security

- Post-Quantum Crypto (PQC) ready
- PPK/SPK authentication, TRNG, PUF...

#### Hard Memory Controller

- LPDDR4x/5 up to 4266 Mb/s
- First UltraScale+ with LPDDR5

#### Small Form Factor

- CSP and BGA packages
- As small as 10x10 mm

# News from the outer world: Reviving the Spartan Technology

<https://www.nextplatform.com/2024/03/08/amd-flexing-spartan-fpga-muscles-in-clouds-and-at-edges/>

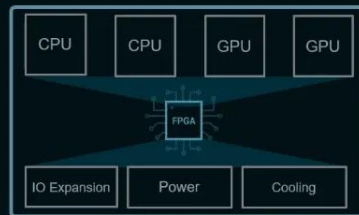
## AMD SPARTAN™ ULTRASCALE+™ FPGA: OPTIMIZED FOR THE EDGE

Proven 16 nm Technology

### BUILT FOR INTERFACING

#### Highest I/O to Logic Cell Ratio

Cost-Effective I/O Expansion & Baseboard Management Controller (BMC)



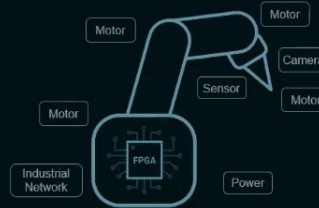
Datacenter BMC

**2.4X** I/O to LC ratio vs. AMD Artix™ 7 FPGA

**3.5X** I/O to LC ratio vs. AMD Spartan™ 7 FPGA

#### Increased Connectivity and Flexibility

Any-to-Any Connectivity for Edge Sensing and Control



Industrial Robotics

**2.5X** Transceiver Bandwidth vs. AMD Artix 7 FPGA

**4X** MIPI Bandwidth vs. AMD Spartan 7 / Artix 7 FPGAs

Industry's highest I/O to Logic Cell Ratio for 28nm and newer FPGAs

#### Future-Ready Capabilities

##### Flexible I/O Interfaces

- Up to 572 I/Os, 3.3V support
- 3.2G MIPI D-PHY

##### State-of-the-Art Security

- Post-Quantum Crypto (PQC) ready
- PPK/SPK authentication, TRNG, PUF...

##### Hard Memory Controller

- LPDDR4x/5 up to 4266 Mb/s
- ECC, parity, and scrubbing

##### Small Form Factor

- CSP and BGA packages
- As small as 10x10 mm

AMD  
together we advance

# News from the outer world: Reviving the Spartan Technology

<https://www.nextplatform.com/2024/03/08/amd-flexing-spartan-fpga-muscles-in-clouds-and-at-edges/>

## POWER EFFICIENCY ACROSS THE PORTFOLIO

### UP TO 30% LOWER TOTAL POWER<sup>1</sup> WITH 16 NM FABRIC

STATIC & DYNAMIC POWER @ 80% UTILIZATION, 200 MHZ

LOWEST DENSITY  
DEVICES



28 nm

Artix™ 7 FPGA  
(7A35T)

16 nm

Spartan™ UltraScale+™  
FPGA (SU35P)

1.9X Fabric Performance<sup>1</sup>

1.2X Total I/O

### UP TO 60% LOWER POWER INTERFACING<sup>1</sup>

TOTAL DEVICE POWER NORMALIZED TO PCIE AND DDR BANDWIDTH

DEVICES WITH  
HARDENED DDR  
AND PCIE



28 nm

Artix 7 FPGA  
(7A200T)

16 nm

Spartan UltraScale+  
(SU200P)

5X Memory Controller Bandwidth

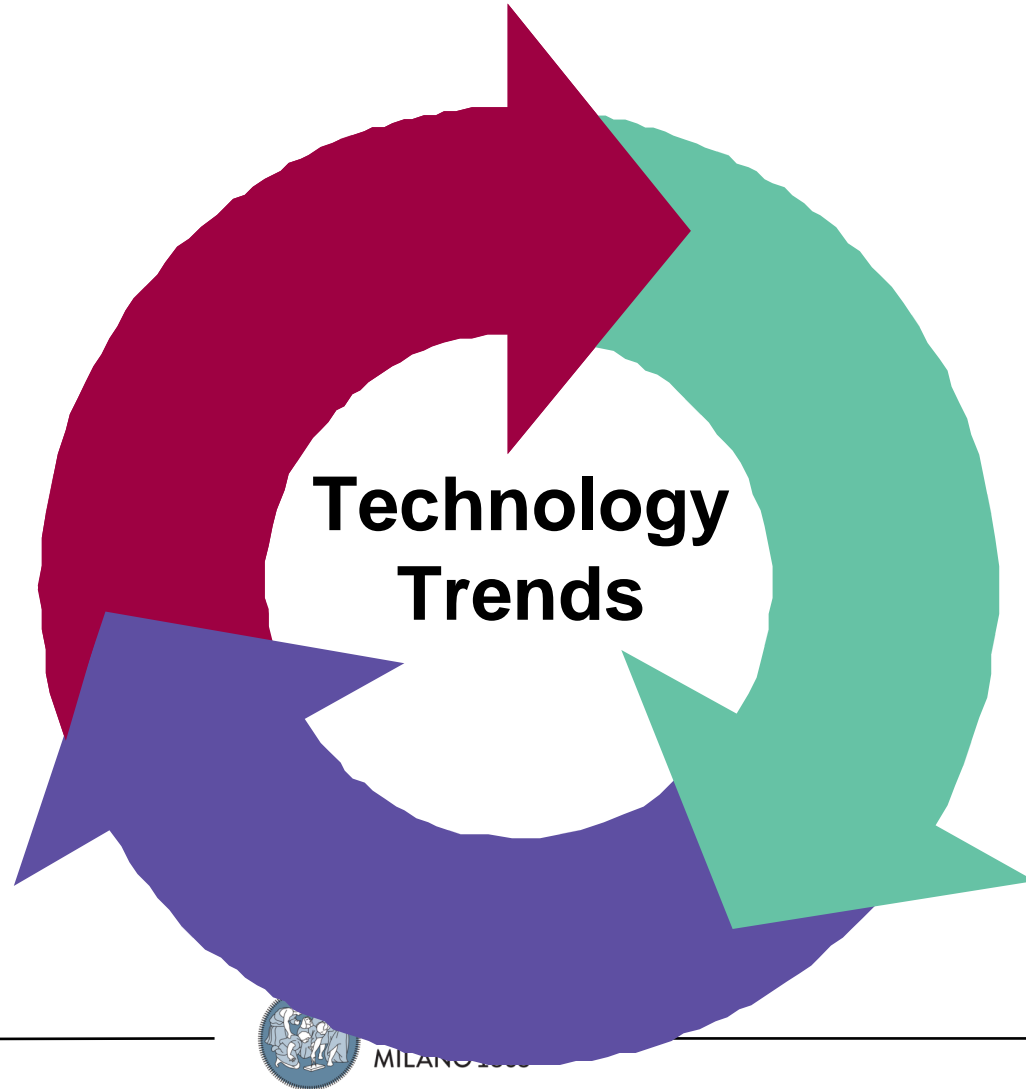
4X PCIe Bandwidth

# Recall: Computer Engineering Methodology



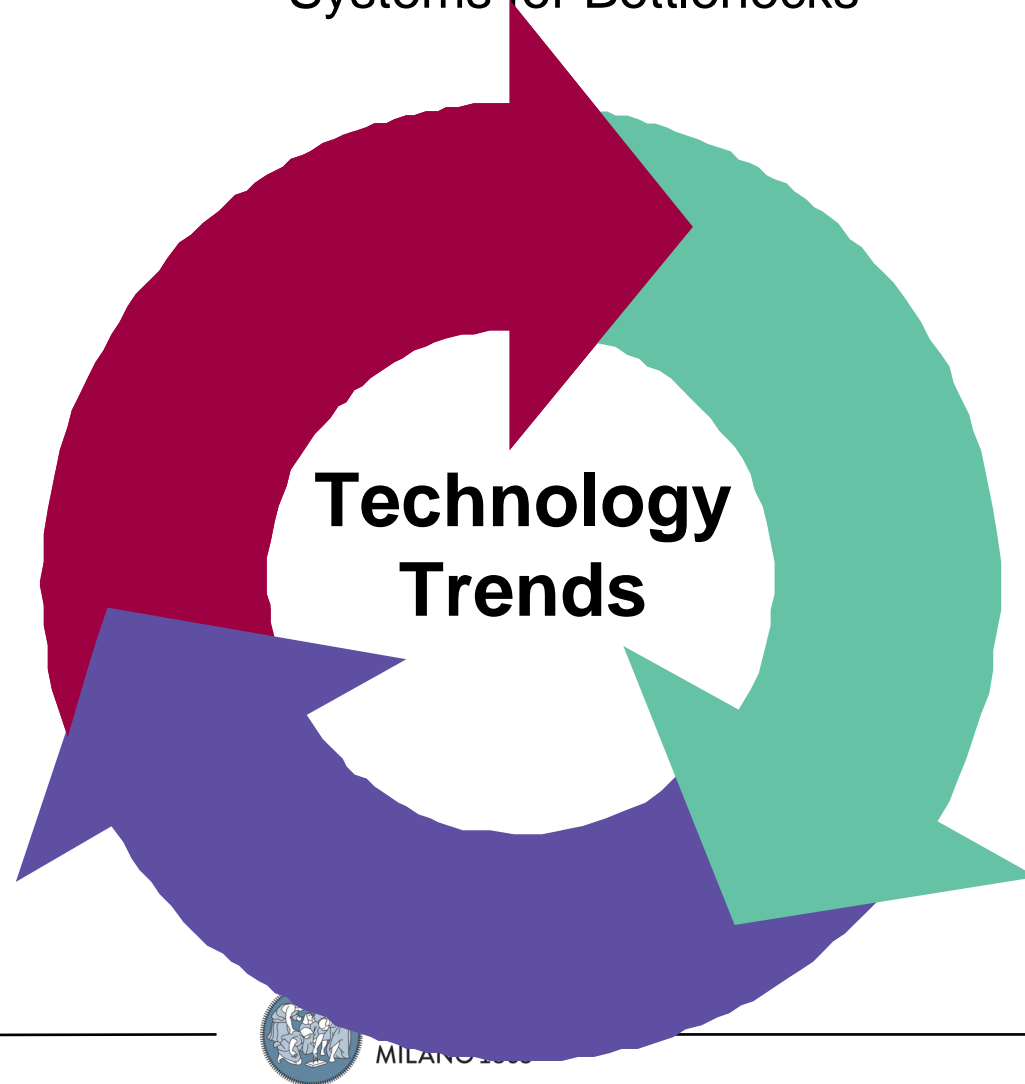


# Recall: Computer Engineering Methodology



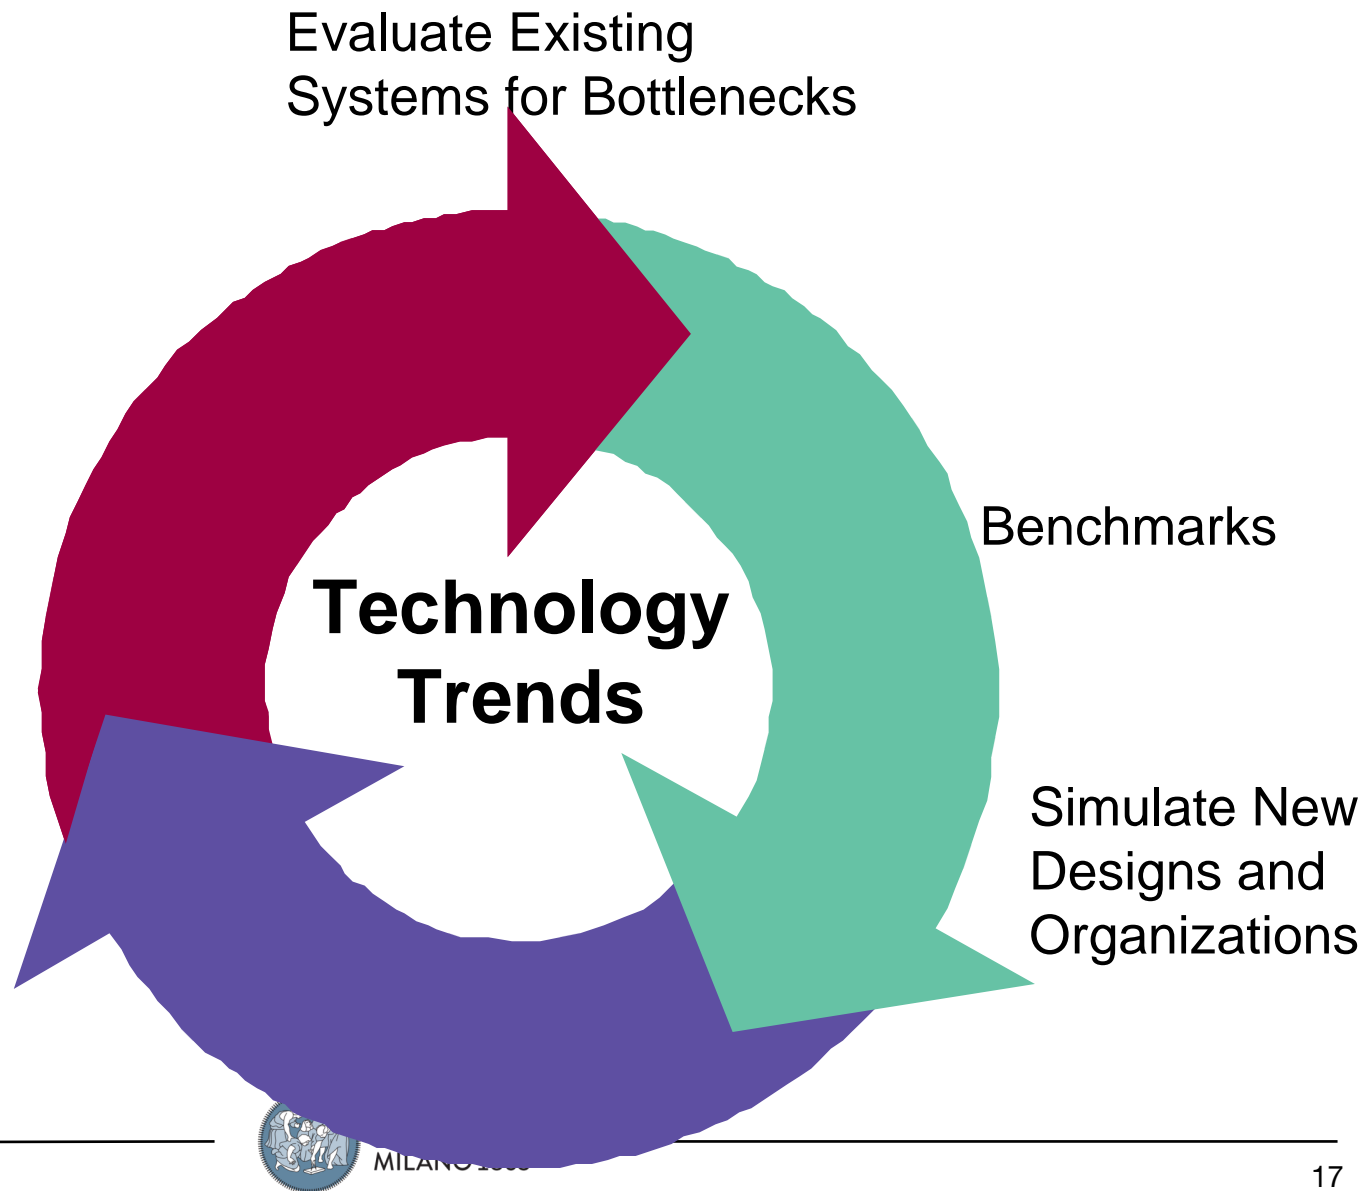
# Recall: Computer Engineering Methodology

Evaluate Existing  
Systems for Bottlenecks

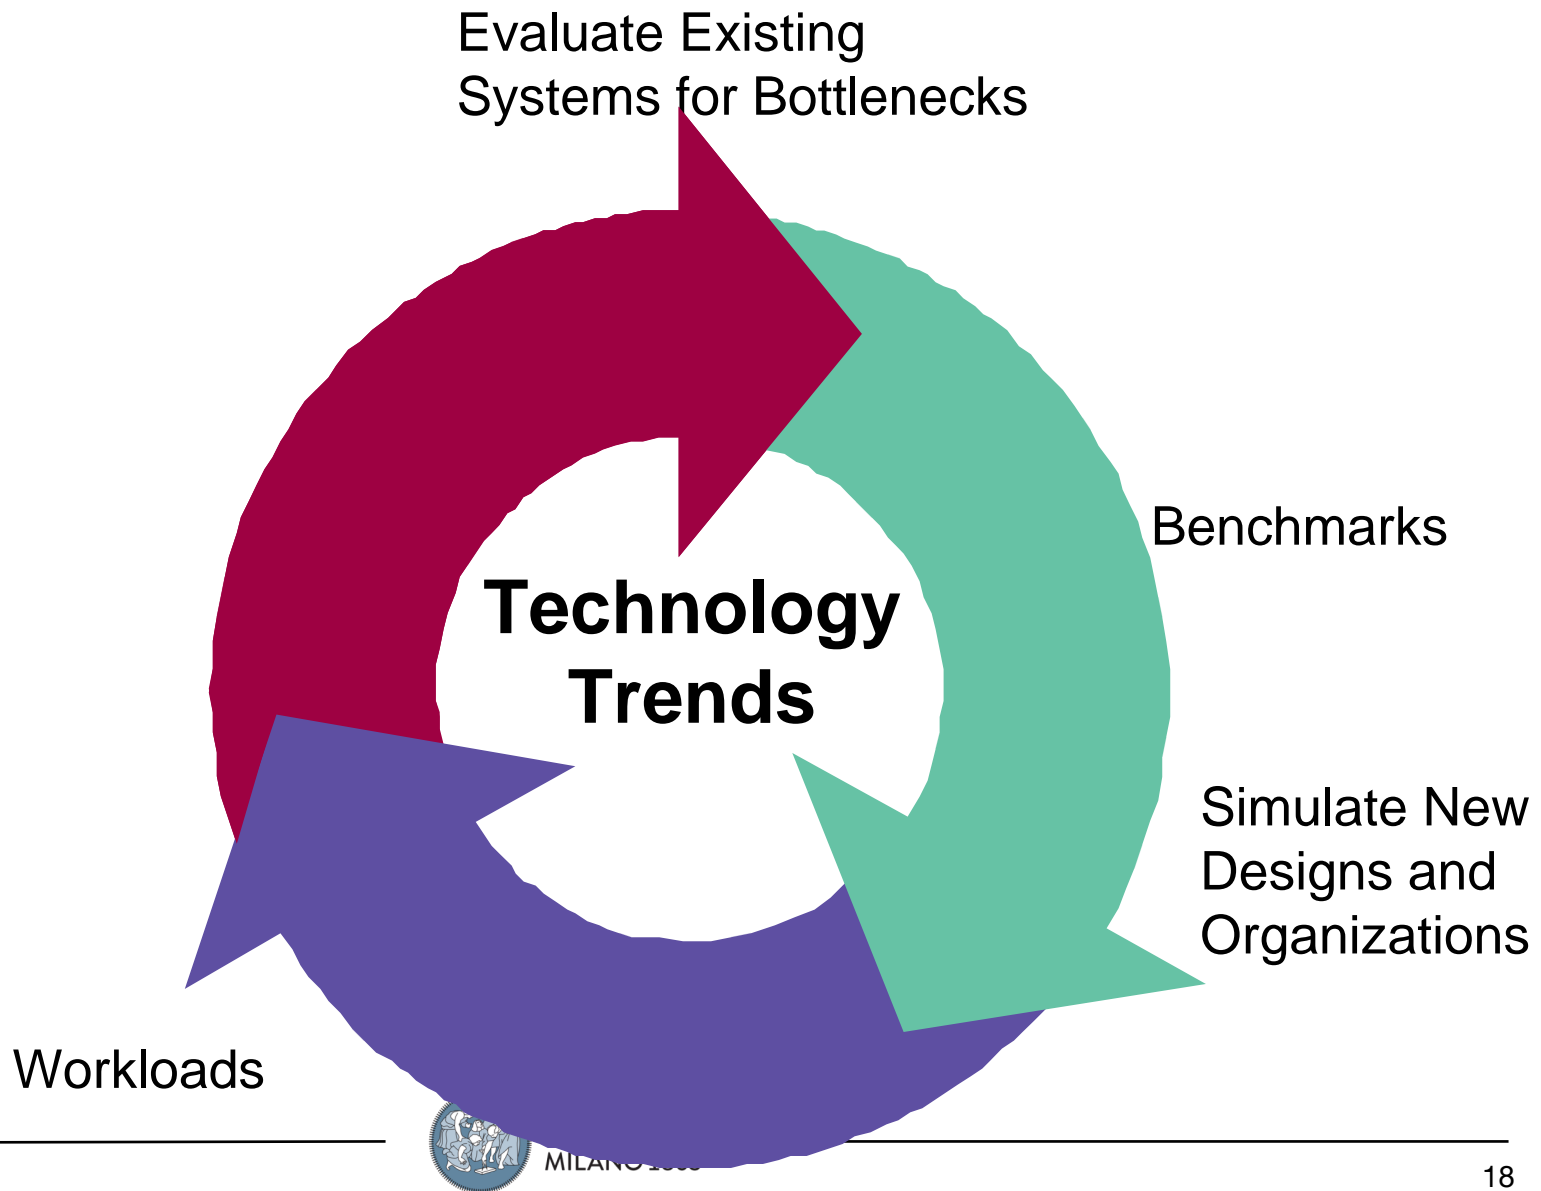




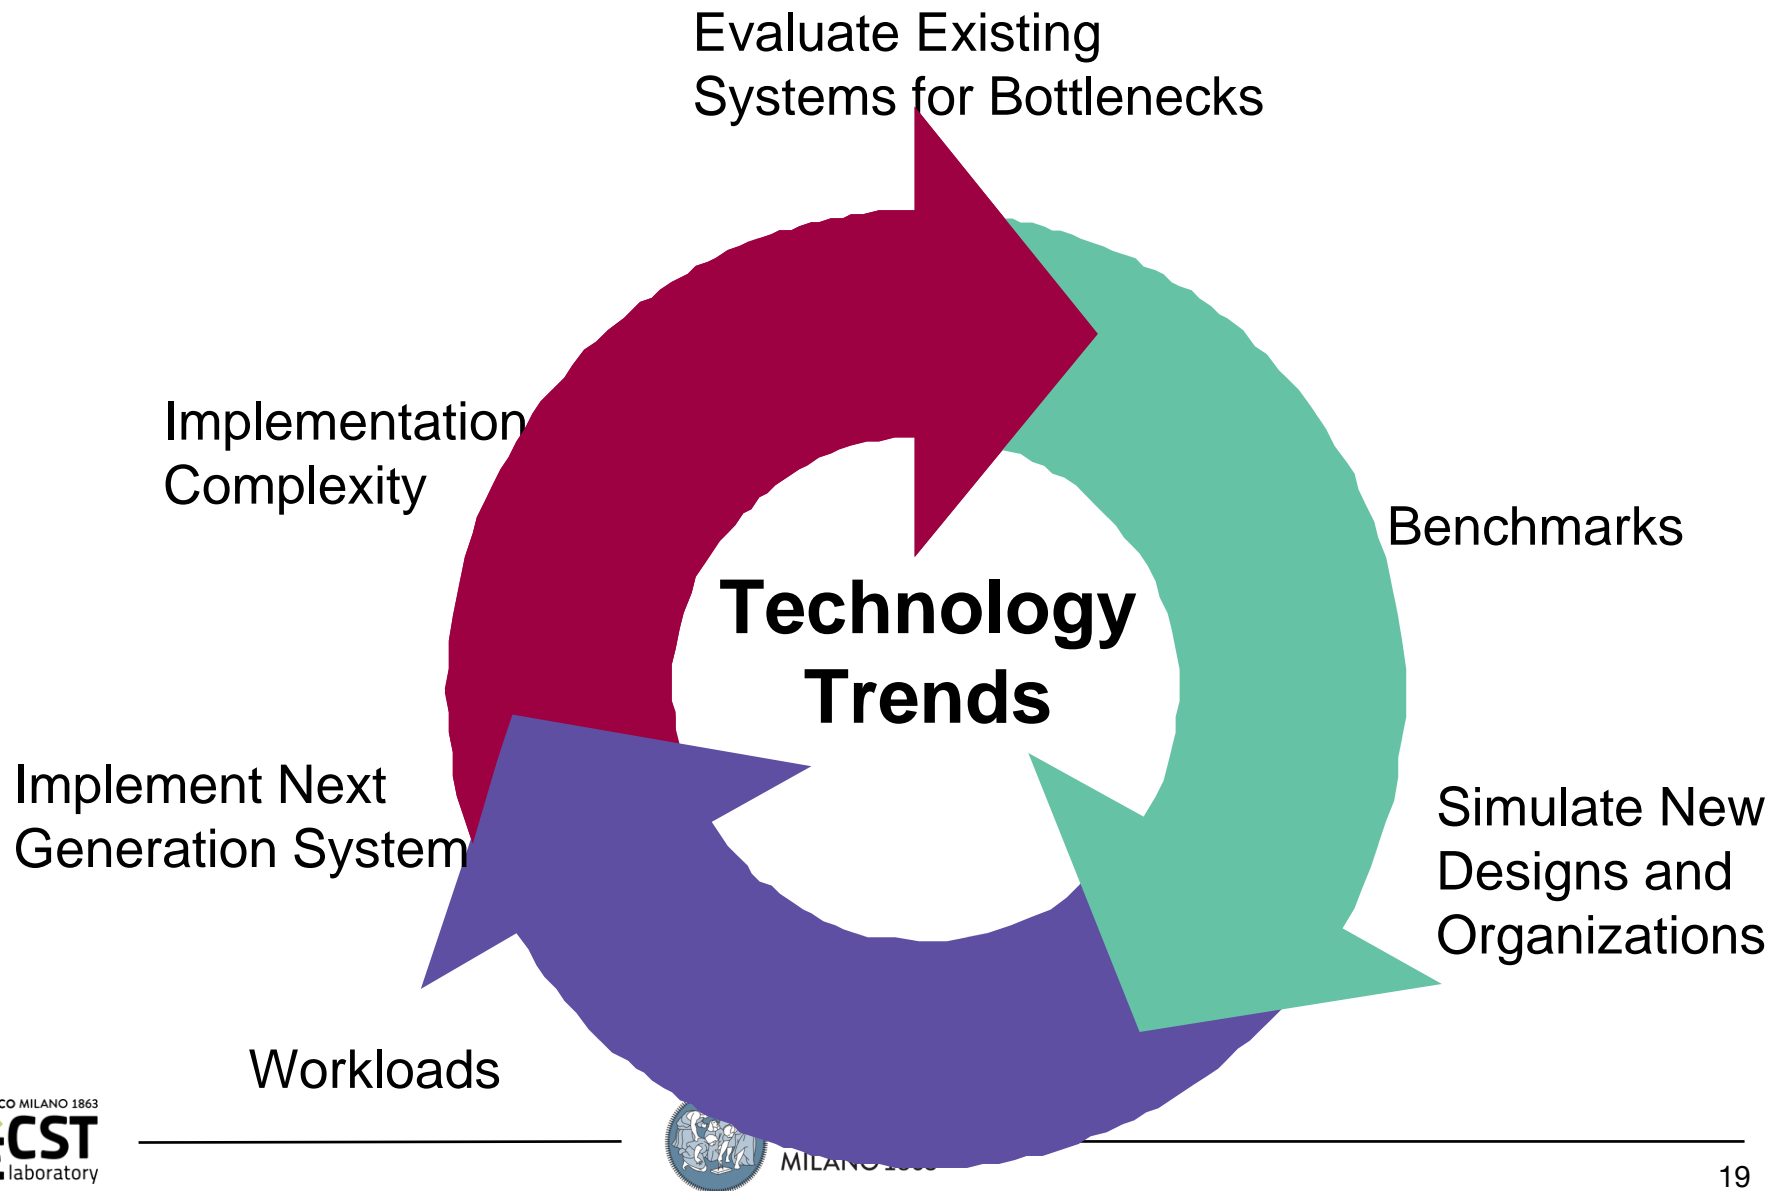
# Recall: Computer Engineering Methodology



# Recall: Computer Engineering Methodology



# Recall: Computer Engineering Methodology



# Exe 1: Throughput vs Response time

# Recall: Throughput vs Response time

- Two Metrics:

- Computer system user

- Minimize elapsed time for program execution:

**response time**:  $\text{execution time} = \text{time\_end} - \text{time\_start}$

- Computer center manager

- Maximize completion rate =  $\# \text{jobs/sec}$

**throughput**: total amount of work done in a given time

# Exe 1: Throughput vs Response time



# Exe 1: Throughput vs Response time



What will happen if...

# Exe 1: Throughput vs Response time



(1)



What will happen if...

(1) we replace with a faster version?



# Exe 1: Throughput vs Response time



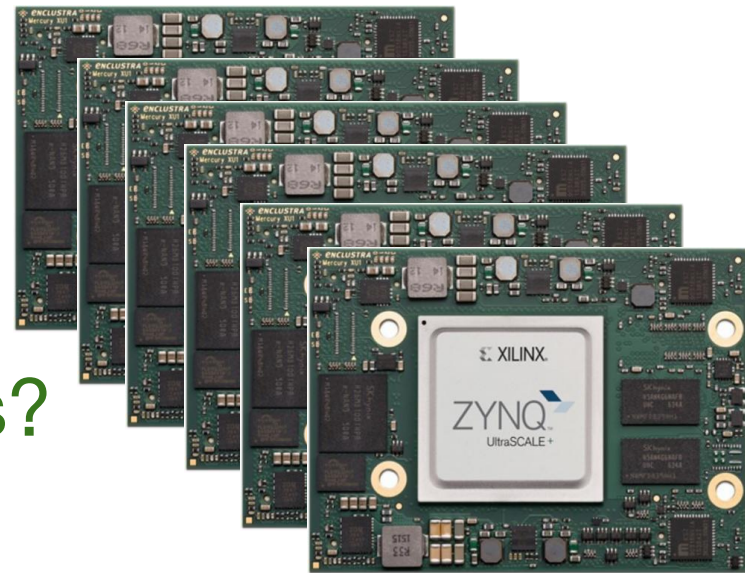
(1)



What will happen if...

(1) we replace with a faster version?

(2) We add multiple parallel systems for independent tasks?



# Case 1: Scale-up

(1)



# Case 1: Scale-up

(1)



decrease response time and  
throughput will increase

# Case 2: Scale-out

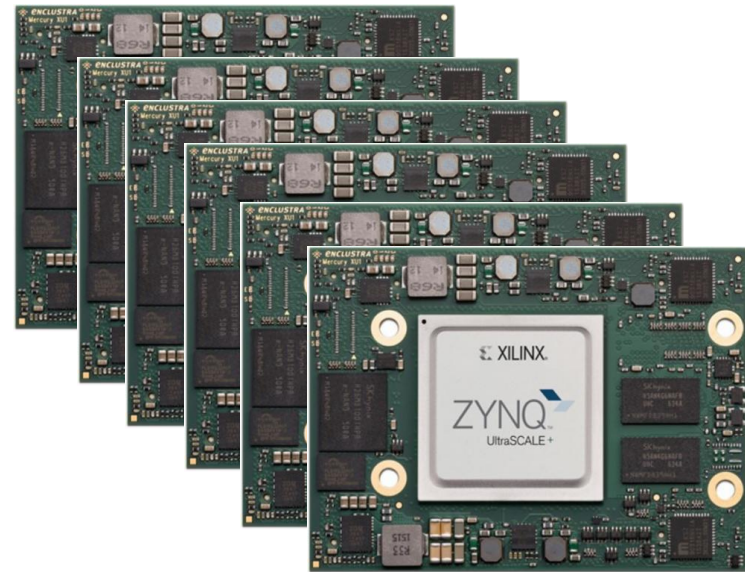
(2)



# Case 2: Scale-out

For sure Throughput will increase

(2)

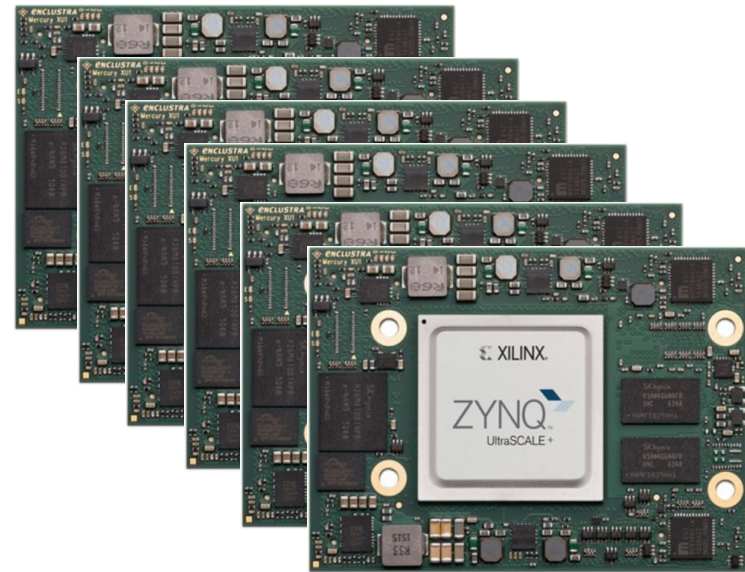


# Case 2: Scale-out

For sure Throughput will increase

Response time?

(2)





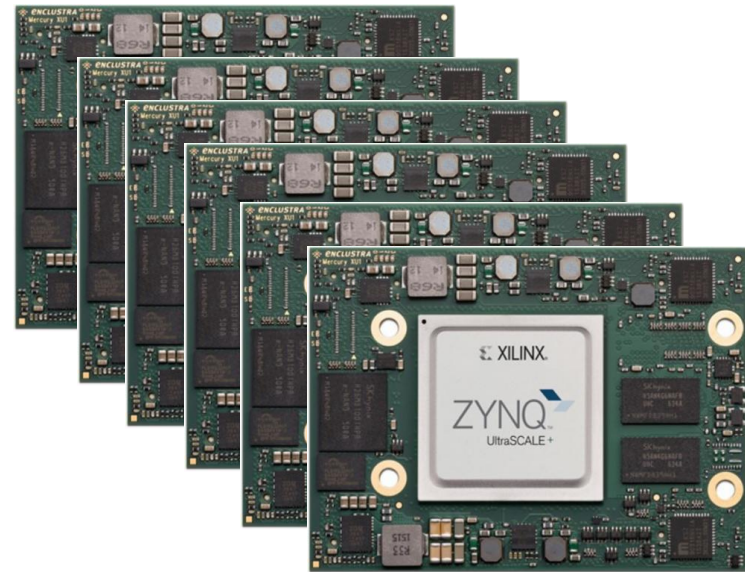
# Case 2: Scale-out

For sure Throughput will increase

(2)

Response time?

Yes, if there were a queue to serve, which was waiting for computing resources





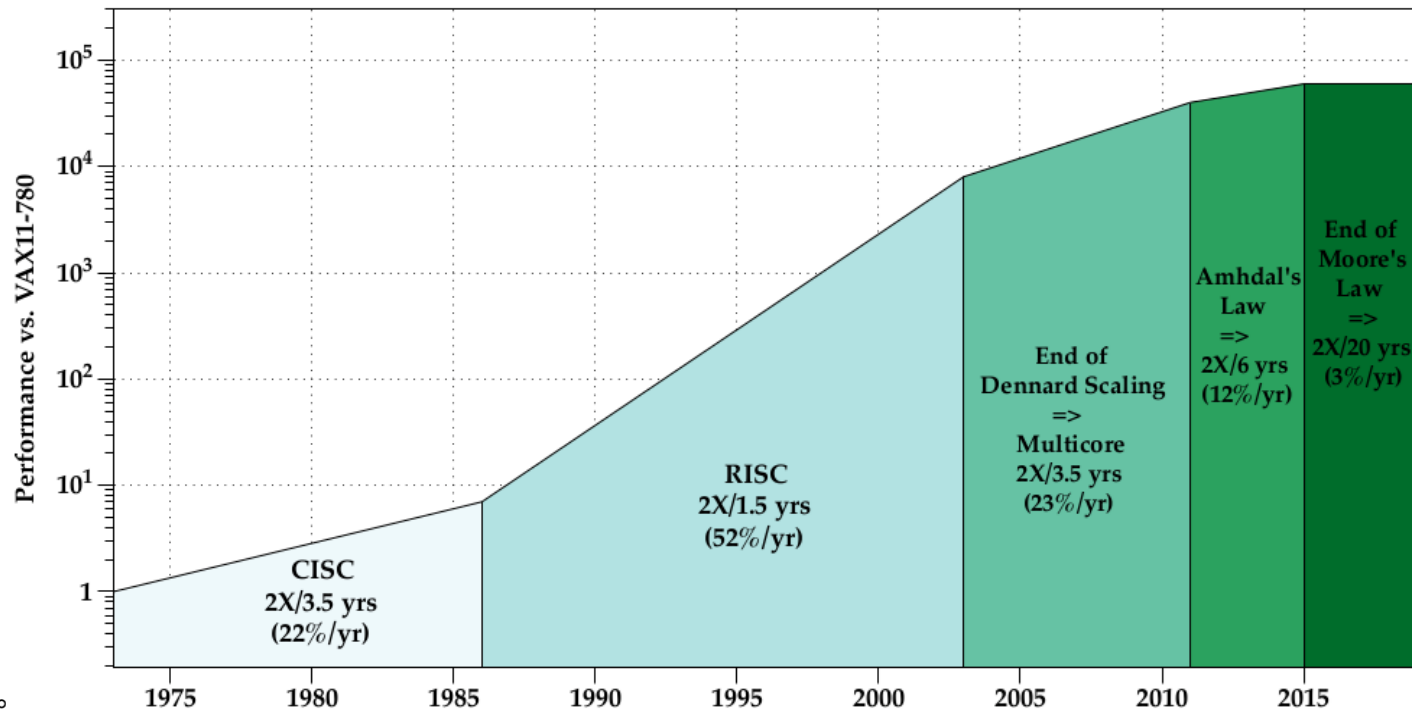


# Recall: Issues as new opportunities

Programming has become very difficult  
Impossible to balance all constraints manually

- More computational horse-power than ever before
  - Cores are free (almost ... <https://doi.org/10.1145/2000064.2000108>)
  - Energy (i.e., perf/joule) is **ALWAYS** a primary concern → Scaling (strong vs weak)

<https://www.kth.se/blogs/pdc/2018/11/scalability-strong-and-weak-scaling/>



# Recall: Some Factors Affecting Performance

Algorithm complexity and data set

Compiler

Instruction set

Available operations

Operating system

Clock rate

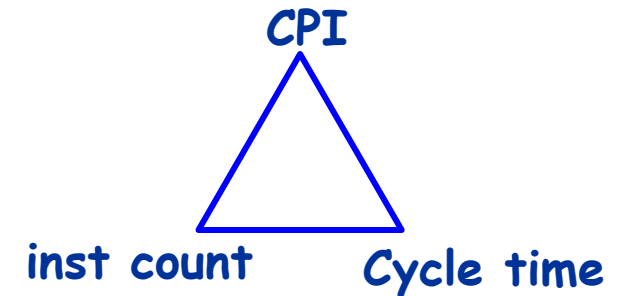
Memory system performance

I/O system performance and overhead

# Recall: CPU time

- Instruction Count, IC
  - Instructions executed, not static code size
  - Determined by algorithm, compiler, Instruction Set Architecture
- Cycles per instructions, CPI
  - Determined by ISA and CPU organization
  - Overlap among instructions (pipelining) reduces this term
- Time/cycle
  - Determined by technology, organization and circuit design

# Recall: Performance equation



	Inst. Count	CPI	Clock Rate
Program	X		
Compiler	X	(X)	
Instr. Set	X	X	
Organization		X	X
Technology			X

# Exe 2: Performance Problem

# Exe 2: Performance Problem

Consider two CPUs: CPU1 and CPU2.

CPU1 has clock cycle of 2 ns while CPU2 has an operating frequency of 700MHz.

# Exe 2: Performance Problem

Consider two CPUs: CPU1 and CPU2.

CPU1 has clock cycle of 2 ns while CPU2 has an operating frequency of 700MHz.

Given the following frequencies of occurrence of the instructions for the two CPUs

# Exe 2: Performance Problem

Consider two CPUs: CPU1 and CPU2.

CPU1 has clock cycle of 2 ns while CPU2 has an operating frequency of 700MHz.

Given the following frequencies of occurrence of the instructions for the two CPUs

<b>Operation type</b>	<b>Frequency</b>	<b>CPU1</b> CYCLE	<b>CPU2</b> CYCLE
A	0.3	2	2
B	0.1	3	3
C	0.2	4	3
D	0.3	2	2
E	0.1	4	3



# Exe 2: Questions

Operation type	Frequency	CPU1 CYCLE	CPU2 CYCLE
A	0.3	2	2
B	0.1	3	3
C	0.2	4	3
D	0.3	2	2
E	0.1	4	3

- A. Compute the average CPI for CPU1 and CPU2
- B. Which is the fastest CPU?

# Exe 2: AVG CPIs

Operation type	Frequency	CPU1 CYCLE	CPU2 CYCLE
A	0.3	2	2
B	0.1	3	3
C	0.2	4	3
D	0.3	2	2
E	0.1	4	3

Recall: 
$$\text{CPI} = \frac{\text{Clock cycles}}{\text{Instruction}}$$

$$\text{CPI} = \sum_{i=1}^n \text{CPI}_i * F_i \quad \text{where} \quad F_i = \frac{I_i}{\text{Instruction Count}}$$

## Exe 2: AVG CPIs

Operation type	Frequency	CPU1 CYCLE	CPU2 CYCLE
A	0.3	2	2
B	0.1	3	3
C	0.2	4	3
D	0.3	2	2
E	0.1	4	3

Recall: 
$$CPI = \frac{\text{Clock cycles}}{\text{Instruction}}$$

$$CPI = \sum_{i=1}^n CPI_i * F_i \quad \text{where} \quad F_i = \frac{I_i}{\text{Instruction Count}}$$

$$CPI_1 = 0.3 * 2 + 0.1 * 3 + 0.2 * 4 + 0.3 * 2 + 0.1 * 4 =$$

## Exe 2: AVG CPIs

Operation type	Frequency	CPU1 CYCLE	CPU2 CYCLE
A	0.3	2	2
B	0.1	3	3
C	0.2	4	3
D	0.3	2	2
E	0.1	4	3

Recall: 
$$CPI = \frac{\text{Clock cycles}}{\text{Instruction}}$$
 
$$CPI = \sum_{i=1}^n CPI_i * F_i \quad \text{where} \quad F_i = \frac{I_i}{\text{Instruction Count}}$$

$$CPI_1 = 0.3 * 2 + 0.1 * 3 + 0.2 * 4 + 0.3 * 2 + 0.1 * 4 =$$

$$= 0.6 + 0.3 + 0.8 + 0.6 + 0.4 = 2.7$$

## Exe 2: AVG CPIs

Operation type	Frequency	CPU1 CYCLE	CPU2 CYCLE
A	0.3	2	2
B	0.1	3	3
C	0.2	4	3
D	0.3	2	2
E	0.1	4	3

Recall: 
$$\text{CPI} = \frac{\text{Clock cycles}}{\text{Instruction}} \quad \text{CPI} = \sum_{i=1}^n \text{CPI}_i * F_i \quad \text{where} \quad F_i = \frac{I_i}{\text{Instruction Count}}$$

$$\begin{aligned} \text{CPI}_1 &= 0.3 * 2 + 0.1 * 3 + 0.2 * 4 + 0.3 * 2 + 0.1 * 4 = \\ &= 0.6 + 0.3 + 0.8 + 0.6 + 0.4 = 2.7 \end{aligned}$$

$$\begin{aligned} \text{CPI}_2 &= 0.3 * 2 + 0.1 * 3 + 0.2 * 3 + 0.3 * 2 + 0.1 * 3 = \\ &= 0.6 + 0.3 + 0.6 + 0.6 + 0.3 = 2.4 \end{aligned}$$

# Exe 2.b: Which is the fastest CPU?

# Exe 2.b: Which is the fastest CPU?

**Recall:** "X is n times faster than Y" means

$$\frac{Performance(X)}{Performance(Y)} = \frac{Exe(Y)}{Exe(X)}$$

$$CPU \text{ time} = \left( \sum_{i=1}^n IC_i \times CPI_i \right) \times \text{Clock cycle time}$$

# Exe 2.b: Which is the fastest CPU?

$$\frac{EXE_{CPU_1}}{EXE_{CPU_2}}$$



## Exe 2.b: Which is the fastest CPU?

$$\frac{EXE_{CPU_1}}{EXE_{CPU_2}} = \left( \frac{IC_1 * CPI_1}{F_1} \right) * \left( \frac{F_2}{IC_2 * CPI_2} \right)$$

## Exe 2.b: Which is the fastest CPU?

$$\frac{EXE_{CPU_1}}{EXE_{CPU_2}} = \left( \frac{IC_1 * CPI_1}{F_1} \right) * \left( \frac{F_2}{IC_2 * CPI_2} \right)$$

$$\frac{IC_1 * CPI_1 * F_2}{IC_2 * CPI_2 * F_1} = \frac{CPI_1 * F_2}{CPI_2 * F_1}$$

$$= \frac{2.7 * 700MHz}{2.4 * 500MHz} = \frac{1890}{1200} = 1.575$$

CPU2 is 1.575 faster than CPU1



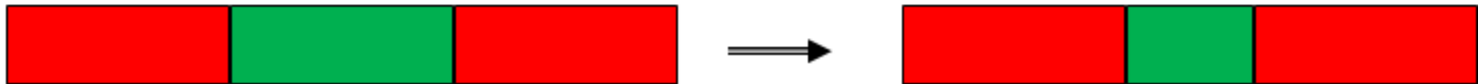
# Exe 3: Amdahl's Law

## Recall

$$Speedup_{overall} = \frac{Executiontime_{old}}{Executiontime_{new}} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

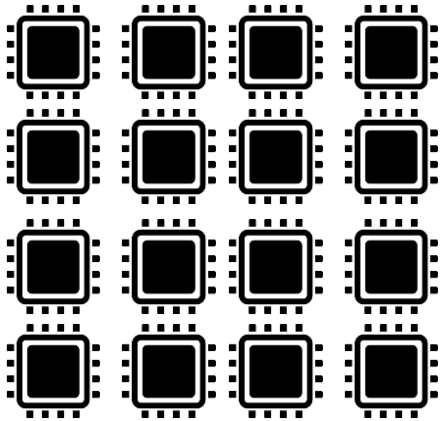
Best you could ever hope to do:

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced})}$$



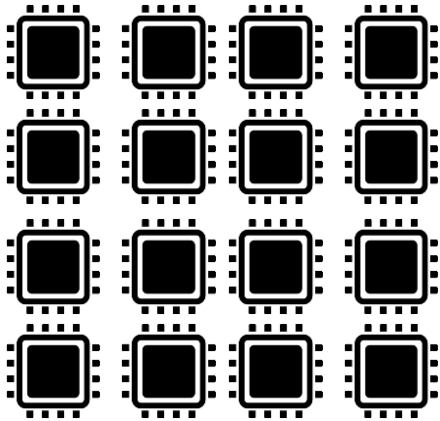
Example of good trade-off: [In-Datcenter Performance Analysis of a Tensor Processing Unit](#)

# Exe 3 : Amdahl's Law





# Exe 3 : Amdahl's Law



# Exe 3 : Amdahl's Law

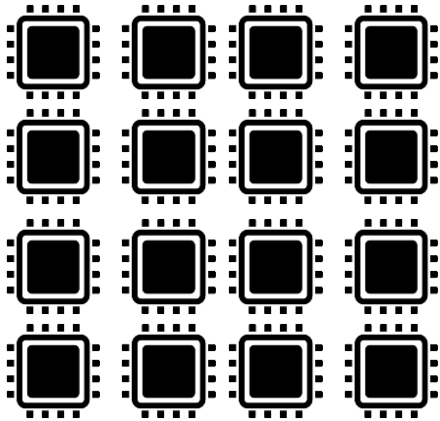


Image Processing task on FPGA is  $2.86x^{[1]}$  times faster  
100W (TDP) vs 30.85 W

# Exe 3 : Questions

A. With what percentage of processing will adding FPGA result in a speedup of 2?

A. (At home, if you want, Enjoy :D ) Draw a graph that plots the speedup as a percentage of the computation spent performing the image processing task. Label y-axis “Net speedup” and x-axis “Percent image processing”



# Exe 3 : Questions

A. With what **percentage of processing** will adding FPGA result in a **speedup of 2**?

A. (At home, if you want, Enjoy :D ) Draw a graph that plots the speedup as a percentage of the computation spent performing the image processing task. Label y-axis “Net speedup” and x-axis “Percent image processing”

# Exe 3.a : Find the processing fraction

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

# Exe 3.a : Find the processing fraction

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

# Exe 3.a : Find the processing fraction

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}}$$

# Exe 3.a : Find the processing fraction

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}}$$

$$\frac{1}{2} = 1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}$$

$$\frac{1}{2} = \frac{2.86 - 2.86Fraction_{enhanced} + Fraction_{enhanced}}{2.86}$$

$$\frac{2.86}{2} = 2.86 - 2.86Fraction_{enhanced} + Fraction_{enhanced}$$

$$1.86Fraction_{enhanced} = 1.43$$

# Exe 3.a : Find the processing fraction

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}}$$

$$\frac{1}{2} = 1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}$$

$$\frac{1}{2} = \frac{2.86 - 2.86Fraction_{enhanced} + Fraction_{enhanced}}{2.86}$$

$$\frac{2.86}{2} = 2.86 - 2.86Fraction_{enhanced} + Fraction_{enhanced}$$

$$1.86Fraction_{enhanced} = 1.43$$

$$Fraction_{enhanced} = 0.768 = 76.8\%$$

# Exe 3.a : Find the processing fraction

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$2 = \frac{1}{1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}}$$

$$\frac{1}{2} = 1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{2.86}$$

$$\frac{1}{2} = \frac{2.86 - 2.86Fraction_{enhanced} + Fraction_{enhanced}}{2.86}$$

$$\frac{2.86}{2} = 2.86 - 2.86Fraction_{enhanced} + Fraction_{enhanced}$$

$$1.86Fraction_{enhanced} = 1.43$$

$$Fraction_{enhanced} = 0.768 = 76.8\%$$

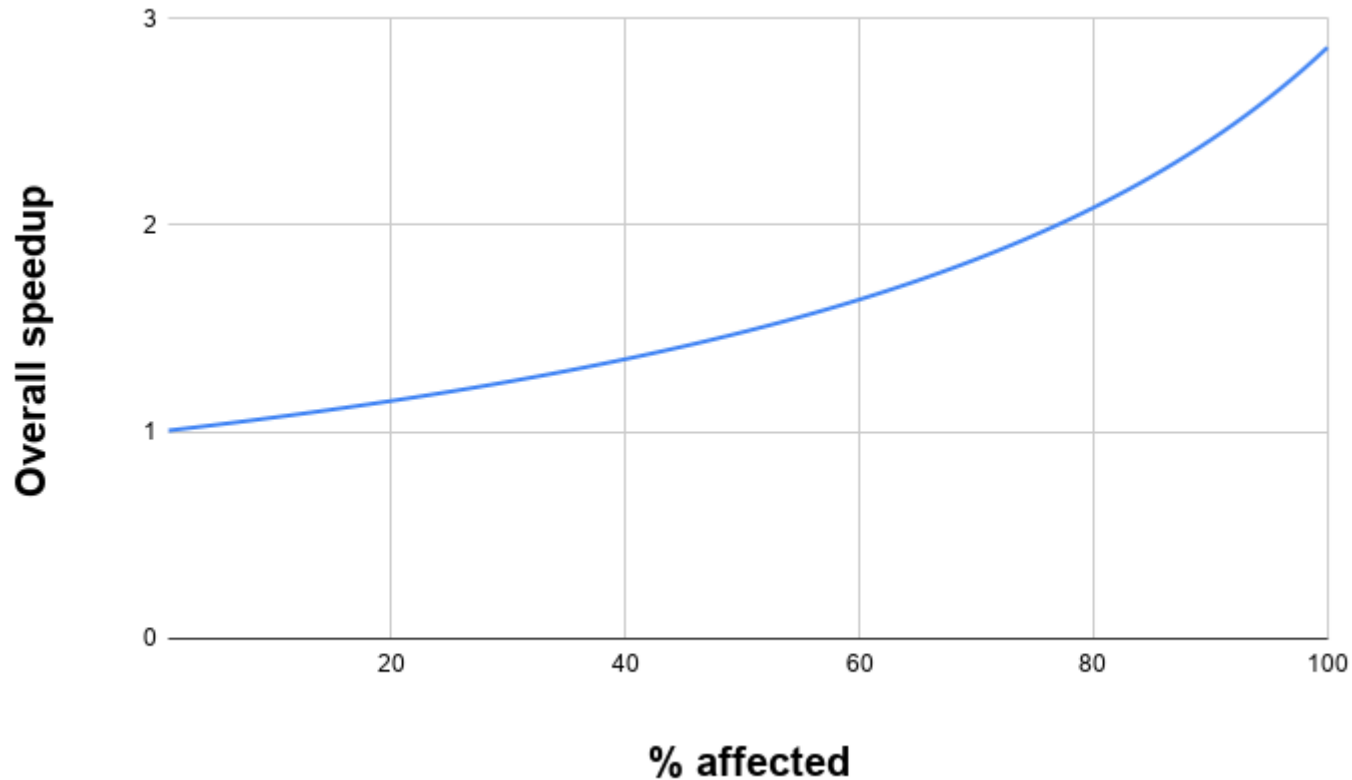
# Exe 3: Questions

A. With what percentage of processing will adding FPGA result in a speedup of 2?

A. (At home, if you want, Enjoy :D ) Draw a graph that plots the speedup as a percentage of the computation spent performing the image processing task. Label y-axis “Net speedup” and x-axis “Percent image processing”



# Exe 3.b : Graph solution



$$100 / ((100-x)+x/2.86)$$

# Performance Scaling

How does the **overall performance scale** if we further **increase the speedup**?

Let's consider an application where the **number** of used **processors/threads linearly increases** the **performance** of the parallelizable portion

# Modern Problems Require Modern Formulae

$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced} + \frac{Fraction_{enhanced}}{Speedup_{enhanced}})}$$

$$\mathbf{speedup_{overall} = 1/(s + p/N)}$$

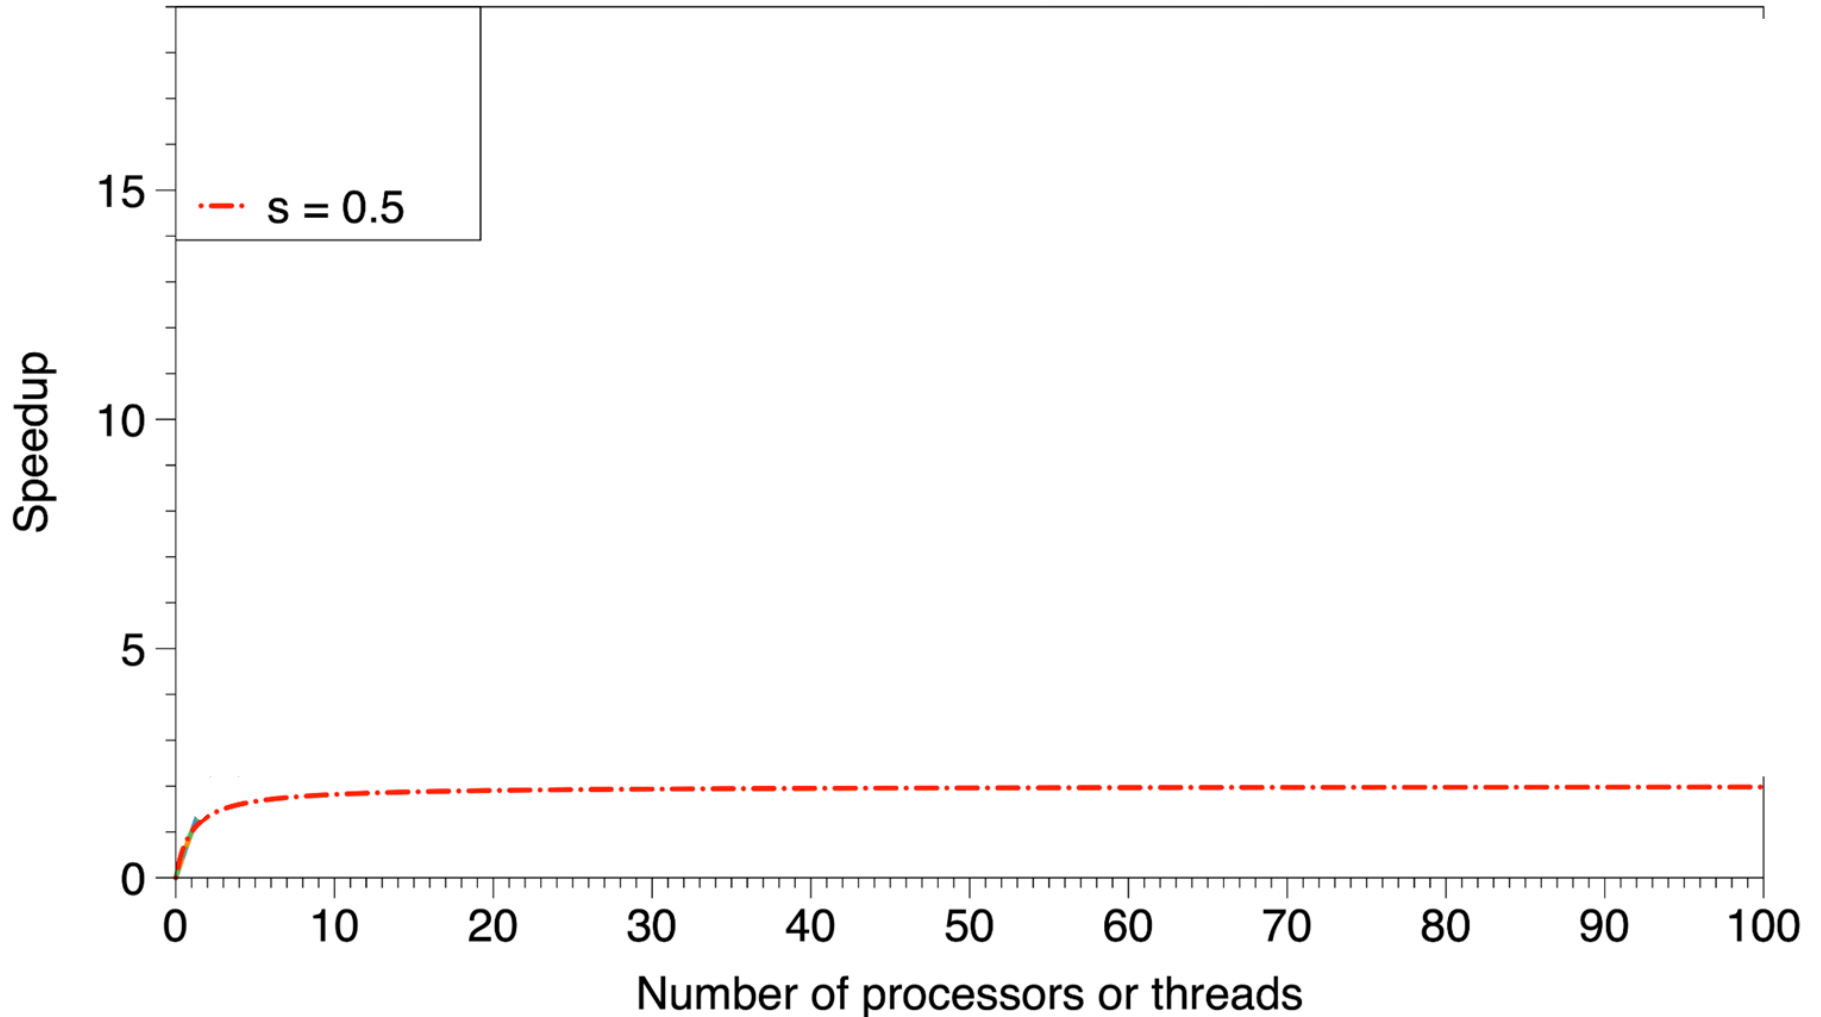
$s$  = serial part =  $1 - Fraction_{enhanced}$

$p = 1 - s$  = parallelizable part =  $Fraction_{enhanced}$

$N$  = number of processors or threads =  $Speedup_{enhanced}$

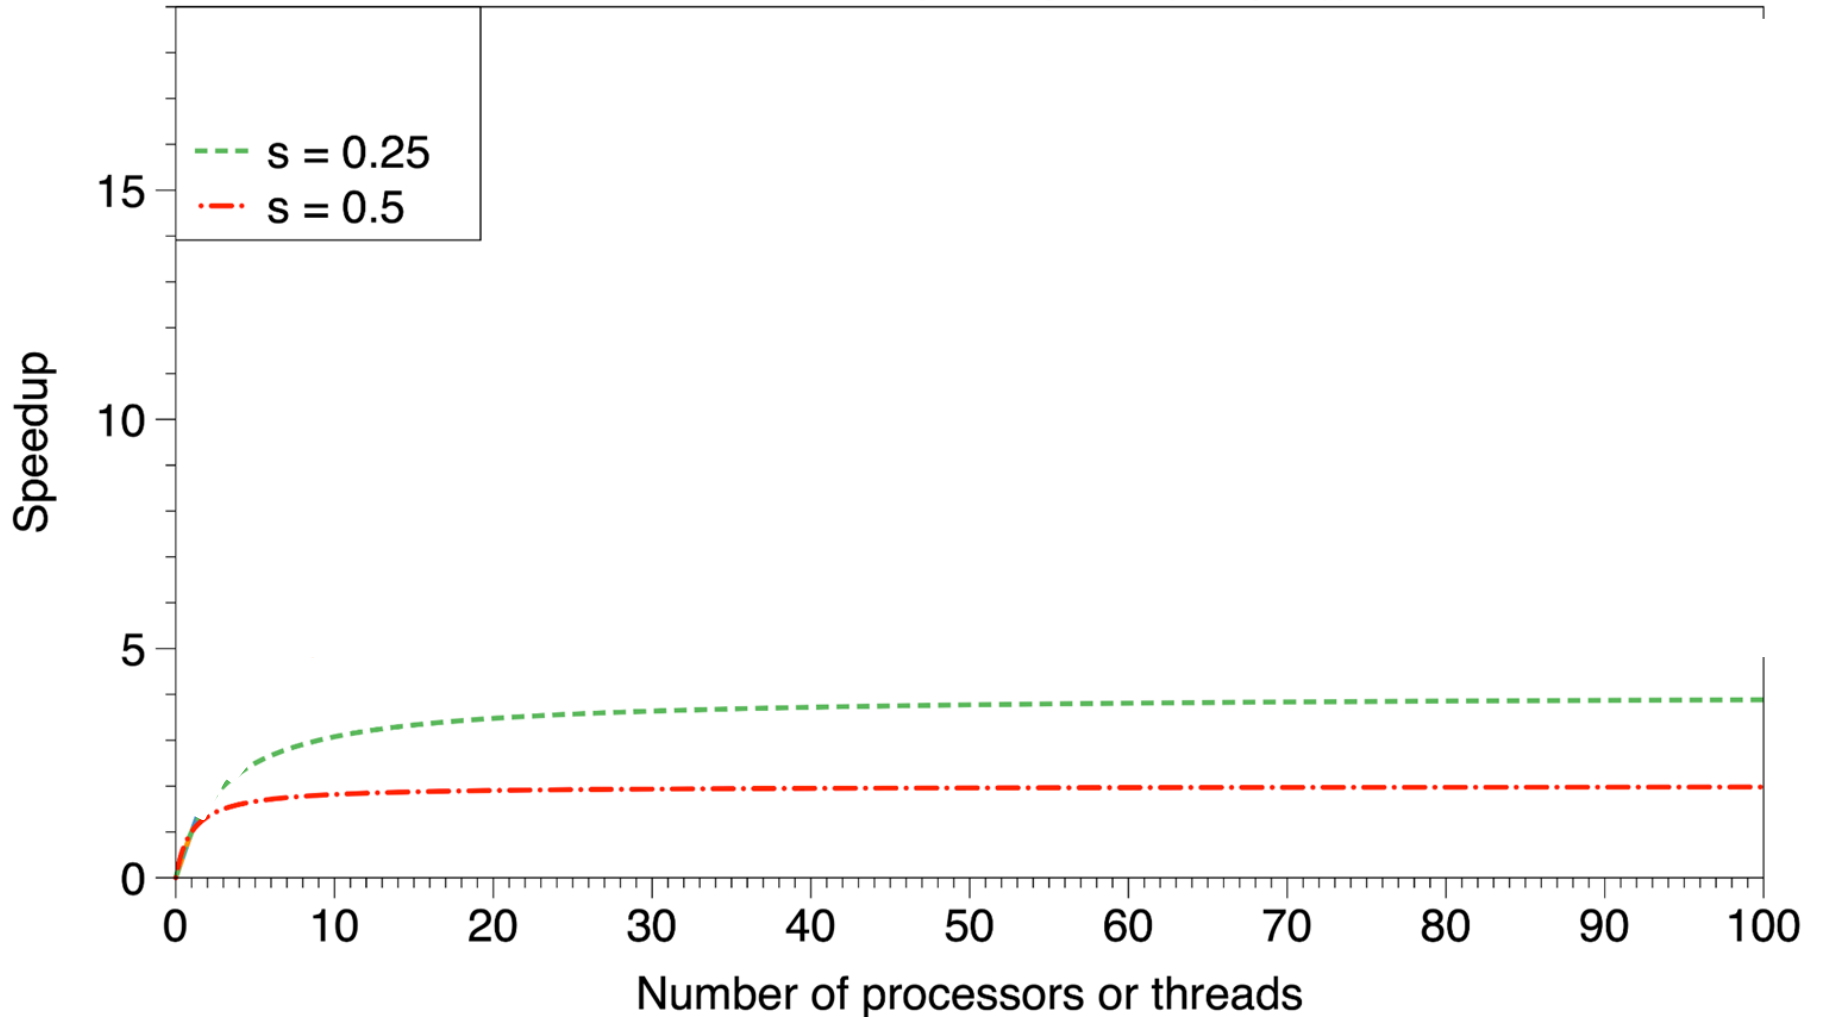
# Speedup Scaling vs $N/s$ Scaling

Amdahl's law - speedup =  $1/(s + p/N)$



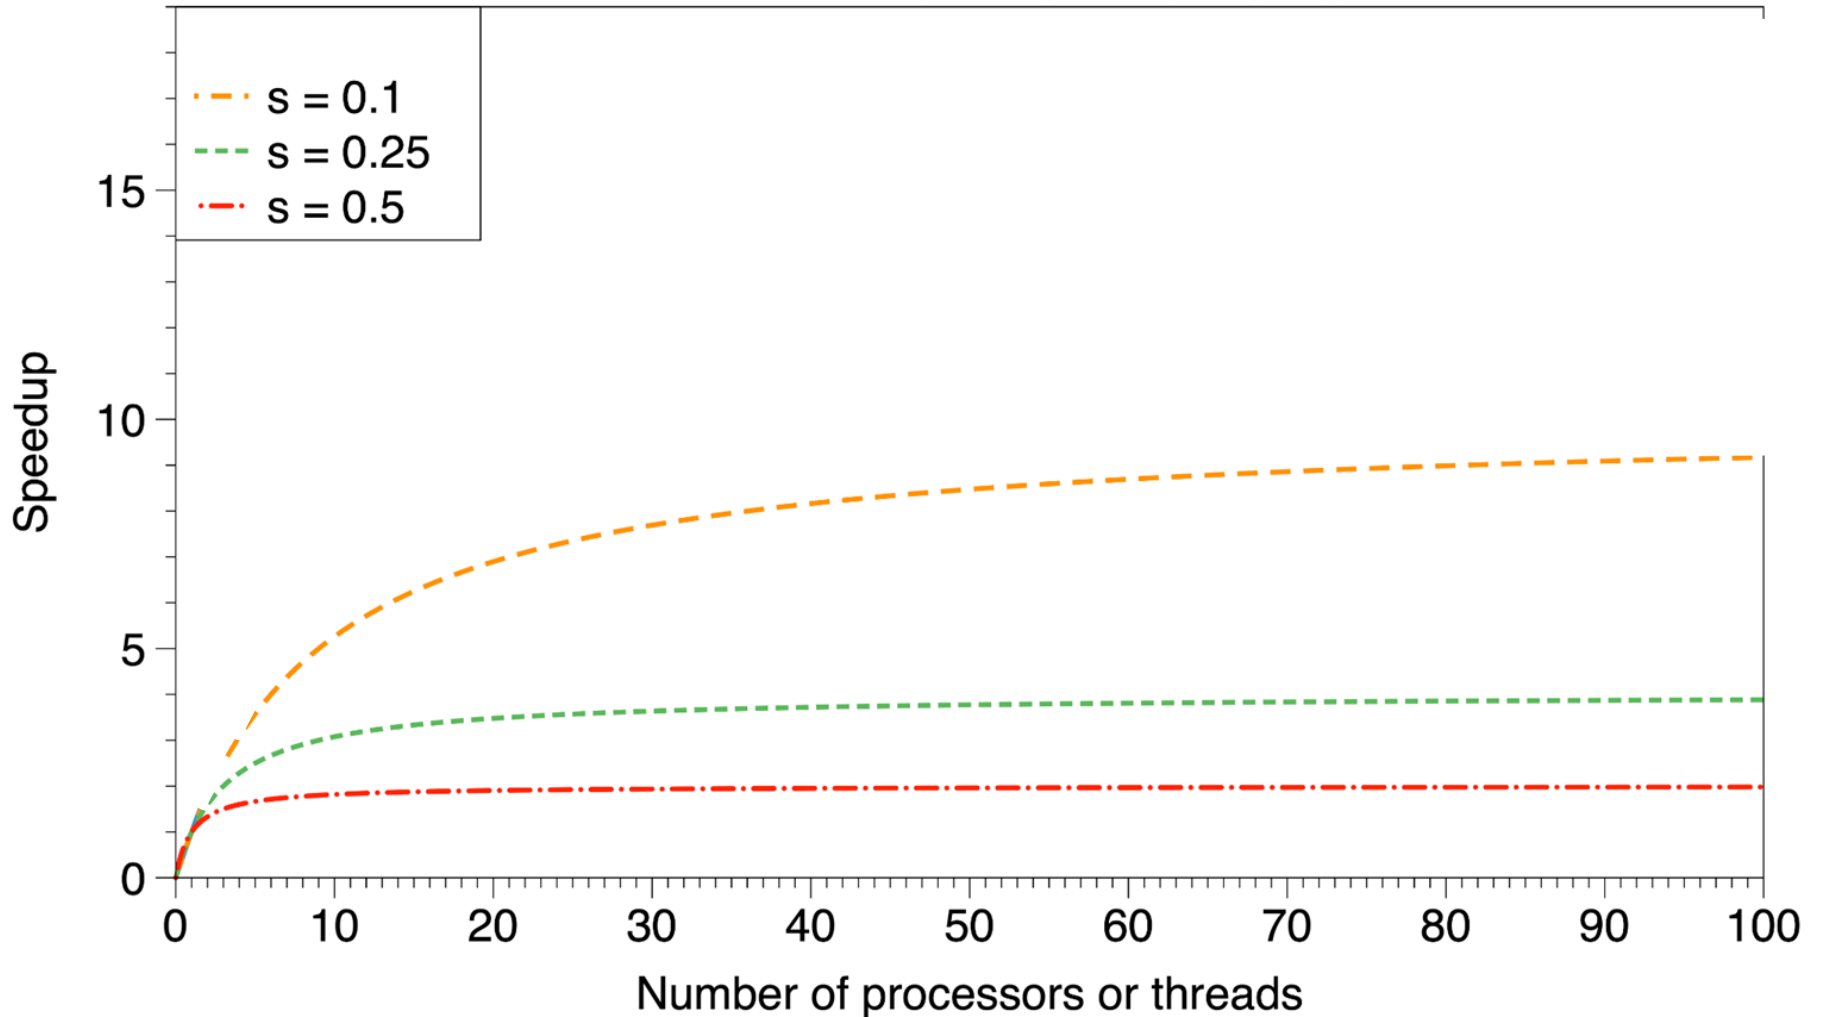
# Speedup Scaling vs $N/s$ Scaling

Amdahl's law - speedup =  $1/(s + p/N)$



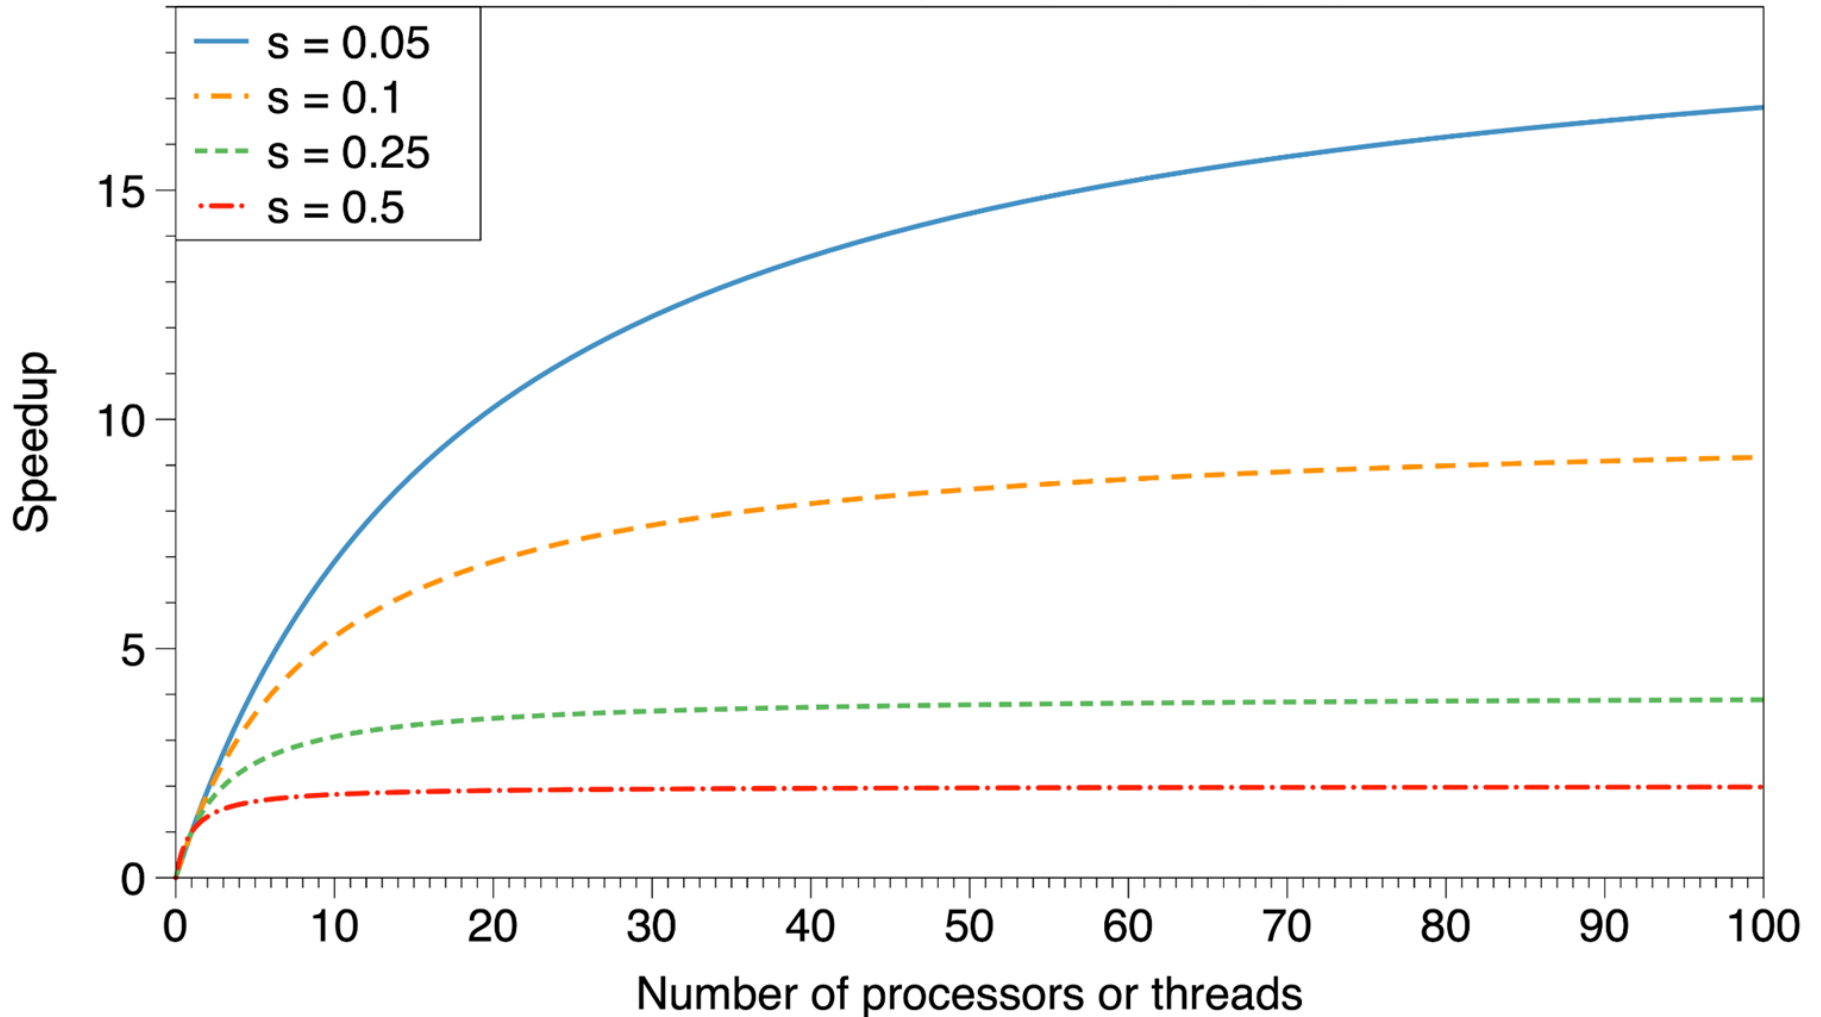
# Speedup Scaling vs $N/s$ Scaling

Amdahl's law - speedup =  $1/(s + p/N)$



# Speedup Scaling vs $N/s$ Scaling

Amdahl's law - speedup =  $1/(s + p/N)$



# To Sum Up

**Amdahl's law** states that, for a **fixed** problem, the upper limit of speedup is determined by the serial fraction of the code -> **strong scaling**

$$speedup = 1/(s + p/N)$$



# Something More about Performance Scaling

Gustafson's law -> weak scaling

$$\text{scaled speedup} = s + p \times N$$

If you want more information **JUST FOR CURIOSITY:**

- <https://www.kth.se/blogs/pdc/2018/11/scalability-strong-and-weak-scaling/>
- <https://dl.acm.org/doi/10.1145/42411.42415>



# Exe 4 : Pipelining and Performance

# Recall and some ref

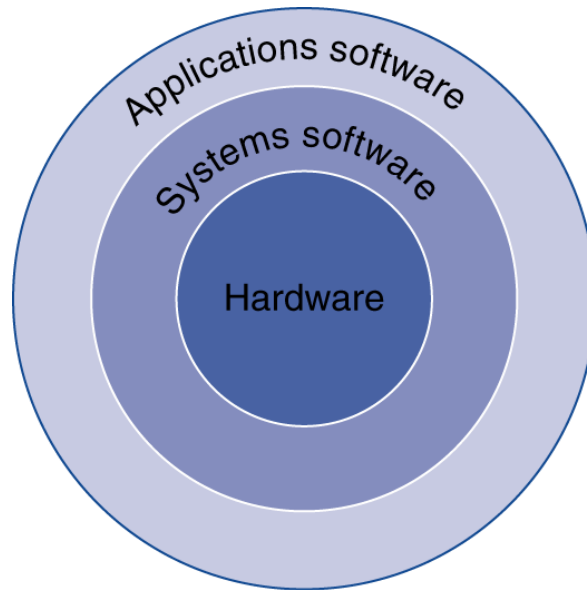
Web simulators: (Not tested) [MIPS simulator from unisi](#) and [RISC-V simulator from unisi](#)

<https://tinyurl.com/aca-grid24>

<https://docs.google.com/spreadsheets/d/1vqTGDF7TNgOfMDSYCP2vqisW3uBdiuXvvZo0hvc4tA/edit?usp=sharing>

# Recall and some ref

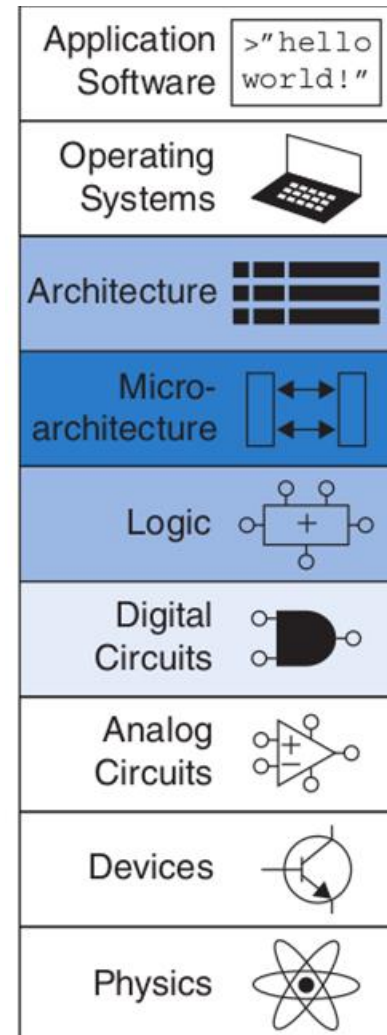
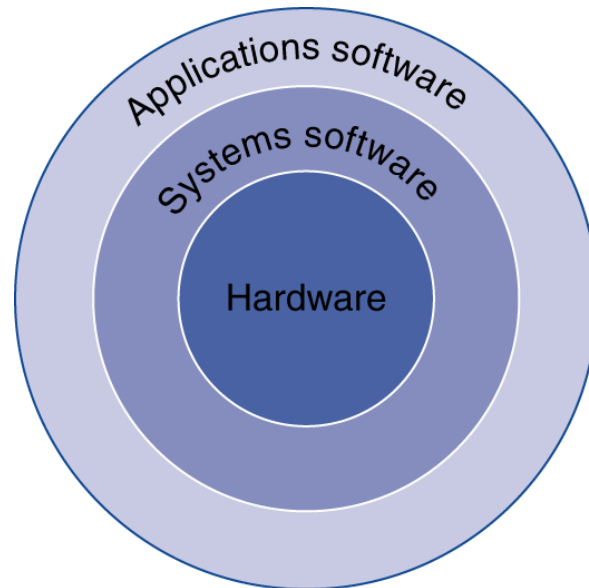
<https://tinyurl.com/aca-grid24>



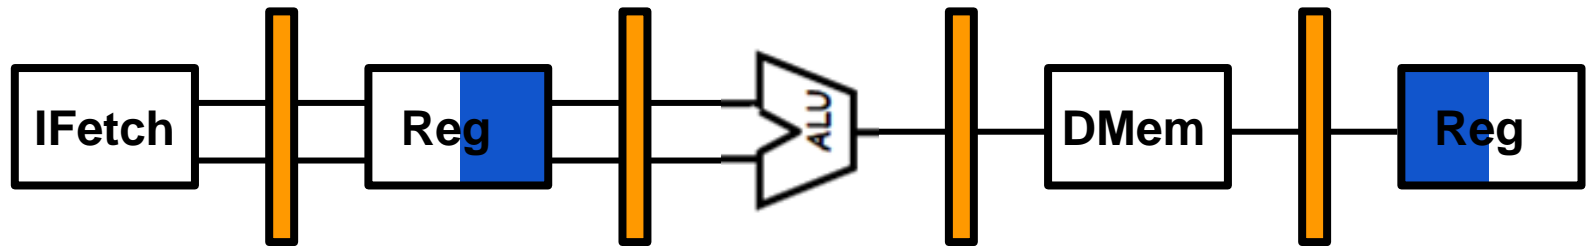


# Recall and some ref

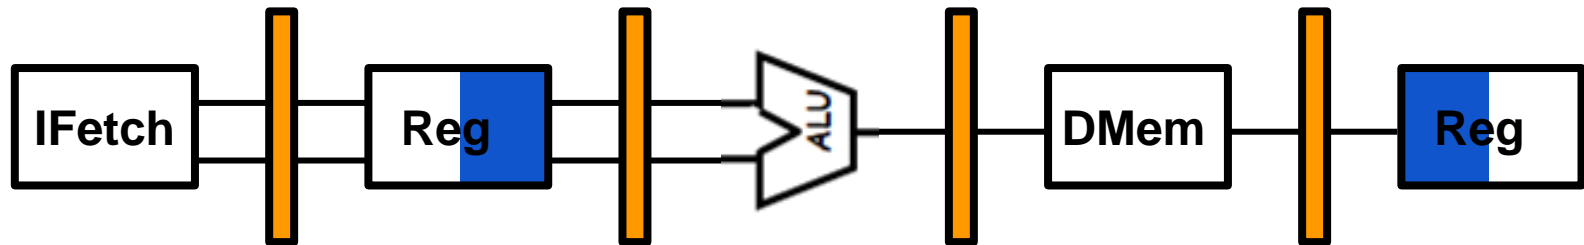
<https://tinyurl.com/aca-grid24>



# Exe 4 : Pipelining and Performance



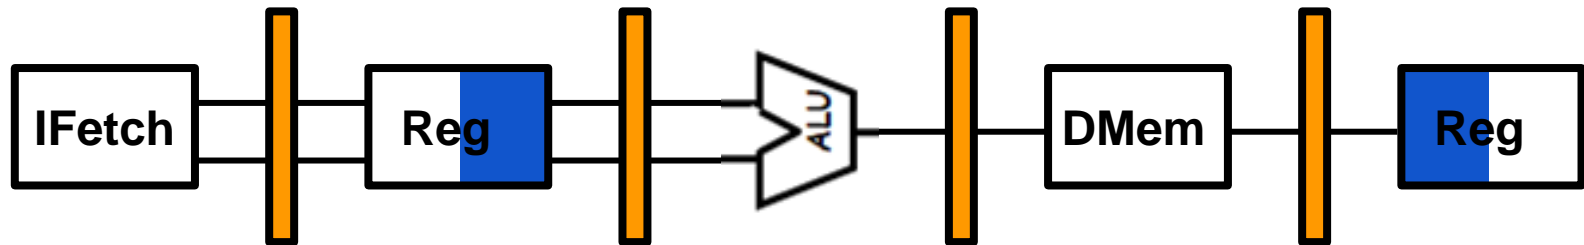
# Exe 4 : Pipelining and Performance



lw      \$1, OFF(\$2)  
addi    \$3, \$1, 4  
sub     \$4, \$1, \$2  
addi    \$2, \$1, -8  
sw      \$4, OFF(\$2)



# Exe 4 : Pipelining and Performance



**lw**      **\$1, OFF(\$2)**  
**addi**   **\$3, \$1, 4**  
**sub**     **\$4, \$1, \$2**  
**addi**   **\$2, \$1, -8**  
**sw**     **\$4, OFF(\$2)**

IF	ID	EX	ME	WB
Instruction Fetch	Instruction Decode	Execution	Memory Access	Write Back

ALU Instructions: **op \$x, \$y, \$z**

Instr. Fetch & PC Increm.	Read of Source Regs. \$y and \$z	ALU Op. (\$y op \$z)		Write Back Destin. Reg. \$x
---------------------------	----------------------------------	----------------------	--	-----------------------------

Load Instructions: **lw \$x, offset(\$y)**

Instr. Fetch & PC Increm.	Read of Base Reg. \$y	ALU Op. (\$y+offset)	Read Mem. M(\$y+offset)	Write Back Destin. Reg. \$x
---------------------------	-----------------------	----------------------	-------------------------	-----------------------------

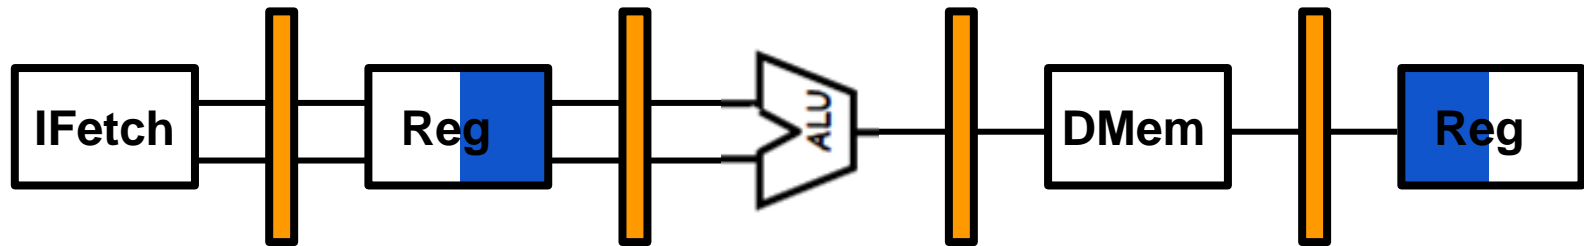
Store Instructions: **sw \$x, offset(\$y)**

Instr. Fetch & PC Increm.	Read of Base Reg. \$y & Source \$x	ALU Op. (\$y+offset)	Write Mem. M(\$y+offset)	
---------------------------	------------------------------------	----------------------	--------------------------	--

Conditional Branches: **beq \$x, \$y, offset**

Instr. Fetch & PC Increm.	Read of Source Regs. \$x and \$y	ALU Op. (\$x-\$y) & (PC+4+offset)	Write PC	
---------------------------	----------------------------------	-----------------------------------	----------	--

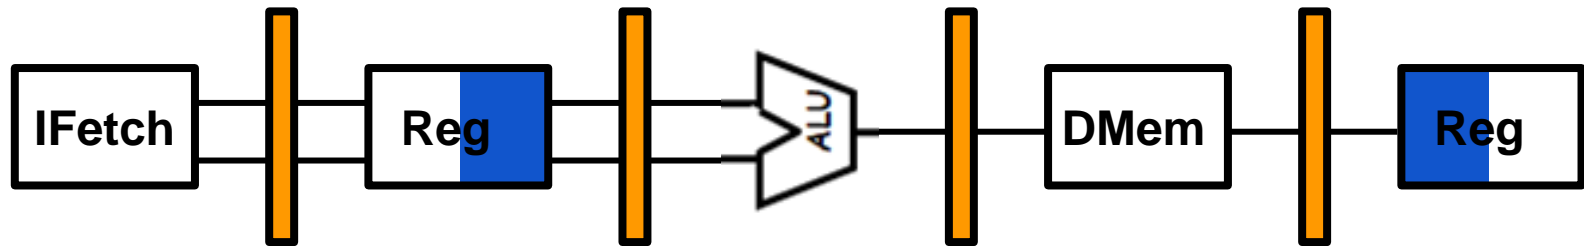
# Exe 4 : Pipelining and Performance



**lw**      **\$1, OFF(\$2)**  
**addi**   **\$3, \$1, 4**  
**sub**     **\$4, \$1, \$2**  
**addi**   **\$2, \$1, -8**  
**sw**      **\$4, OFF(\$2)**

No optimization in the MIPS pipeline (e.g., forwarding paths) just our “optimization” (i.e., RF access R/W)

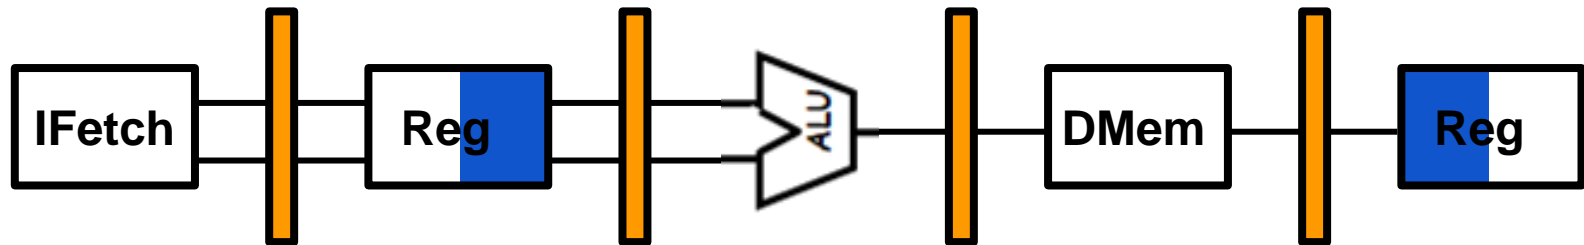
# Exe 4 : Pipelining and Performance



**lw**      **\$1, OFF(\$2)**  
**addi**    **\$3, \$1, 4**  
**sub**      **\$4, \$1, \$2**  
**addi**    **\$2, \$1, -8**  
**sw**      **\$4, OFF(\$2)**

No optimization in the MIPS pipeline (e.g., forwarding paths) just our “optimization” (i.e., RF access R/W)  
The processor has a clock cycle of 2ns

# Exe 4 : Pipelining and Performance



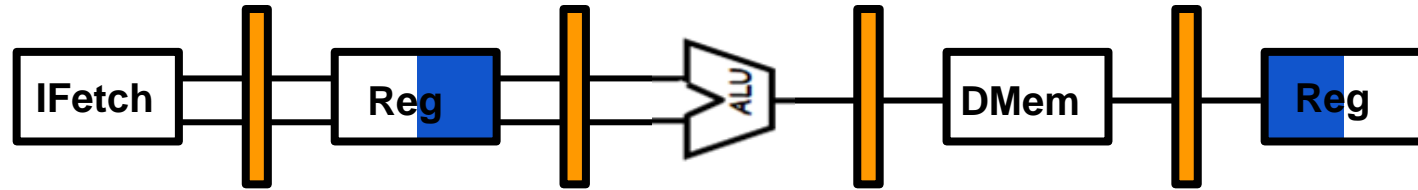
**lw**      **\$1, OFF(\$2)**  
**addi**   **\$3, \$1, 4**  
**sub**     **\$4, \$1, \$2**  
**addi**   **\$2, \$1, -8**  
**sw**      **\$4, OFF(\$2)**

No optimization in the MIPS pipeline (e.g., forwarding paths) just our “optimization” (i.e., RF access R/W)

The processor has a clock cycle of 2ns

- A. Draw the pipeline schema and highlight possible hazards
- B. Represent the real execution (Insert the stalls )
- C. Calculate IC, CPI, MIPS

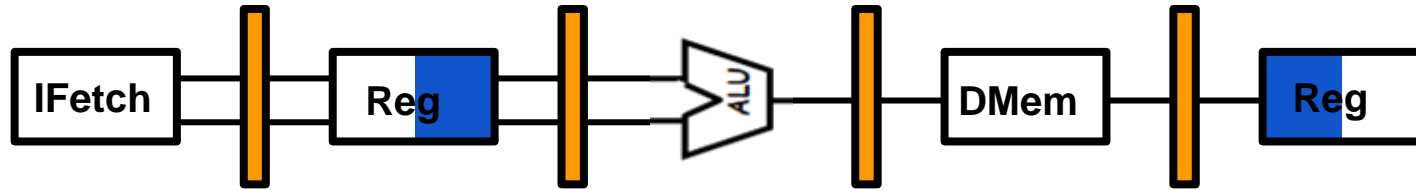
# Exe 4.a : Ideal case



CC 0

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)															
addi \$3, \$1, 4															
sub \$4, \$1, \$3															
addi \$2, \$1, -8															
sw \$5, OFF(\$2)															

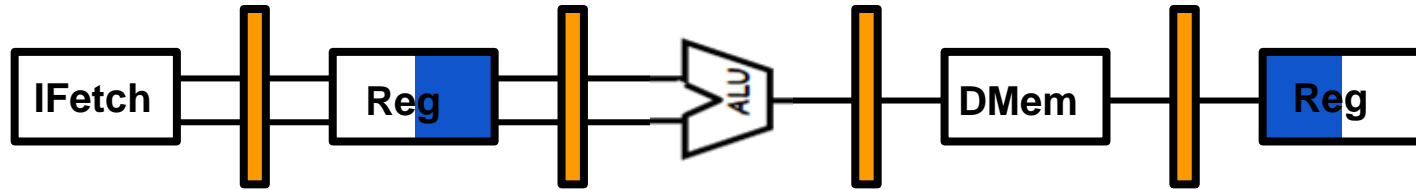
# Exe 4.a : Ideal case



CC 2

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D													
addi \$3, \$1, 4		F													
sub \$4, \$1, \$3															
addi \$2, \$1, -8															
sw \$5, OFF(\$2)															

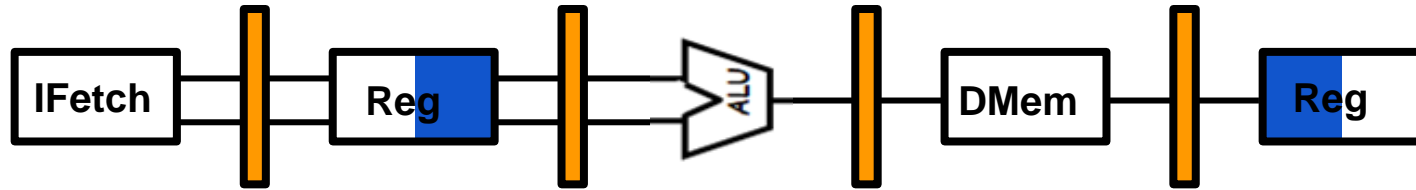
# Exe 4.a : Ideal case



CC 3

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E												
addi \$3, \$1, 4		F	D												
sub \$4, \$1, \$3			F												
addi \$2, \$1, -8															
sw \$5, OFF(\$2)															

# Exe 4.a : Ideal case



CC 9


Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	D	E	M	W									
sub \$4, \$1, \$3			F	D	E	M	W								
addi \$2, \$1, -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						



# Recall: Type of Data Hazard

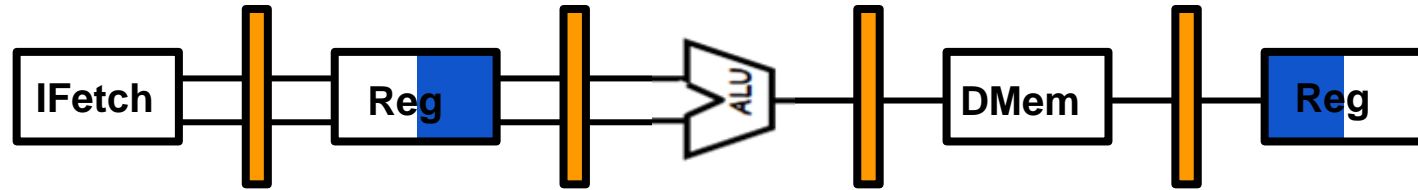
## Read After Write (RAW)

Instr<sub>j</sub> tries to read operand before Instr<sub>i</sub> writes it

 I: add **r1**, r2, r3  
J: sub r4, **r1**, r3

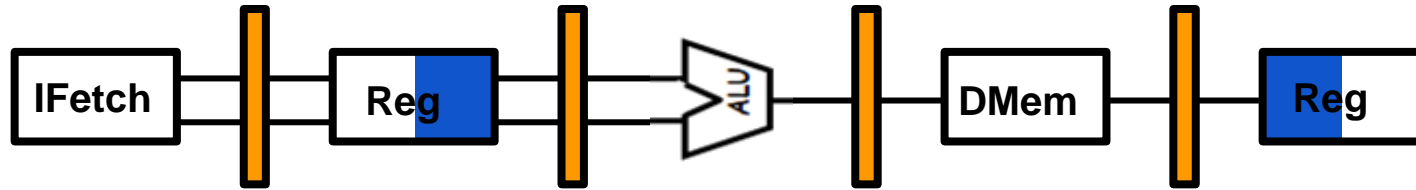
Caused by a “**Dependence**” (in compiler nomenclature). This hazard results from an actual need for communication.

# Exe 4.a : Ideal case



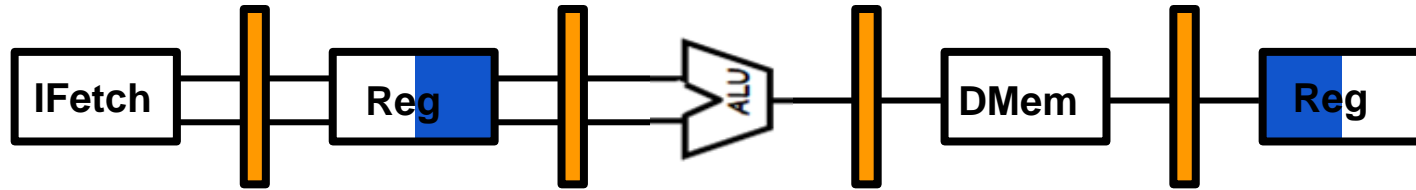
Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw <b>\$1</b> , OFF(\$2)	F	D	E	M	W										
addi \$3, <b>\$1</b> , 4		F	D	E	M	W									
sub \$4, \$1, \$3			F	D	E	M	W								
addi \$2, \$1, -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						

# Exe 4.a : The Hazard



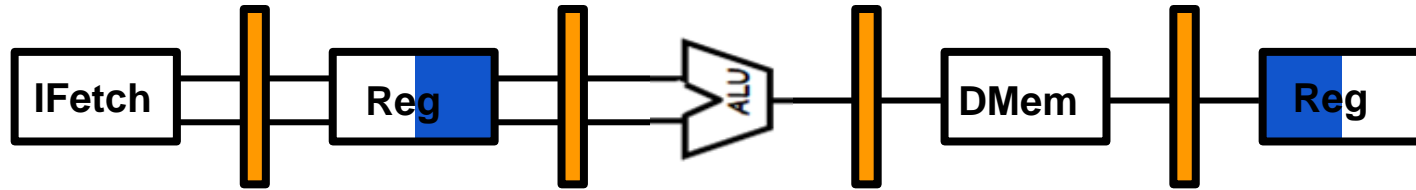
Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw <b>\$1</b> , OFF(\$2)	F	D	E	M	W										
addi \$3, <b>\$1</b> , 4		F	D	E	M	W									
sub \$4, <b>\$1</b> , \$3			F	D	E	M	W								
addi \$2, \$1, -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						

# Exe 4.a : The Hazard



Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw <b>\$1</b> , OFF(\$2)	F	D	E	M	W										
addi \$3, <b>\$1</b> , 4		F	D	E	M	W									
sub \$4, <b>\$1</b> , \$3			F	D	E	M	W								
addi \$2, <b>\$1</b> , -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						

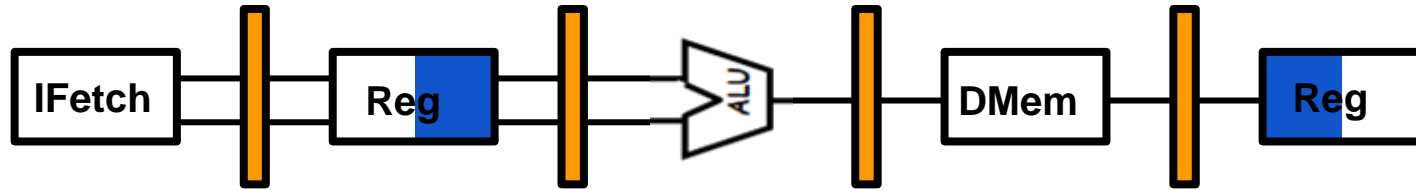
# Exe 4.a : The Hazard



Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	D	E	M	W									
sub \$4, \$1, \$3			F	D	E	M	W								
addi \$2, \$1, -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						

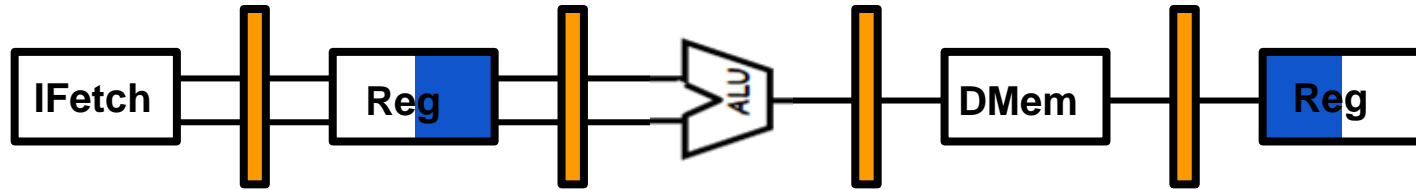
Not a real Hazard

# Exe 4.a : The Hazard



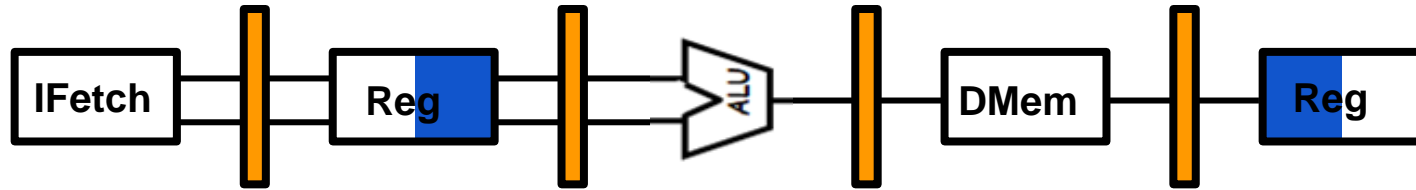
Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw <b>\$1</b> , OFF(\$2)	F	D	E	M	W										
addi \$3, <b>\$1</b> , 4		F	D	E	M	W									
sub \$4, <b>\$1</b> , \$3			F	D	E	M	W								
addi \$2, <b>\$1</b> , -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						

# Exe 4.a : The Hazard



Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw <b>\$1</b> , OFF(\$2)	F	D	E	M	W										
addi <b>\$3</b> , <b>\$1</b> , 4		F	D	E	M	W									
sub \$4, <b>\$1</b> , <b>\$3</b>			F	D	E	M	W								
addi \$2, <b>\$1</b> , -8				F	D	E	M	W							
sw \$5, OFF(\$2)					F	D	E	M	W						

# Exe 4.a : The Hazard



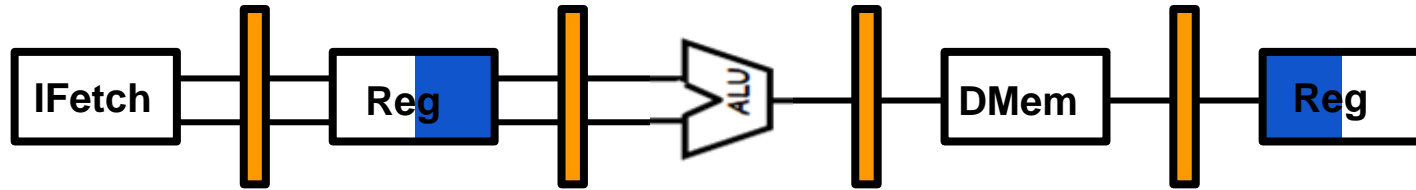
Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw <b>\$1</b> , OFF(\$2)	F	D	E	M	W										
addi <b>\$3</b> , <b>\$1</b> , 4		F	D	E	M	W									
sub \$4, <b>\$1</b> , <b>\$3</b>			F	D	E	M	W								
addi <b>\$2</b> , <b>\$1</b> , -8				F	D	E	M	W							
sw <b>\$5</b> , OFF( <b>\$2</b> )					F	D	E	M	W						



# Recall: Data Hazards: Possible Solutions

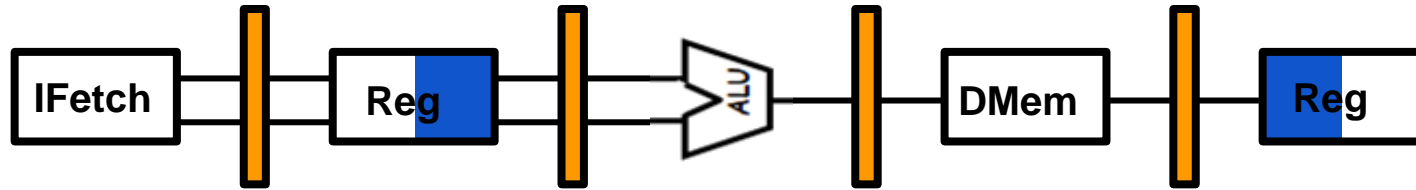
- **Compilation Techniques:**
  - Insertion of nop (no operation) instructions
  - Instructions Scheduling to avoid that correlating instructions are too close
    - The compiler tries to insert independent instructions among correlating instructions
    - When the compiler does not find independent instructions, it insert nops.
- **Hardware Techniques:**
  - Insertion of “bubbles” or stalls in the pipeline
  - Data Forwarding or Bypassing

# Exe 4.b : Bubble insertion



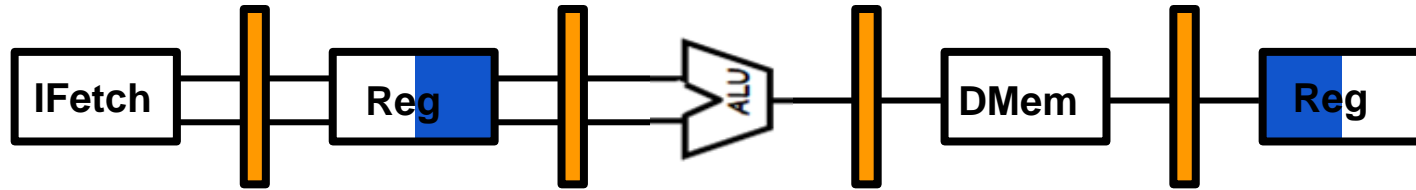
Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	D	E	M	W						
addi \$2, \$1, -8						F	D	E	M	W					
sw \$5, OFF(\$2)							F	D	E	M	W				

# Exe 4.b : Bubble insertion



Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	D	E	M	W		

# Exe 4.b : Bubble insertion



Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

# Exe 4.c : Performance

Calculate IC, CPI, MIPS

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

# Exe 4.c : Performance

Calculate IC, CPI, MIPS

$$IC = 5$$

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

# Exe 4.c : Performance

Calculate IC, CPI, MIPS

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

$$IC = 5$$

$$CPI = \frac{CCs}{IC}$$

# Exe 4.c : Performance

Calculate IC, CPI, MIPS

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

$$IC = 5$$

$$CPI = \frac{CCs}{IC} = \frac{15}{5}$$



# Recall: MIPS - Exe 4.c : Performance

Calculate IC, CPI, MIPS

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

$$IC = 5$$

$$CPI = \frac{CCs}{IC} = \frac{15}{5}$$

$$MIPS = \frac{ClockFrequency}{CPI * 10^6}$$

# Exe 4.c : Performance

Calculate IC, CPI, MIPS

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$1, OFF(\$2)	F	D	E	M	W										
addi \$3, \$1, 4		F	Stall	Stall	D	E	M	W							
sub \$4, \$1, \$3					F	Stall	Stall	D	E	M	W				
addi \$2, \$1, -8								F	D	E	M	W			
sw \$5, OFF(\$2)									F	Stall	Stall	D	E	M	W

$$IC = 5$$

$$CPI = \frac{CCs}{IC} = \frac{15}{5}$$

$$MIPS = \frac{ClockFrequency}{CPI * 10^6} = \frac{\frac{1}{2} * 10^9}{3 * 10^6} = 166$$



# ACA24 Research Projects

Deliverables:

1) Repository with source code; 2) Presentation; 3) Report.

Other details differ for each project

Main themes:

- FPGA-based Accelerator Design
- Hardware/Software System Co-Design
- System Design for High-Performance Memory Subsystems
- Superscalar RISC-V Architectures on FPGA prototypes
- High-Performance Regex Matching with Vector or Spatial Architectures
- Static Parallelization Techniques on VLIW AIE
- System design for Future Datacenter Networks
- ... ???



# Thank you for your attention

## Questions?

Davide Conficconi <davide.conficconi@polimi.it>

### Acknowledgements

E. Del Sozzo, Marco D. Santambrogio, D. Sciuto

Part of this material comes from:

- Hennessy and Patterson [Turing Lecture](#)
- “Computer Organization and Design” and “Computer Architecture A Quantitative Approach” Patterson and Hennessy books
- “Digital Design and Computer Architecture” Harris and Harris

and are ***properties of their respective owners***