POLITECNICO DI MILANO

Computing Infrastructures





The Datacenter as a Computer: Building and Infrastructures

The topics of the course: what are we going to see today?

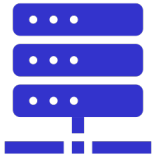

A. **HW Infrastructures**:
- **System-level**: Computing Infrastructures and Data Center Architectures, Rack/Structure;
- **Node-level**: Server (computation, HW accelerators), Storage (Type, technology), Networking (architecture and technology);
- **Building-level**: **Cooling systems, power supply, failure recovery**

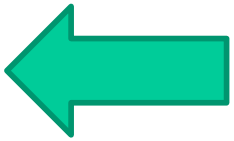

B. **SW Infrastructures**:
- **Virtualization**: Process/System VM, Virtualization Mechanisms (Hypervisor, Para/Full virtualization)
- **Computing Architectures**: Cloud Computing (types, characteristics), Edge/Fog Computing, X-as-a service
- **Machine and deep learning-as-a-service**

C. **Methods**:
- **Reliability and availability of datacenters** (definition, fundamental laws, RBDs)
- **Disk performance** (Type, Performance, RAID)
- **Scalability and performance of datacenters** (definitions, fundamental laws, queuing network theory)

BUILDING AND INFRASTRUCTURE

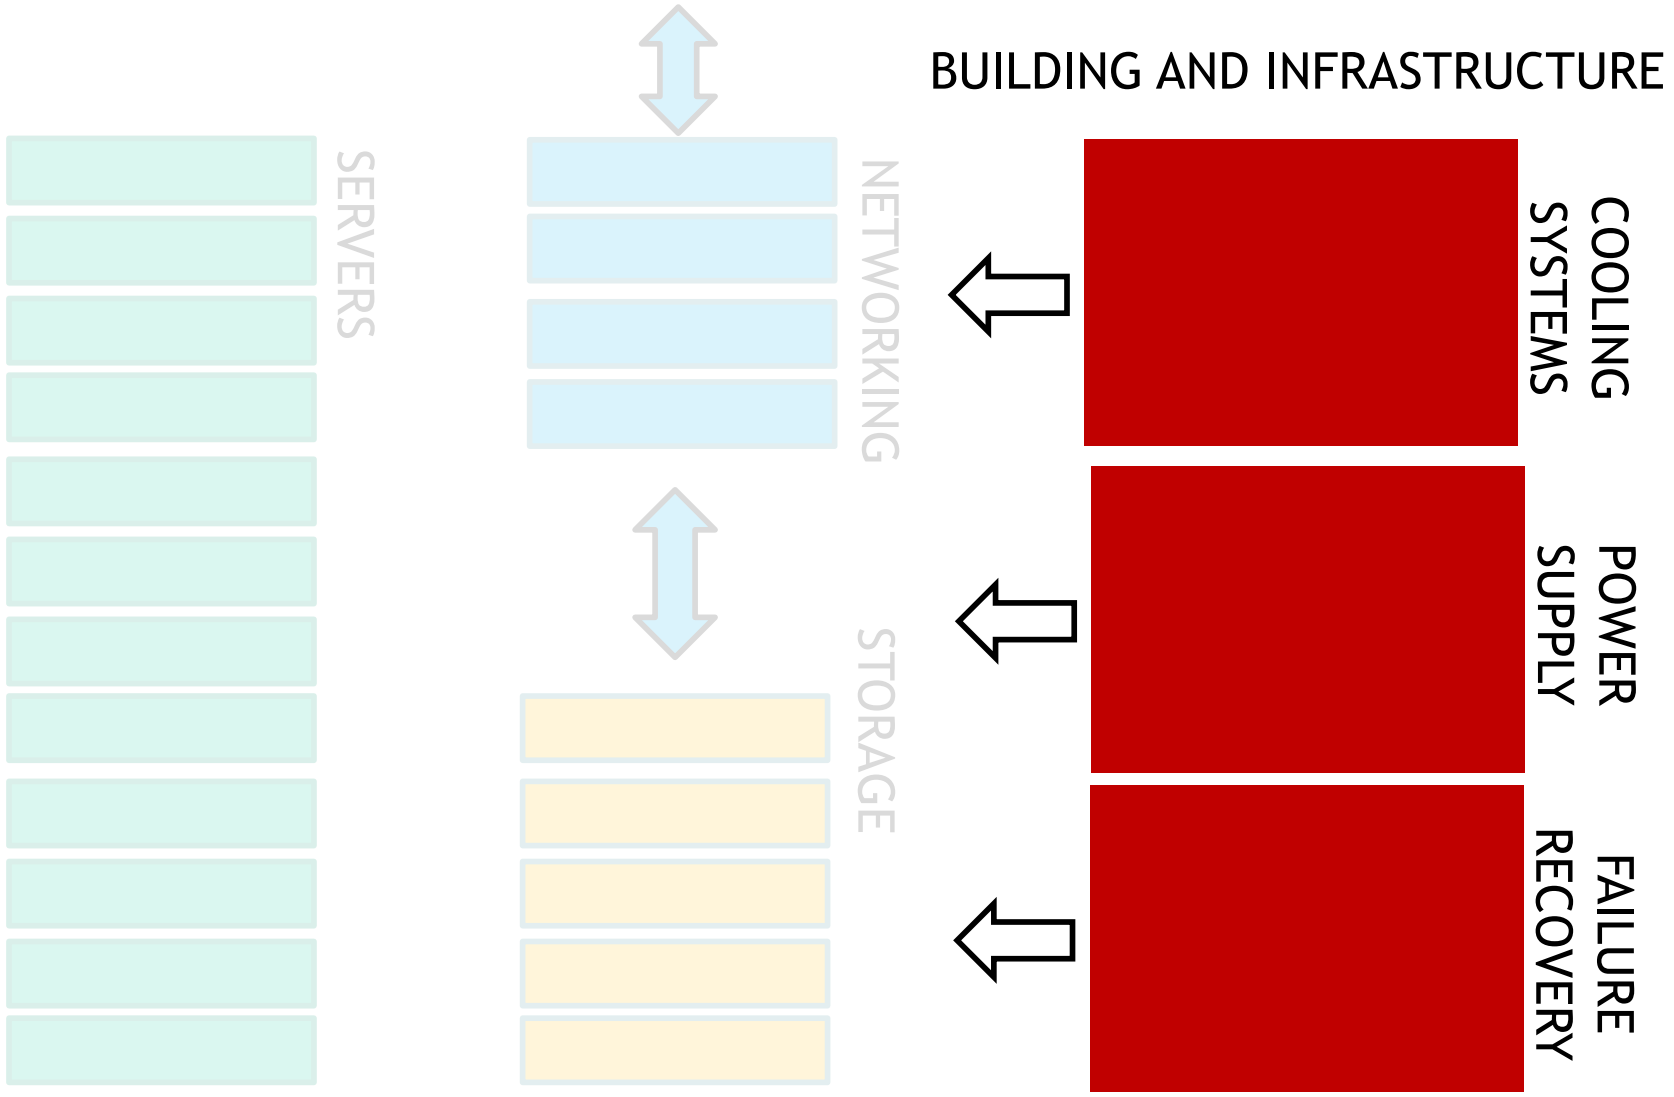

Aerial view of a Google data center campus in Iowa (US)

A Google data center building

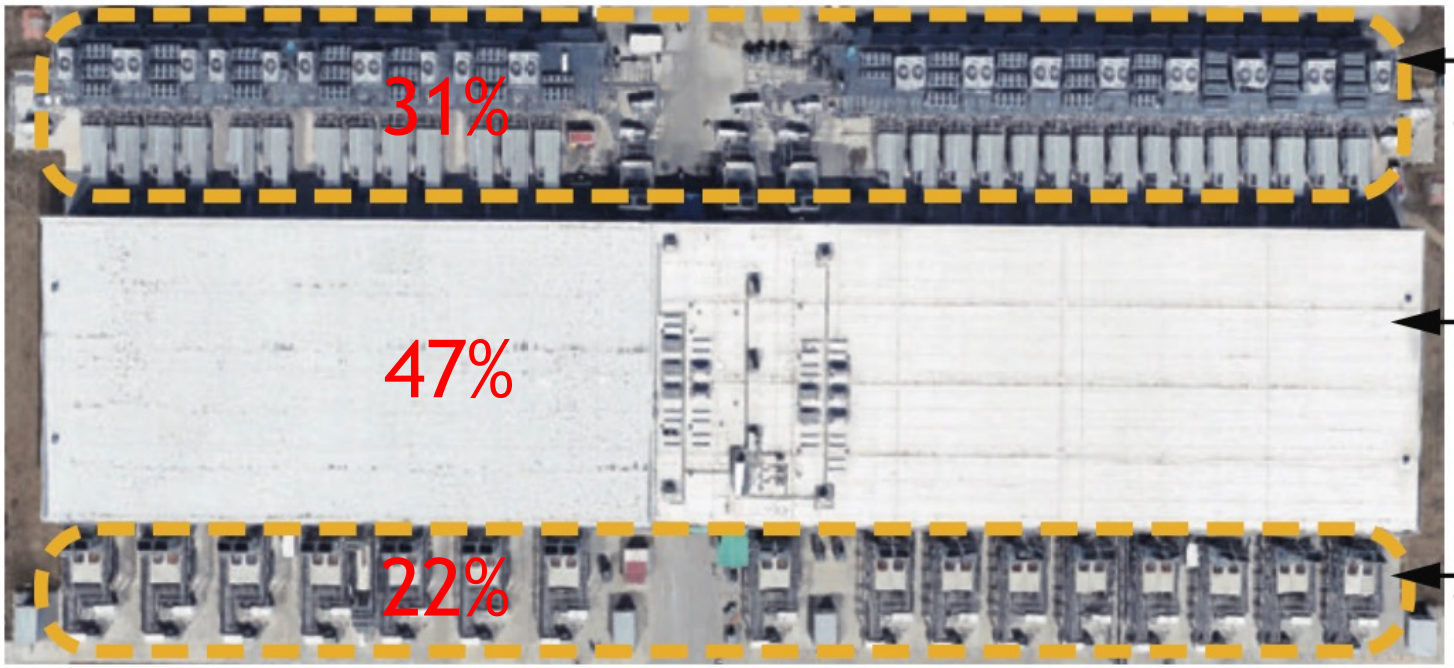

The main components of a typical data center

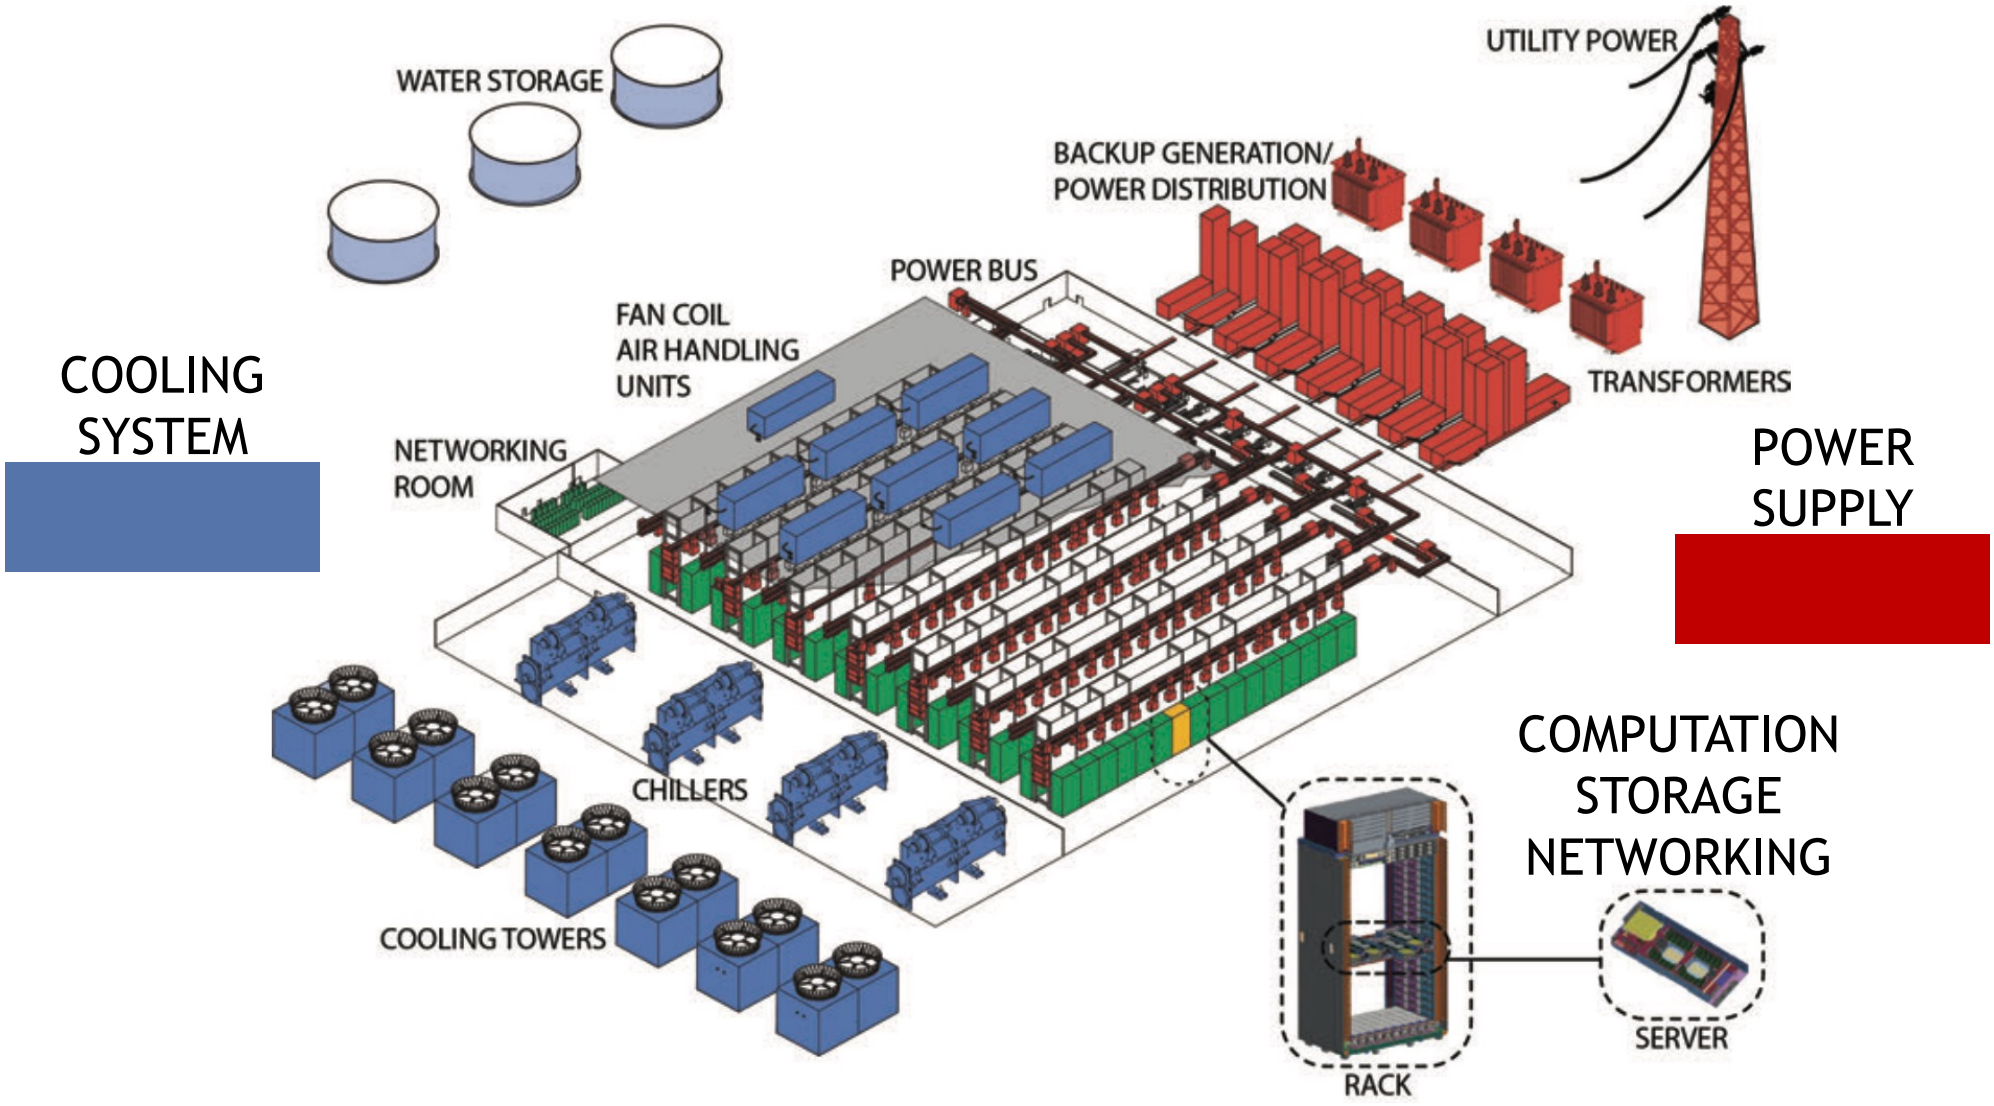

Image taken from Barroso

Not just computation, storage and networking

WSC has other important components related to **power delivery**, **cooling**, and **building infrastructure** that also need to be considered

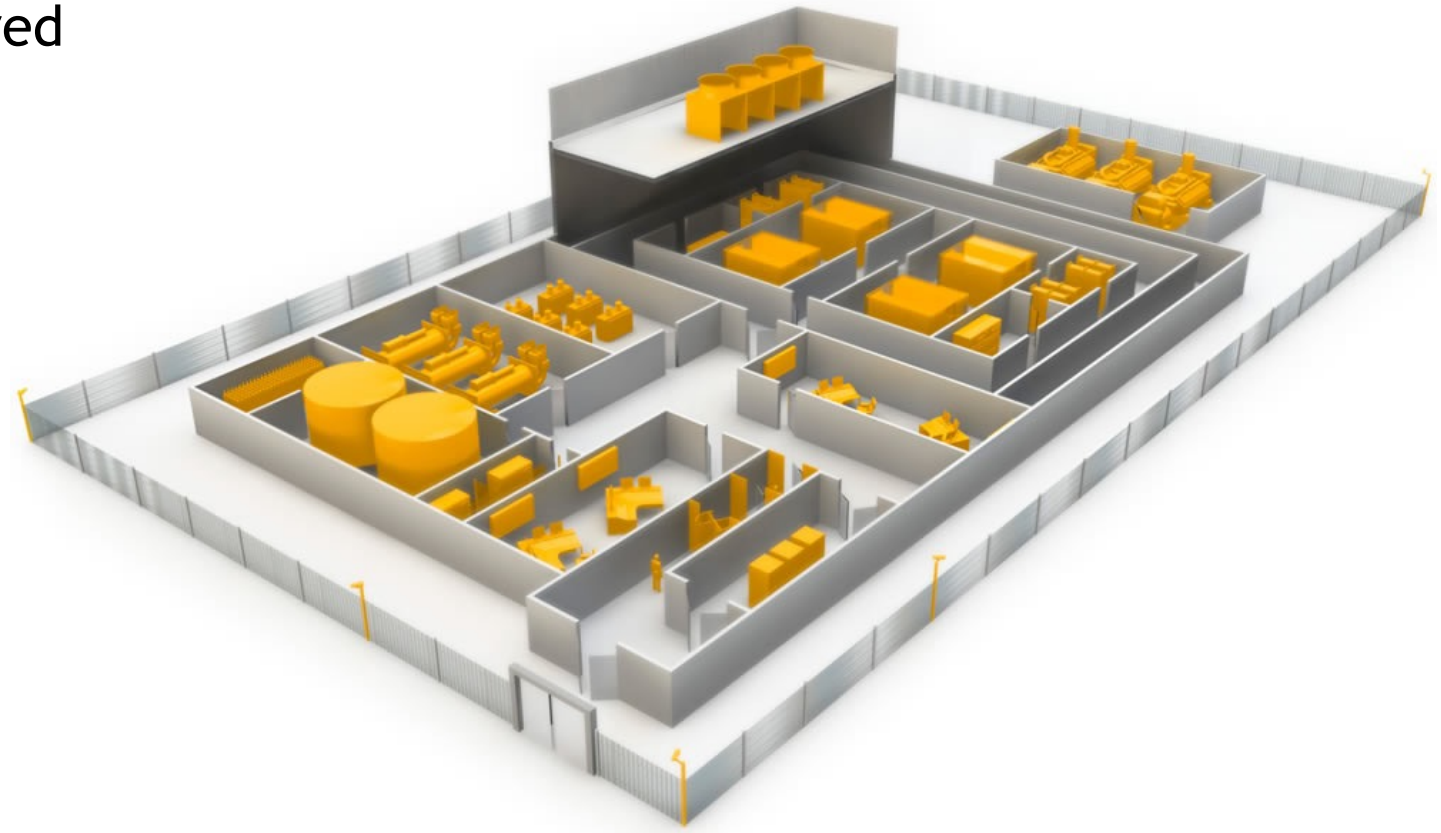

DATA CENTER POWER SYSTEMS

In order to protect against power failure, battery and diesel generators are used to backup the external supply.

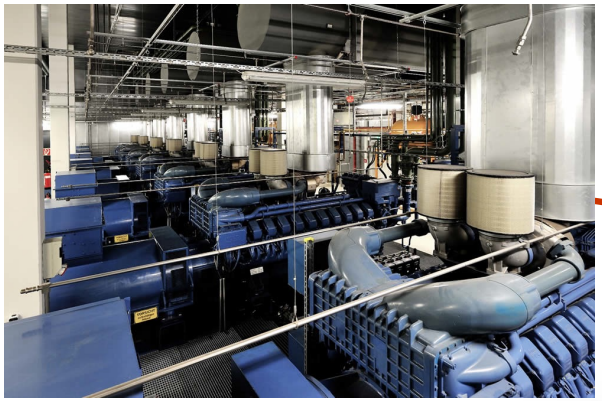

Diesel generators

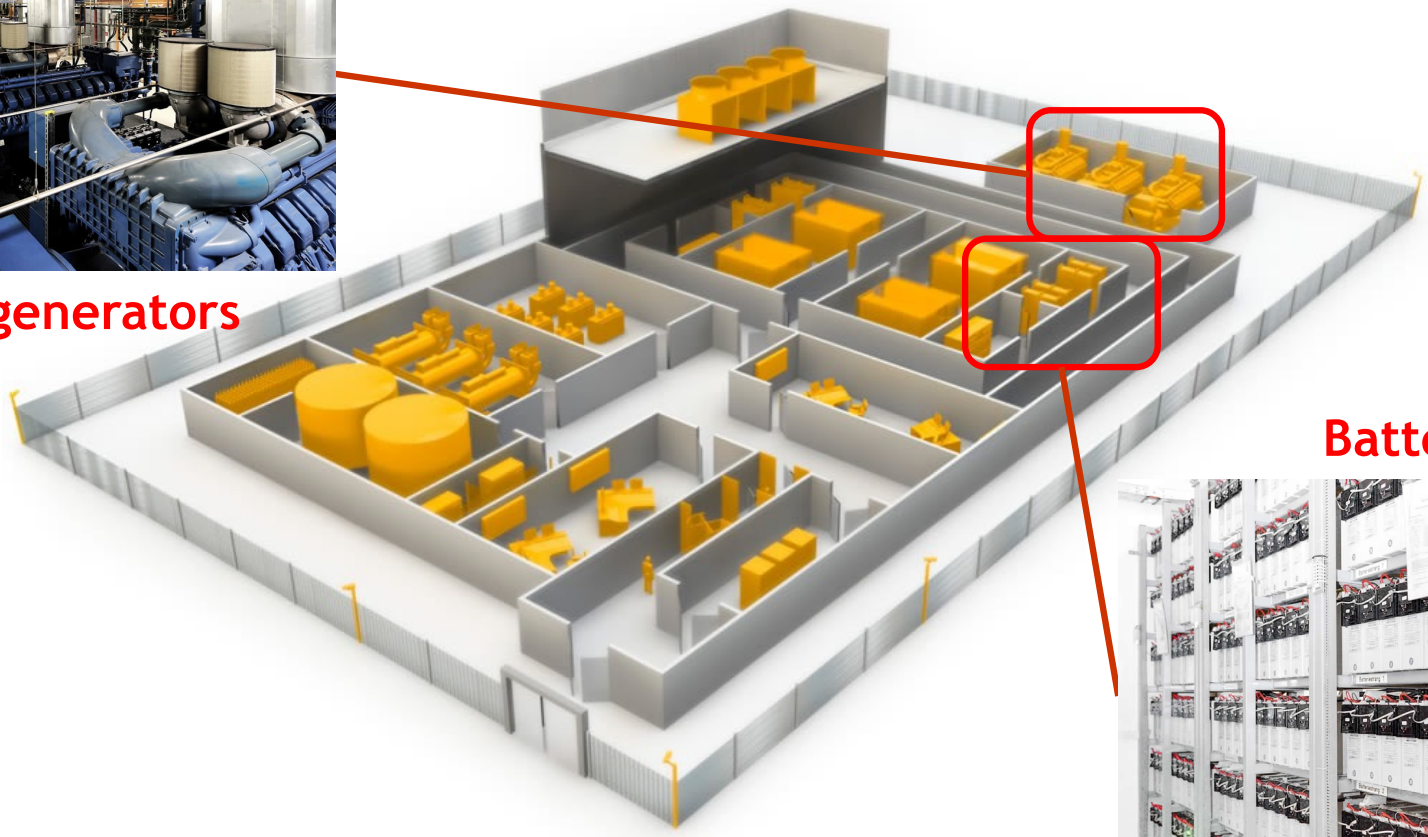

Batteries

DATA CENTER POWER SYSTEMS

In order to protect against power failure, battery and diesel generators are used to backup the external supply.

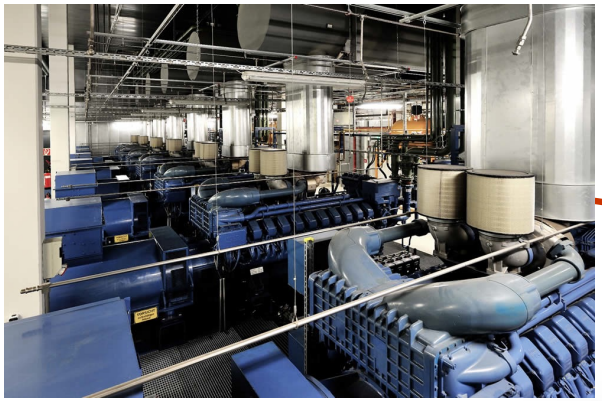

Diesel generators

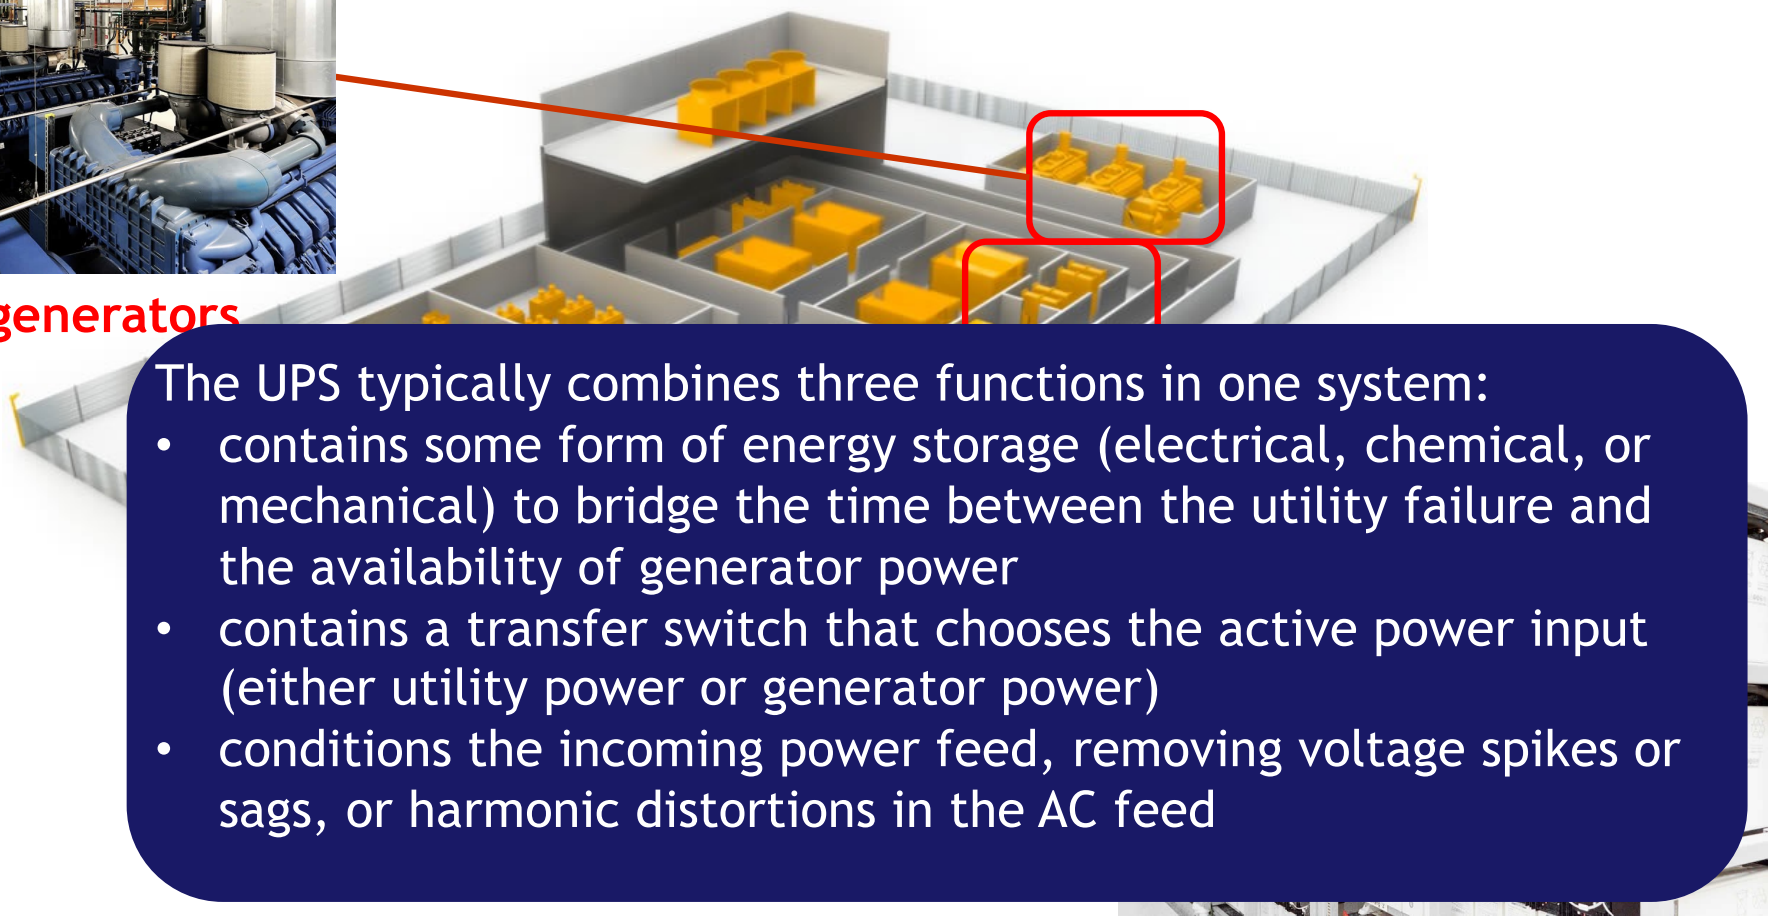

The UPS typically combines three functions in one system:

- contains some form of energy storage (electrical, chemical, or mechanical) to bridge the time between the utility failure and the availability of generator power
- contains a transfer switch that chooses the active power input (either utility power or generator power)
- conditions the incoming power feed, removing voltage spikes or sags, or harmonic distortions in the AC feed

DATA CENTER COOLING SYSTEMS

IT equipment generates **a lot of heat**: the **cooling system** is usually a very expensive component of the datacenter, and it is composed by coolers, heat-exchangers and cold water tanks.

Turbo coolers

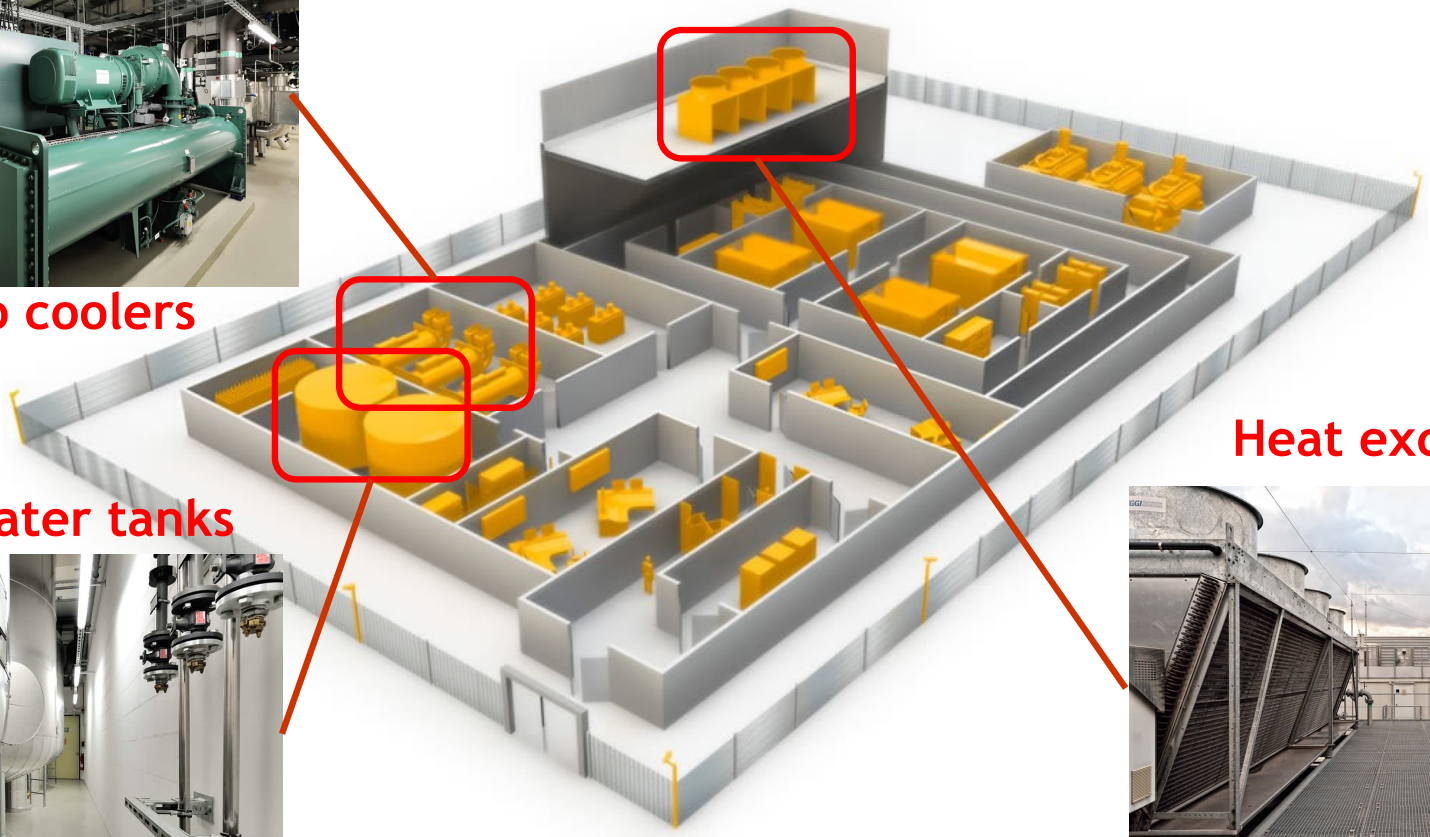

Heat exchangers

Cold water tanks

Data-center corridors

- Server Racks are **NEVER BACK-to-BACK**
- Corridors where servers are located are split into *cold aisle*, where the front panels of the equipment is reachable, and *warm aisle*, where the back connections are located
- Cold air flows from the front (cool aisle), cools down the equipment, and leave the room from the back (warm aisle)

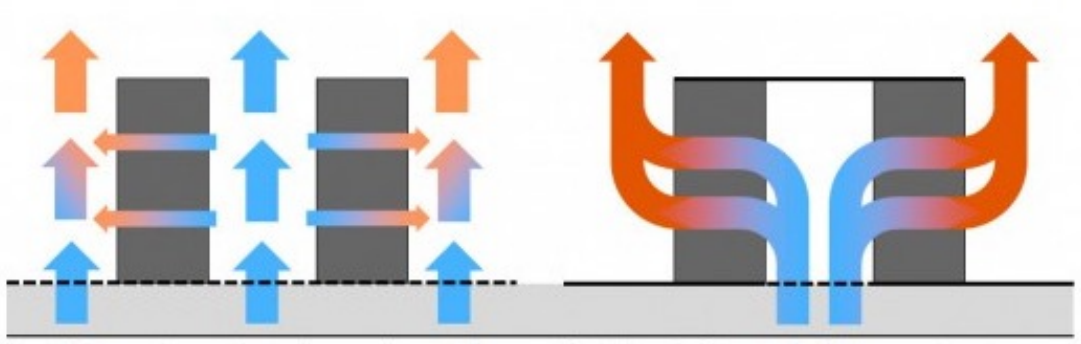

Not Unique Solution

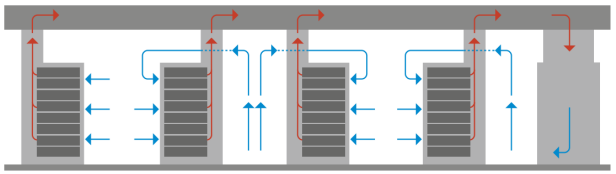

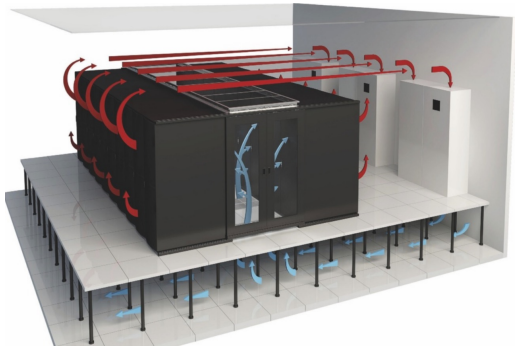

Open-Loop

- The simplest topology is fresh air cooling (or air economization)—essentially, opening the windows
- This is a single «open-loop» system

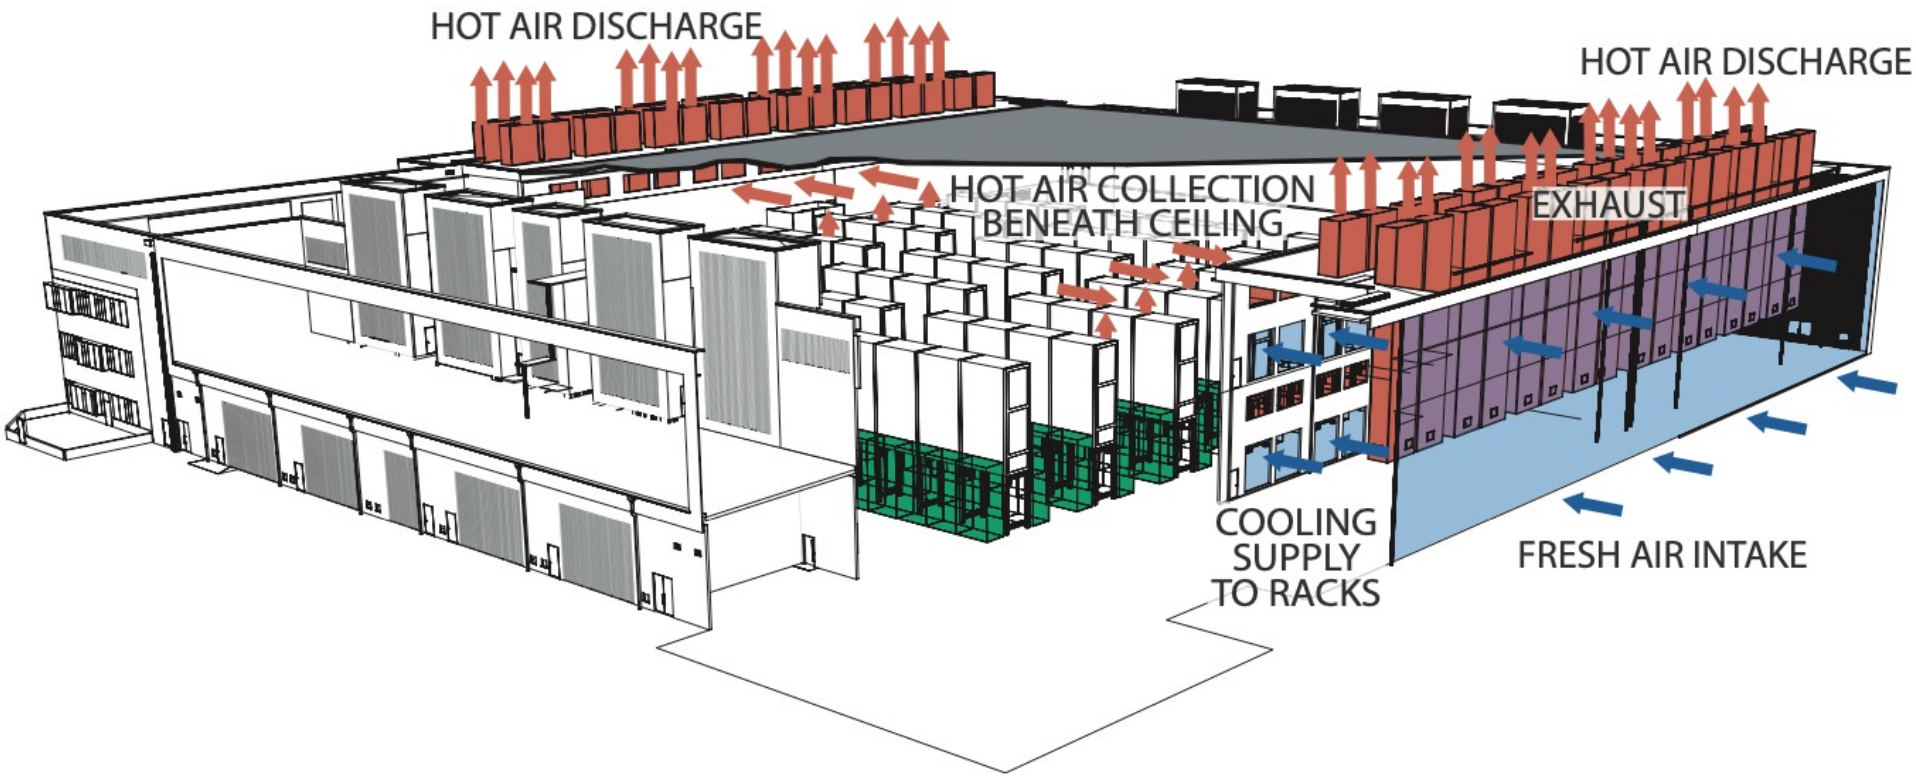

Open vs Closed Loop

- Free cooling, i.e., **open-loop**, refers to the use of cold outside air to either help the production of chilled water or directly cool servers. It is not completely free in the sense of zero cost, but it involves very low-energy costs compared to chillers

- **Closed-loop** systems come in many forms, the most common being the air circuit on the data center floor
  - The goal is to isolate and remove heat from the servers and transport it to a heat exchanger
  - Cold air flows to the servers, heats up, and eventually reaches a heat exchanger to cool it down again for the next cycle through the servers

Closed-loop with two loops

- ✓ The airflow through the underfloor plenum, the racks, and back to the CRAC (a 1960s term for *computer room air conditioning*) defines the primary air circuit, i.e., the first loop
- ✓ The second loop (the liquid supply inside the CRACs units) leads directly from the CRAC to external heat exchangers (typically placed on the building roof) that discharge the heat to the environment

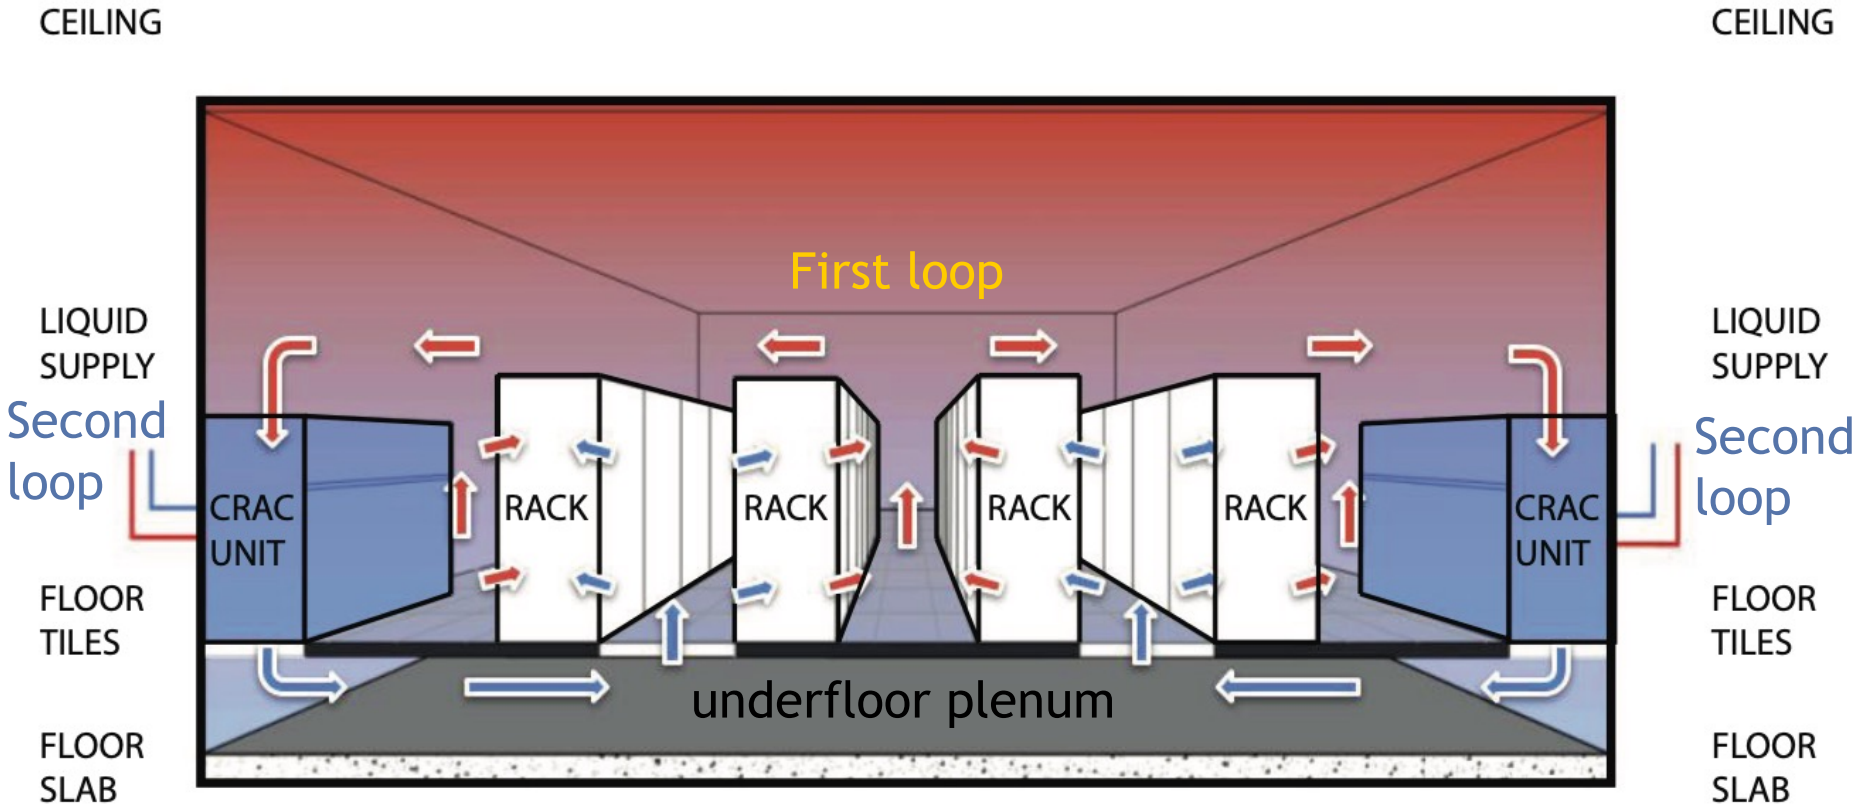

A three-loop system commonly used in large-scale data center

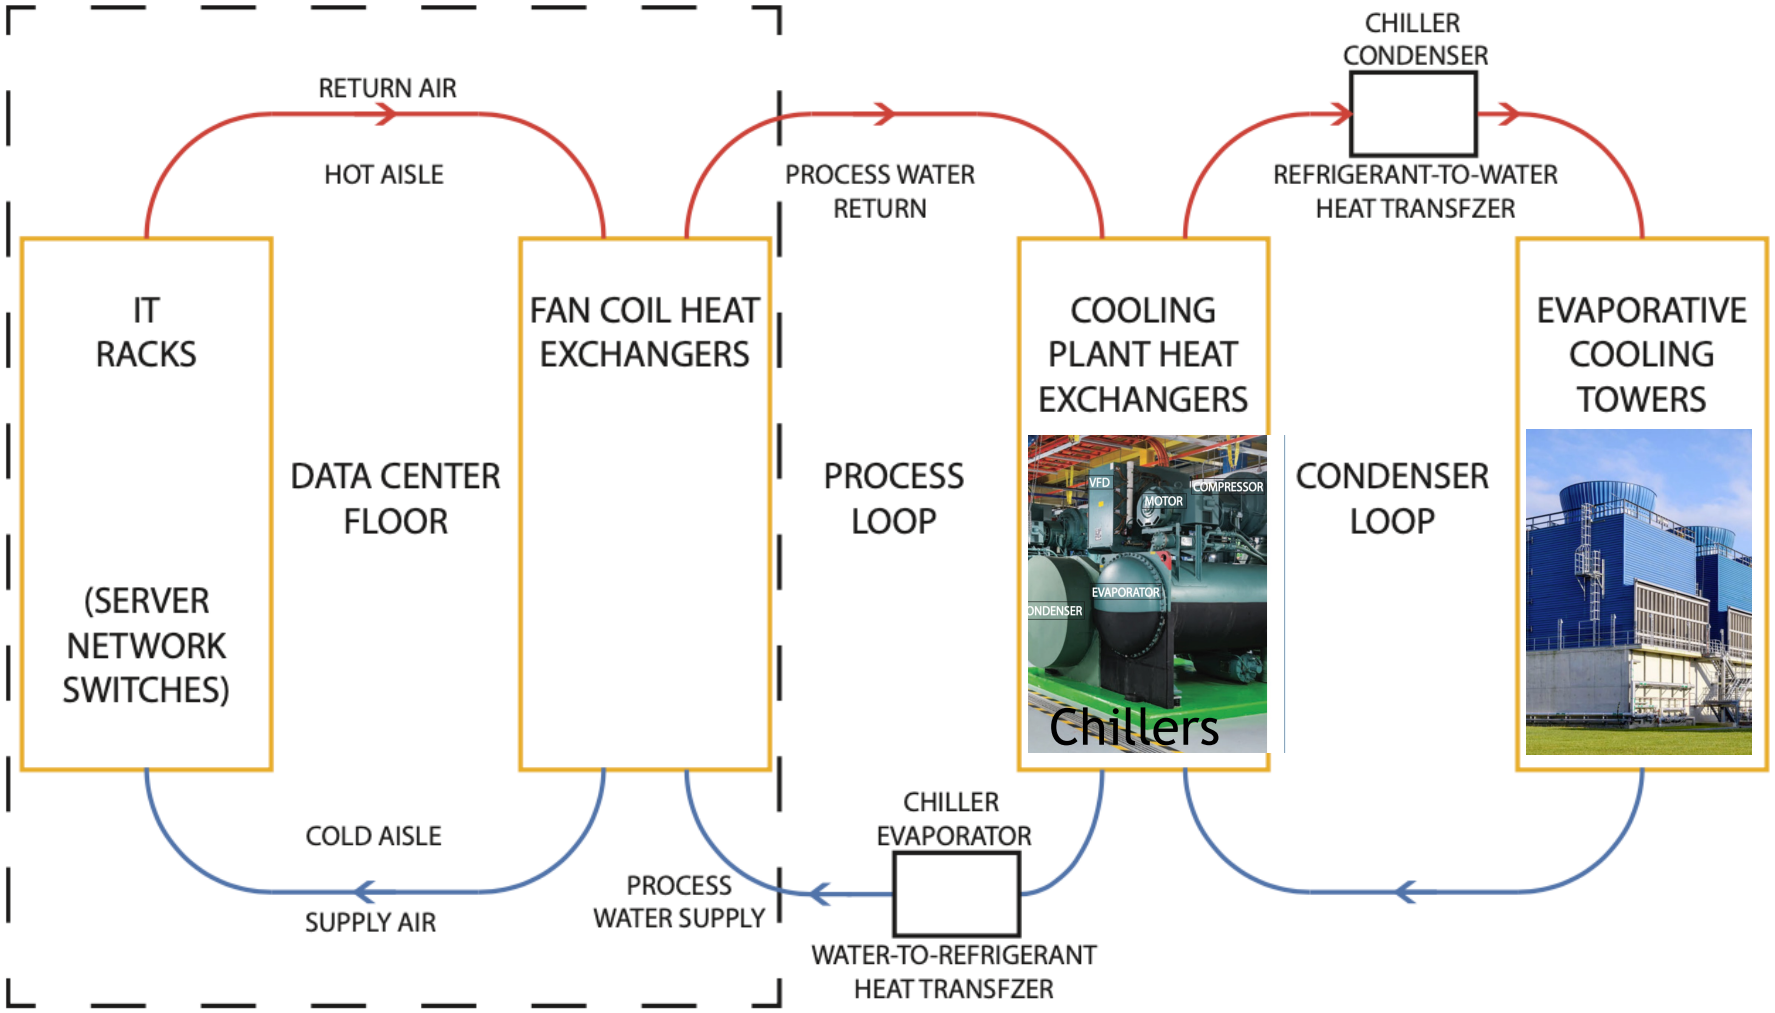

First loop Second loop Third loop

Chillers and Cooling Towers

A water-cooled chiller can be thought of as a water-cooled air conditioner

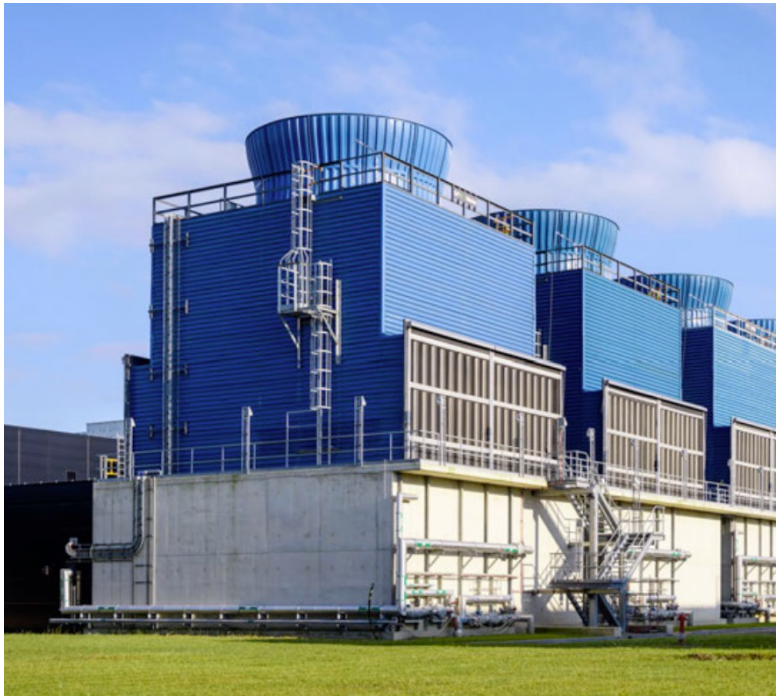

Cooling towers cool a water stream by evaporating a portion of it into the atmosphere. They do not work as well in very cold climates because they need additional mechanisms to prevent ice formation

A critical comparison

Each topology presents tradeoffs in complexity, efficiency, and cost:

- ✓ **Fresh air cooling** can be very efficient but does not work in all climates, requires filtering of airborne particulates, and can introduce complex control problems
- ✓ **Two-loop systems** are easy to implement, relatively inexpensive to construct, and offer isolation from external contamination, but typically have lower operational efficiency
- ✓ **A three-loop system** is the most expensive to construct and has moderately complex controls, but offers contaminant protection and good efficiency

What's next?
In-rack, In-row, and Liquid Cooling

- **In-rack cooler** adds an air-to-water heat exchanger at the back of a rack so the hot air exiting the servers immediately flows over coils cooled by water, essentially reducing the path between server exhaust and CRAC input
- **In-row cooling** works like in-rack cooling except the cooling coils are not in the rack, but adjacent to the rack.

In-rack cooling

What's next?
In-rack, In-row, and Liquid Cooling

- **In-rack cooler** adds an air-to-water heat exchanger at the back of a rack so the hot air exiting the servers immediately flows over coils cooled by water, essentially reducing the path between server exhaust and CRAC input
- **In-row cooling** works like in-rack cooling except the cooling coils are not in the rack, but adjacent to the rack.

In-row cooling

Liquid cooling

- We can directly cool server components using cold plates, i.e., local liquid-cooled heat sinks:
  - Impractical to cool all compute components with cold plates
  - Components with the highest power dissipation are targeted for liquid cooling while other components are air-cooled
- The liquid circulating through the heat sinks transports the heat to a liquid-to-air or liquid-to-liquid heat exchanger that can be placed close to the tray or rack, or be part of the data center building (such as a cooling tower)

Google TPUs

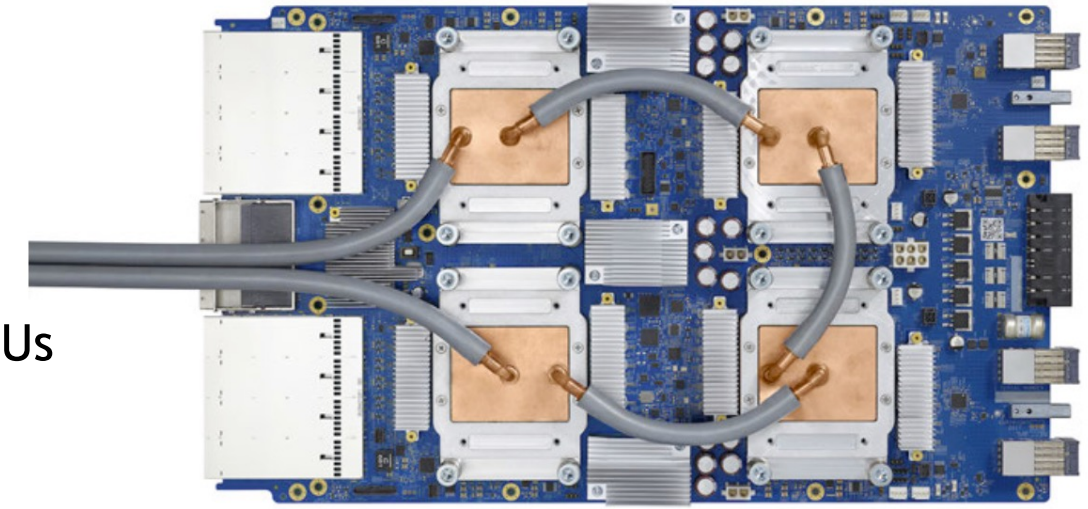

Liquid cooling

- We can directly cool server components using cold plates, i.e., local liquid-cooled heat sinks:
  - Impractical to cool all compute components with cold plates
  - Components with the highe
    for liquid cooling while oth
- The liquid circulating through t
  to a liquid-to-air or liquid-to-li
  placed close to the tray or racl
  building (such as a cooling tow

*Nvidia Hopper H100 – TDP: 700W – T_{core} max: 84 °C

Google TPUs

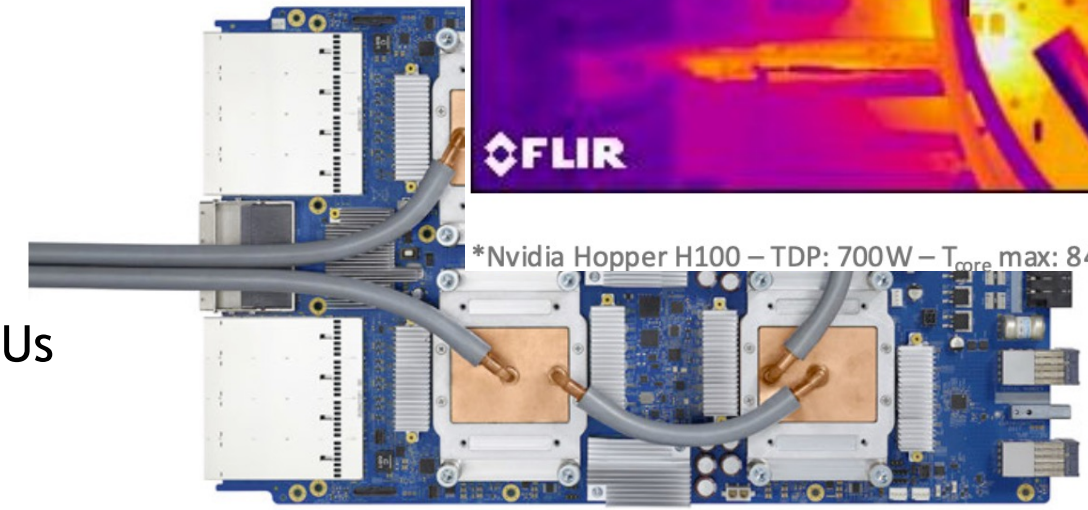

Liquid cooling - two phase cooling

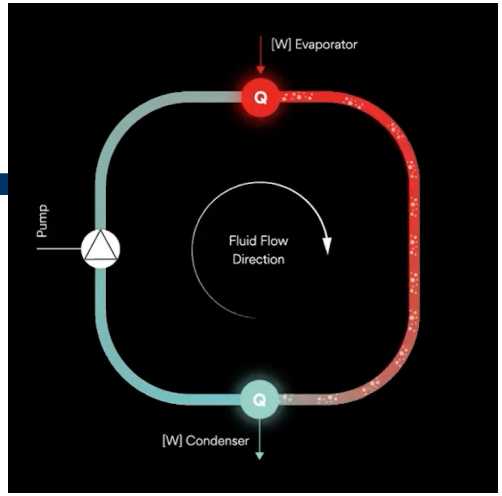

- Advanced thermal management system that leverages the principle of phase change to efficiently dissipate heat
- Uses a working fluid (dielectric liquid) that absorbs heat from components, causing it to change phase from liquid to gas
- The phase change enables the absorption of large amounts of heat due to the latent heat of vaporization
- The vaporized coolant carries the absorbed heat away from the source and moves
- The vapor releases its heat and condenses back into liquid form
- The condensed liquid is returned to a reservoir or accumulator and pumped back to repeat the cycle

Liquid cooling - immersive cooling

- Blades are fully immerged in a non conductive liquid
- Single or two phase liquid
- Noise Reduction: Fans are no longer required, leading to quieter operations
- About 50% energy consumption saving at DC level
- Less space required for the cooling system in the DC
- 60% cheaper data center build
- Faster installation

https://youtu.be/fQfolP7NoNc

Container-based Data Centers

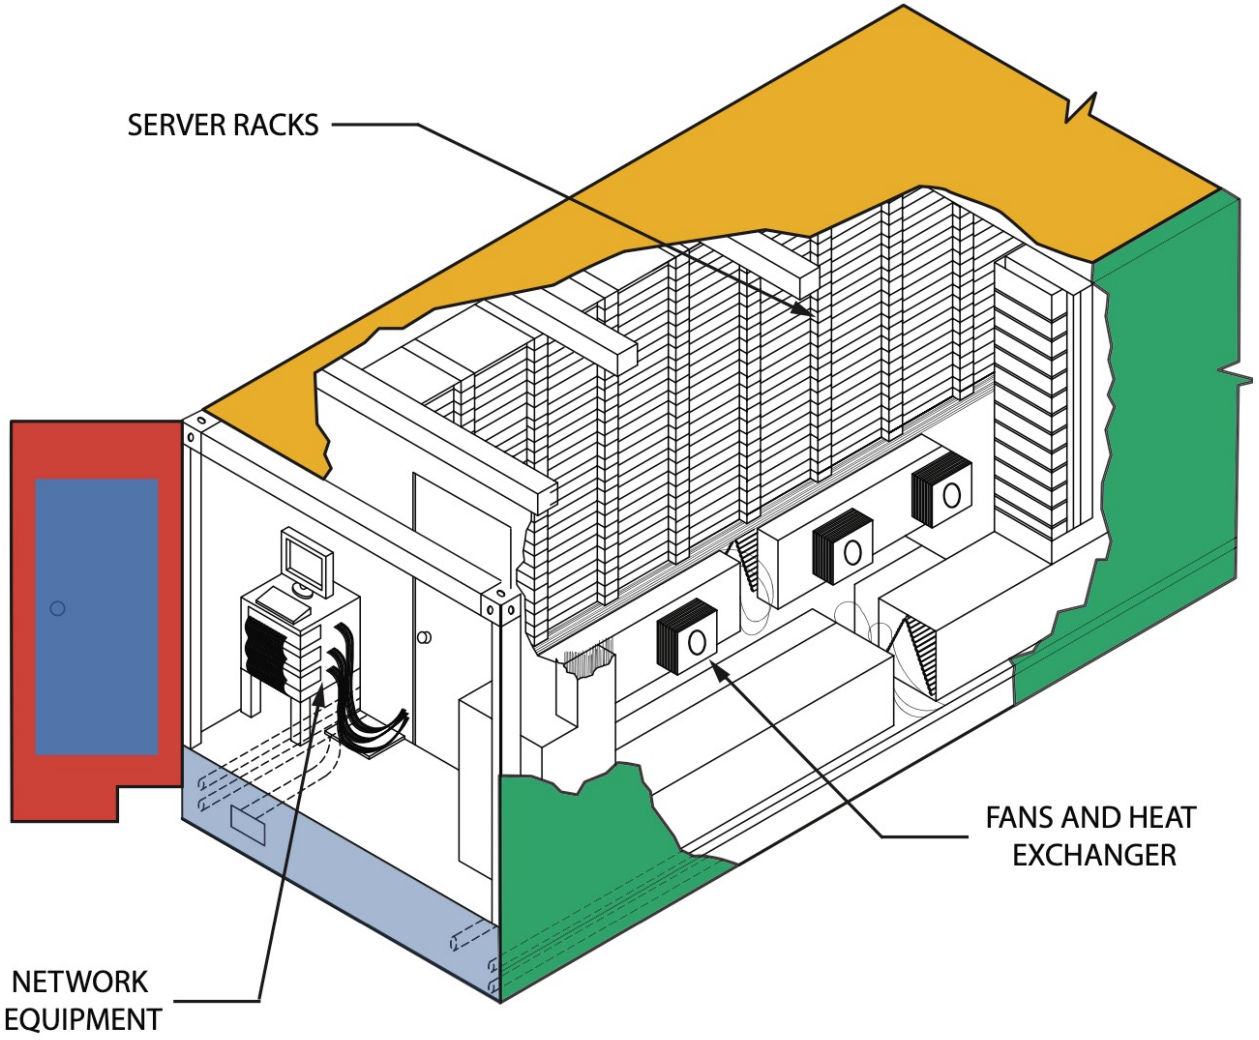

Container-based data centers go one step beyond in-row cooling by placing the server racks inside a container (typically 6 to 12 mt long) and integrating heat exchange and power distribution into the container as well.

Container-based Data Centers

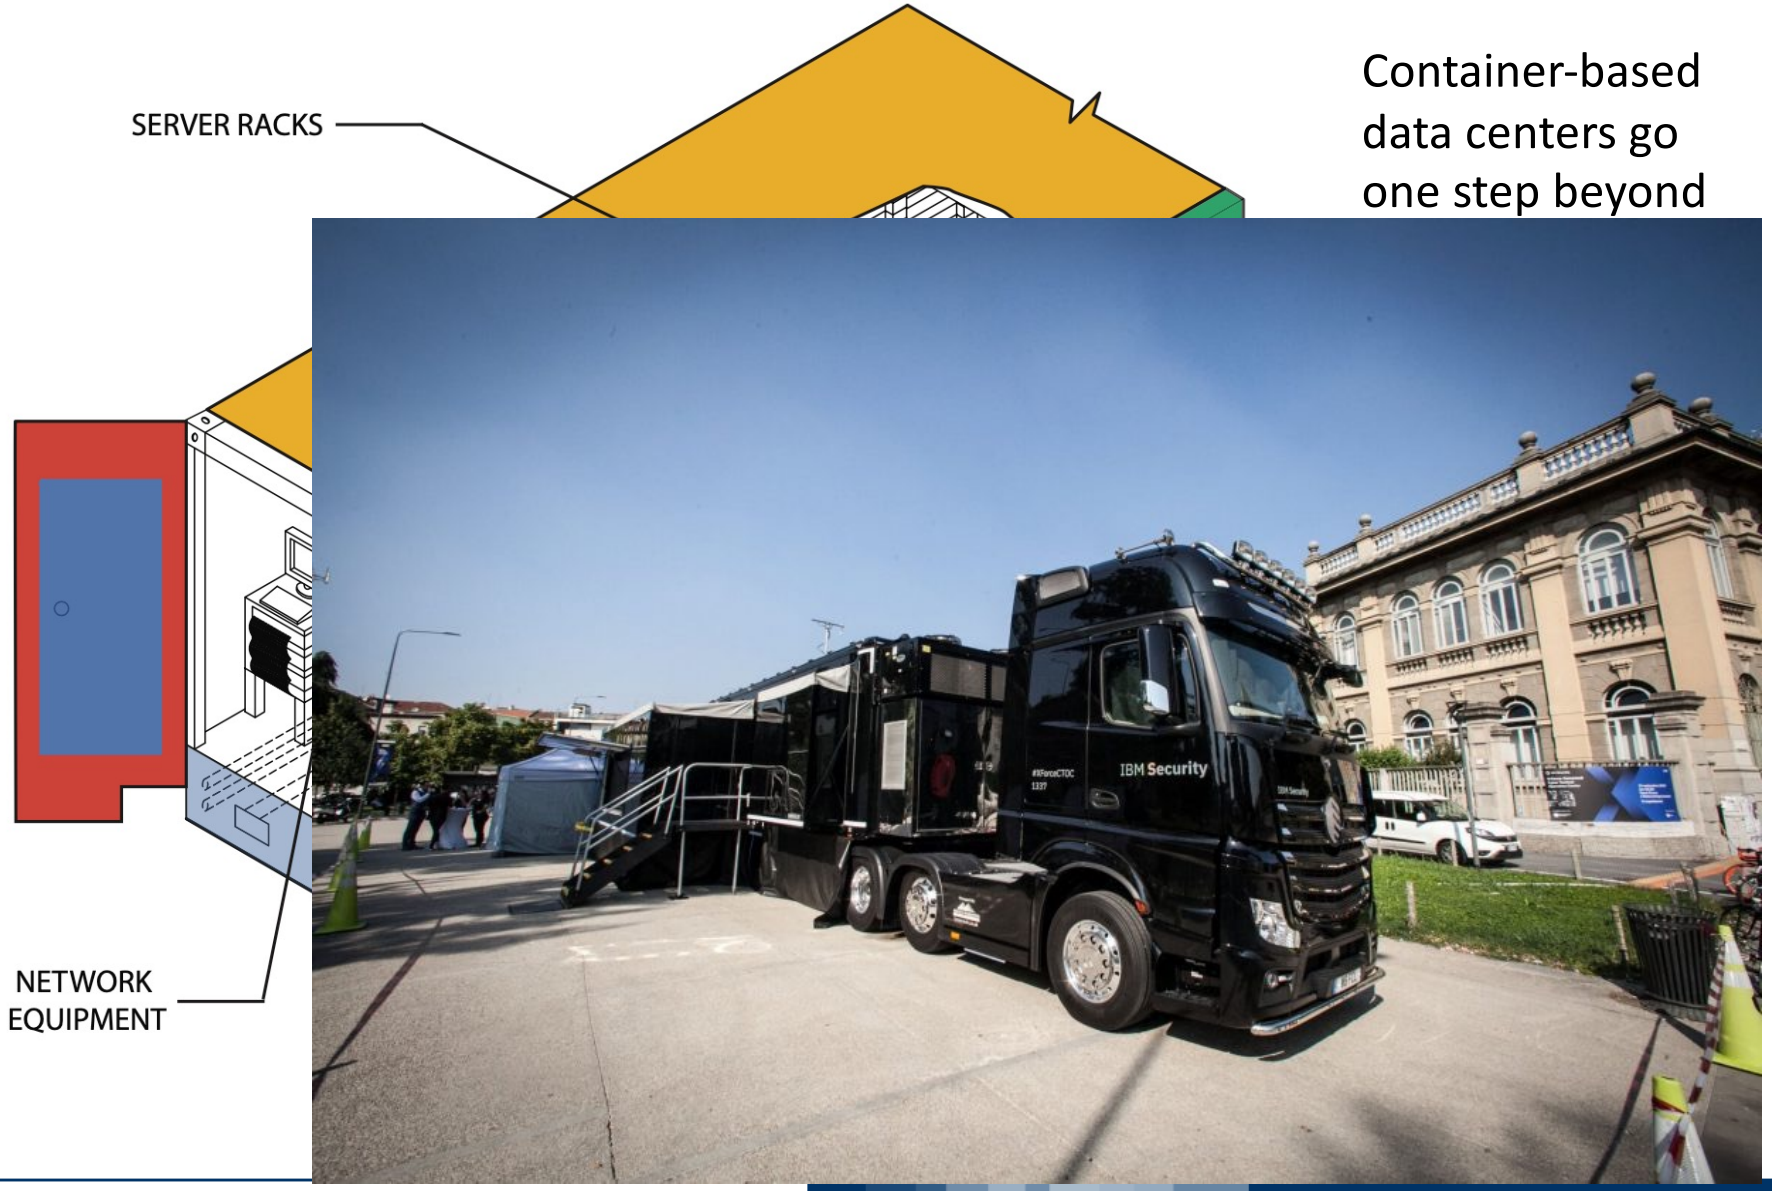

Container-based data centers go one step beyond

Data-center power consumption

- Data-center power consumption is an issue, since it can reach several MWs
- Cooling usually requires about half the energy required by the IT equipment (servers + network + disks)
- Energy transformation creates also a large amount of energy wasted for running a datacenter
- DCs consume 3% of global electricity supply (416.2 TWh > UK's 300 TWh)
- DCs produce 2% of total greenhouse gas emissions (same as worldwide air traffic pre-pandemic)
  - Predictions in the short term (next couple of years): up to 4%
- DCs produce as much CO2 as The Netherlands or Argentina

Data-center power consumption

- Data-center power consumption is an issue, since it can reach several MWs
- Cooling usually requires about half the energy required by the IT equipment (servers + network + disks)
- Energy transformation creates also a large amount of energy wasted for running a datacenter

Amortized Cost	Component	Sub-Components
~45%	Servers	CPU, memory, disk
~25%	Infrastructure	UPS, cooling, power distribution
~15%	Power draw	Electrical utility costs
~15%	Network	Switches, links, transit

PUE Metric

- **Power usage effectiveness** (PUE) is the ratio of the total amount of energy used by a DC facility to the energy delivered to the computing equipment

$$PUE = \frac{Total\ Facility\ Power}{IT\ Equipment\ power}$$

- Total facility power = covers IT systems (servers, network, storage) + other equipment (cooling, UPS, switch gear, generators, lights, fans, etc.)

- Data Center infrastructure Efficiency (DCiE): PUE inverse

PUE Metric

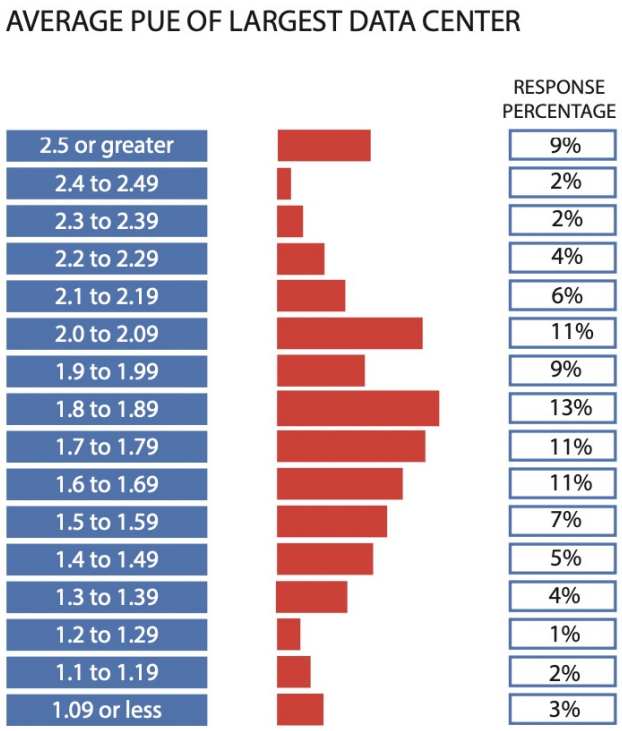

2012 Study

AVERAGE PUE 1.8 - 1.89

PUE	DCiE	Level of Efficieny
3.0	33%	Very Inefficient
2.5	40%	Inefficient
2.0	50%	Average
1.5	67%	Efficient
1.2	83%	Very Efficient

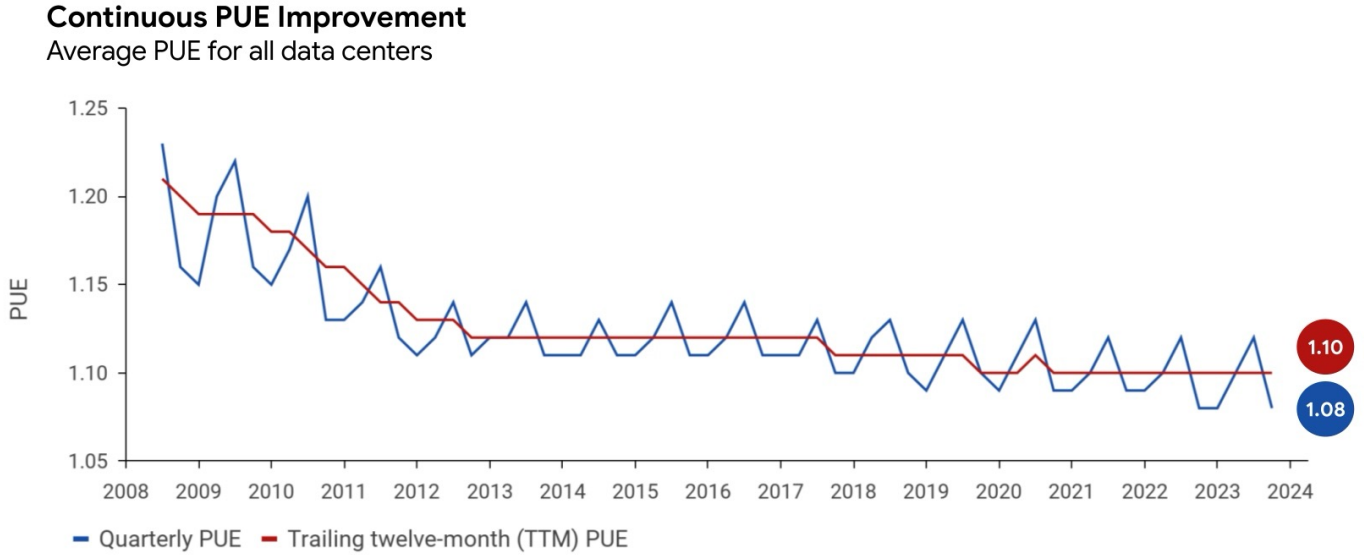

Google DCs

PUE Metric

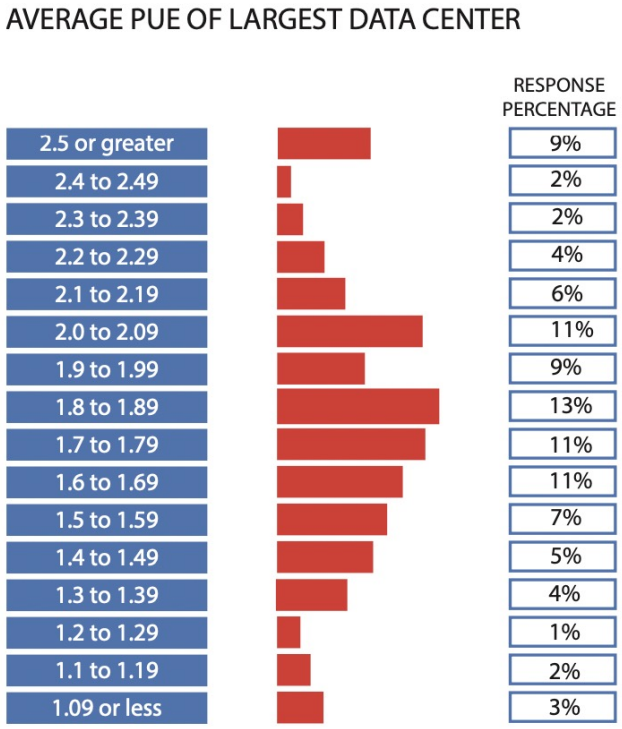

2012 Study

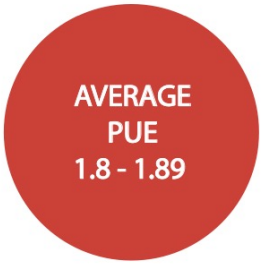

PUE	DCiE	Level of Efficieny
3.0	33%	Very Inefficient
2.5	40%	Inefficient
2.0	50%	Average
1.5	67%	Efficient
1.2	83%	Very Efficient

Additional readings:
https://www.google.com/about/datacenters/efficiency/

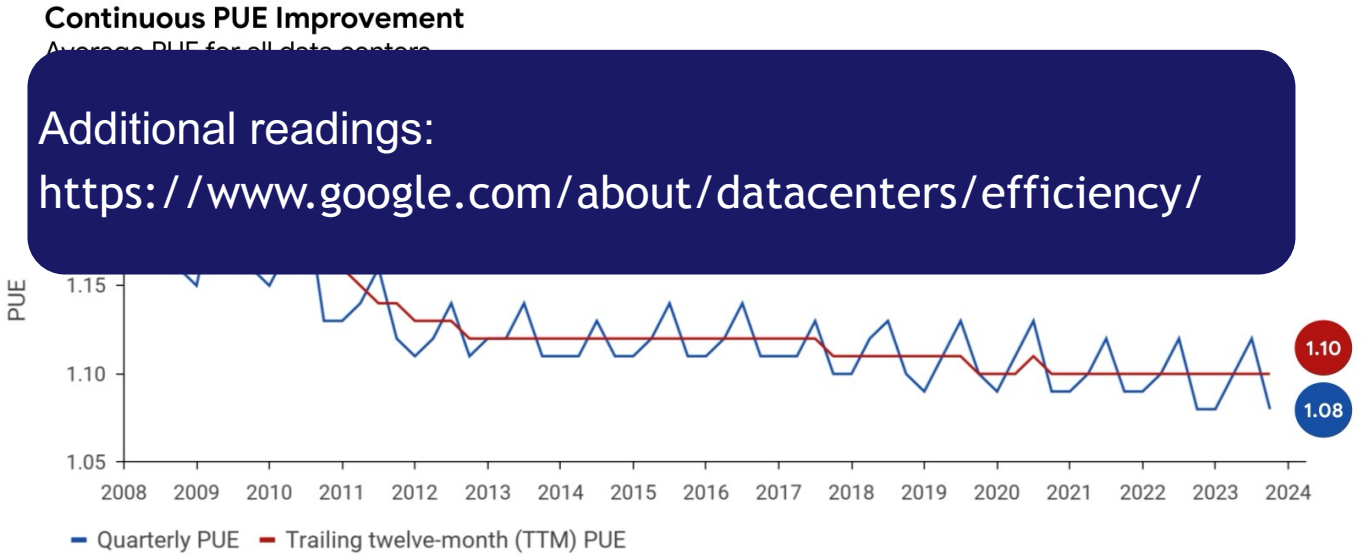

Google DCs

Carbon computing at Google

Keynote 1 – November 9, 5.30pm-6.30pm Milan, 11.30am-12.30am NYC

Carbon Aware Computing at Google

Ana Radovanovic has been a research scientist at Google since early 2008, after she earned her PhD Degree in Electrical Engineering from Columbia University (2005) and worked for 3 years as a Research Staff Member in the Mathematical Sciences Department at IBM TJ Watson Research Center. For the last 8 years, Ana Radovanovic has focused all her research efforts at Google on building innovative technologies and business models with two goals in mind: (i) to deliver more reliable, affordable and clean electricity to everyone in the world, and (ii) to help Google become a thought leader in decarbonizing the electricity grid. Nowadays, Ana is widely recognized as a technical lead and research entrepreneur. She is a Senior Staff Research Scientist, serving as a Technical Lead for Energy Analytics and Carbon Aware Computing at Google.

https://www.performance2021.deib.polimi.it/www.performance2021.deib.polimi.it/keynote-lectures/index.html

https://youtu.be/a0oya25Tir8

Data-center tiers

Data-center availability is defined by in four different tier level. Each one has its own requirements.

Tier Level	Requirements
1	• Single non-redundant distribution path serving the IT equipment • Non-redundant capacity components • Basic site infrastructure with expected availability of 99.671%
2	• Meets or exceeds all Tier 1 requirements • Redundant site infrastructure capacity components with expected availability of 99.741%
3	• Meets or exceeds all Tier 2 requirements • Multiple independent distribution paths serving the IT equipment • All IT equipment must be dual-powered and fully compatible with the topology of a site's architecture • Concurrently maintainable site infrastructure with expected availability of 99.982%
4	• Meets or exceeds all Tier 3 requirements • All cooling equipment is independently dual-powered, including chillers and heating, ventilating and air-conditioning (HVAC) systems • Fault-tolerant site infrastructure with electrical power storage and distribution facilities with expected availability of 99.995%