



POLITECNICO
MILANO 1863

DIPARTIMENTO DI ELETTRONICA
INFORMAZIONE E BIOINGEGNERIA

POLITECNICO MILANO 1863
NECST
laboratory

Exercise Session 2

Pipelining, Static Branch Prediction, Dynamic Branch Prediction

Advanced Computer Architectures

17th March 2025

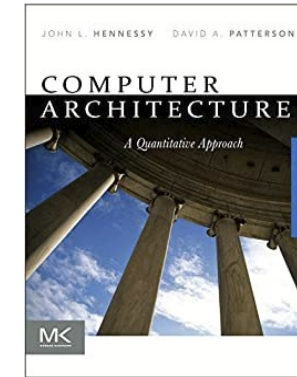
Davide Conficconi <davide.conficconi@polimi.it>

Recall: Material (EVERYTHING OPTIONAL)

<https://webeep.polimi.it/course/view.php?id=14754>



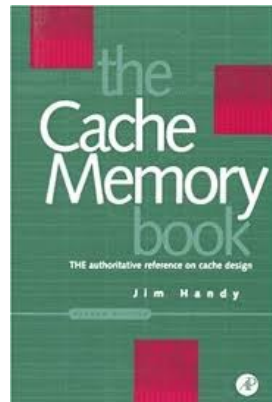

<https://tinyurl.com/aca-grid25>

Textbook: Hennessy and Patterson, Computer Architecture: A Quantitative Approach





Application Software	<code>>"hello world!"</code>
Operating Systems	
Architecture	
Micro-architecture	
Logic	
Digital Circuits	
Analog Circuits	
Devices	
Physics	

Other Interesting Reference



News from the outer world: new interesting books incoming!



News from the outer world: Championship Branch Prediction@ISCA2025

"the goal of this competition is to encourage researchers and practitioners to push the envelope in branch prediction."



<https://www.sigarch.org/call-contributions/championship-branch-prediction-cbp2025-isca-2025/>

A wide banner image showing a cityscape of Tokyo with the Tokyo Tower in the center. Overlaid on the image is the text 'ISCA 2025' in large white letters, a Japanese flag, the dates 'June 21-25, 2025', and 'Tokyo, Japan' in white text.

ISCA 2025



June 21-25, 2025

Tokyo, Japan

News from the outer world: LLMs use prediction (speculation) for decoding


Speculative decoding has proven to be an effective technique for **faster** and **cheaper inference** from **LLMs** without compromising quality.

WITHOUT SPECULATIVE DECODING



My favorite thing about fall is the

WITH SPECULATIVE DECODING




My favorite thing about fall is the change in the leaves. The trees

It has also proven to be an effective paradigm for a range of optimization techniques.

<https://research.google/blog/looking-back-at-speculative-decoding/>

Recall: 5-Stage Pipelining



IF	ID	EX	ME	WB
Instruction Fetch	Instruction Decode	Execution	Memory Access	Write Back

ALU Instructions: `op $x, $y, $z`

Instr. Fetch & PC Increm.	Read of Source Regs. $\$y$ and $\$z$	ALU Op. ($\$y \text{ op } \z)		Write Back Destin. Reg. $\$x$
---------------------------	--------------------------------------	-----------------------------------	--	-------------------------------

Load Instructions: `lw $x, offset($y)`

Instr. Fetch & PC Increm.	Read of Base Reg. $\$y$	ALU Op. ($\$y + \text{offset}$)	Read Mem. $M(\$y + \text{offset})$	Write Back Destin. Reg. $\$x$
---------------------------	-------------------------	-----------------------------------	------------------------------------	-------------------------------

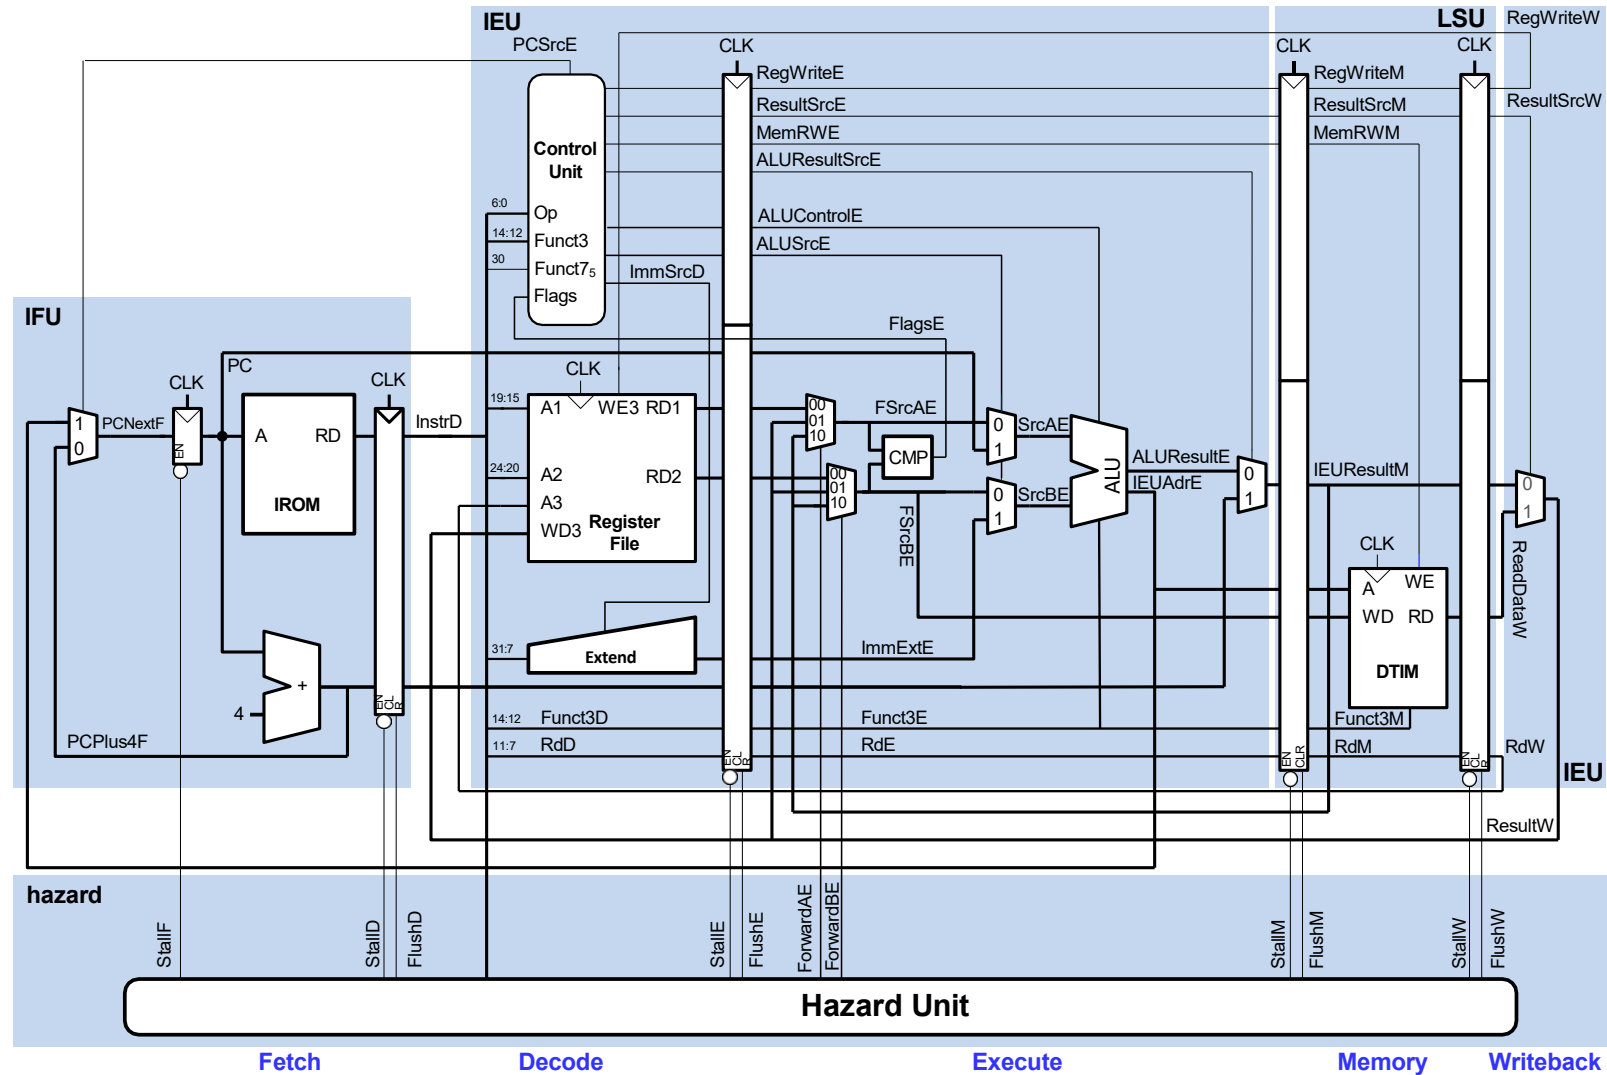
Store Instructions: `sw $x, offset($y)`

Instr. Fetch & PC Increm.	Read of Base Reg. $\$y$ & Source $\$x$	ALU Op. ($\$y + \text{offset}$)	Write Mem. $M(\$y + \text{offset})$	
---------------------------	--	-----------------------------------	-------------------------------------	--


Conditional Branches: `beq $x, $y, offset`

Instr. Fetch & PC Increm.	Read of Source Regs. $\$x$ and $\$y$	ALU Op. ($\$x - \y) & ($\text{PC} + 4 + \text{offset}$)	Write PC	
---------------------------	--------------------------------------	---	----------	--

5-Stage Pipelining RISC-V Wally Educational core



7



Recall: Pipeline performance

Pipeline CPI = Ideal pipeline CPI + Structural Stalls + Data Hazard Stalls + Control Stalls

Ideal pipeline CPI: measure of the maximum performance attainable by the implementation

Structural hazards: HW cannot support this combination of instructions

Data hazards: Instruction depends on result of prior instruction still in the pipeline

Control hazards: Caused by delay between the fetching of instructions and decisions about changes in control flow (branches, jumps, exceptions)

Recall: Three Classes of Hazards

Structural Hazards: Attempt to use the same resource from different instructions simultaneously

Example: Single memory for instructions and data


Data Hazards: Attempt to use a result before it is ready

Example: Instruction depending on a result of a previous instruction still in the pipeline

Control Hazards: Attempt to make a decision on the next instruction to execute before the condition is evaluated

Example: Conditional branch execution

Recall: Data Hazards possible solutions



sub **\$2**, \$1, \$3

and \$12, **\$2**, \$5

or \$13, \$6, **\$2**

add \$14, **\$2**, \$2

sw \$15, 100(**\$2**)

add \$4, \$10, \$11

and \$7, \$8, \$9

lw \$16, 100(\$18)

lw \$17, 200(\$19)



sub **\$2**, \$1, \$3

add \$4, \$10, \$11

and \$7, \$8, \$9

lw \$16, 100(\$18)

lw \$17, 200(\$19)


and \$12, **\$2**, \$5

or \$13, \$6, **\$2**


add \$14, **\$2**, \$2

sw \$15, 100(**\$2**)


Recall: Data Hazards possible solutions



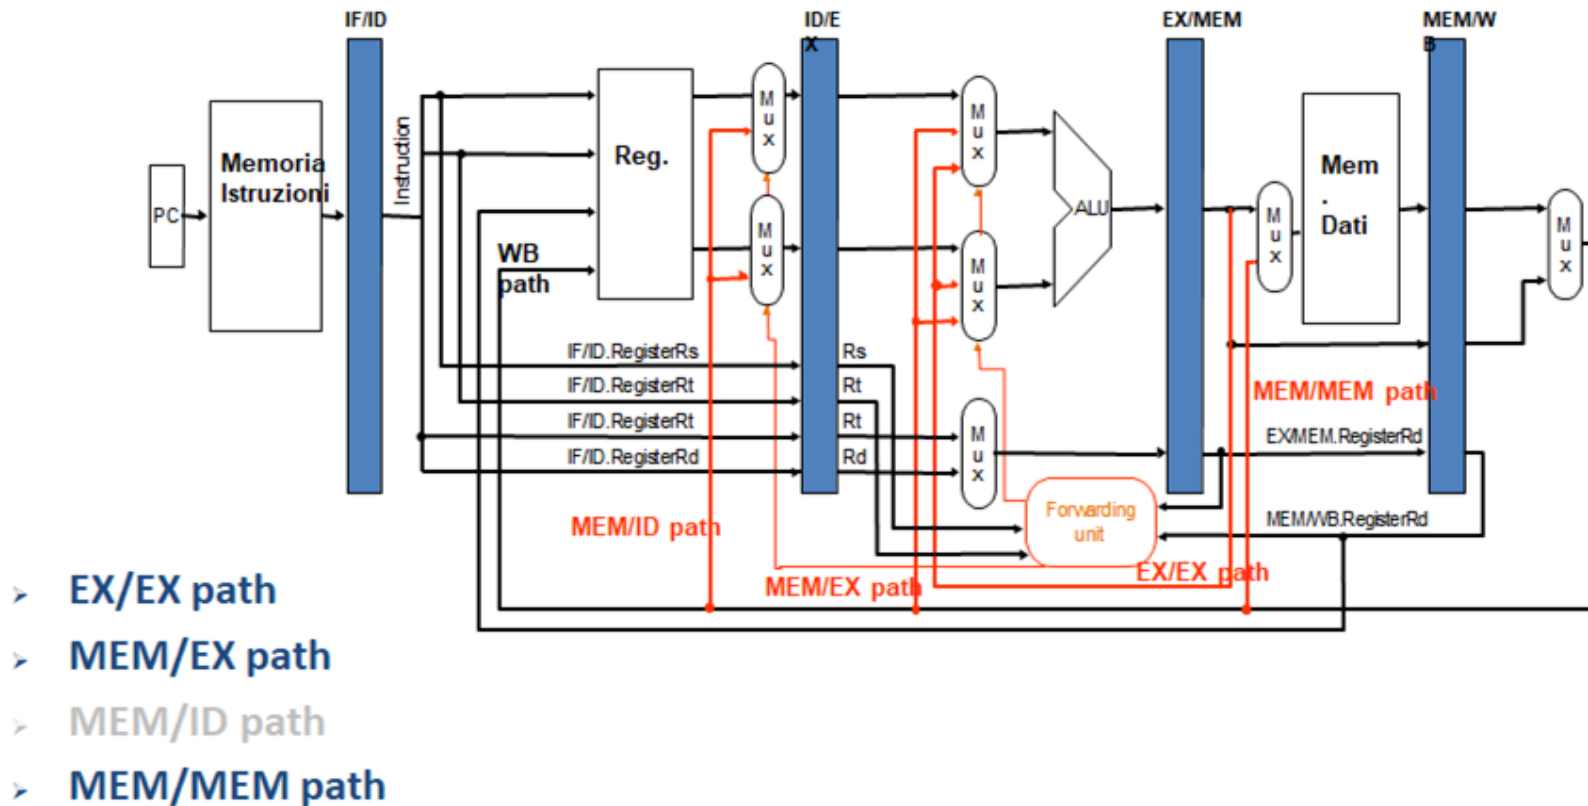

sub **\$2**, \$1, \$3
 and \$12, **\$2**, \$5
 or \$13, \$6, **\$2**
 add \$14, **\$2**, \$2
 sw \$15, 100(**\$2**)
 add \$4, \$10, \$11
 and \$7, \$8, \$9
 lw \$16, 100(\$18)
 lw \$17, 200(\$19)




sub **\$2**, \$1, \$3
 add \$4, \$10, \$11
 and \$7, \$8, \$9
 lw \$16, 100(\$18)
 lw \$17, 200(\$19)
 and \$12, **\$2**, \$5
 or \$13, \$6, **\$2**
 add \$14, **\$2**, \$2
 sw \$15, 100(**\$2**)




Recall: Pipelining and Forwarding



Exe : Pipelining




Exe : Pipelining



i1: add \$t1, \$t0, \$t1
i2: add \$t2, \$t1, \$t2
i3: subi \$t0, \$t2, 1
i4: sw \$t0, 0x00BB(\$t2)
i5: beq \$t0, \$t2, 0x0089

Exe : Pipelining



i1: add \$t1, \$t0, \$t1
i2: add \$t2, \$t1, \$t2
i3: subi \$t0, \$t2, 1
i4: sw \$t0, 0x00BB(\$t2)
i5: beq \$t0, \$t2, 0x0089

IF	ID	EX	ME	WB
Instruction Fetch	Instruction Decode	Execution	Memory Access	Write Back

ALU Instructions: **op \$x, \$y, \$z**

Instr. Fetch & PC Increm.	Read of Source Regs. \$y and \$z	ALU Op. (\$y op \$z)		Write Back Destin. Reg. \$x
---------------------------	----------------------------------	----------------------	--	-----------------------------

Load Instructions: **lw \$x, offset(\$y)**

Instr. Fetch & PC Increm.	Read of Base Reg. \$y	ALU Op. (\$y+offset)	Read Mem. M(\$y+offset)	Write Back Destin. Reg. \$x
---------------------------	-----------------------	----------------------	-------------------------	-----------------------------


Store Instructions: **sw \$x, offset(\$y)**

Instr. Fetch & PC Increm.	Read of Base Reg. \$y & Source \$x	ALU Op. (\$y+offset)	Write Mem. M(\$y+offset)	
---------------------------	------------------------------------	----------------------	--------------------------	--

Conditional Branches: **beq \$x, \$y, offset**

Instr. Fetch & PC Increm.	Read of Source Regs. \$x and \$y	ALU Op. (\$x-\$y) & (PC+4+offset)	Write PC	
---------------------------	----------------------------------	-----------------------------------	----------	--


Exe : Pipelining



i1: add \$t1, \$t0, \$t1
i2: add \$t2, \$t1, \$t2
i3: subi \$t0, \$t2, 1
i4: sw \$t0, 0x00BB(\$t2)
i5: beq \$t0, \$t2, 0x0089

- No forwarding paths
- RF access **R/W optimization**
- **Control Hazard** solved in ID

Exe : Pipelining



i1: add \$t1, \$t0, \$t1

i2: add \$t2, \$t1, \$t2

i3: subi \$t0, \$t2, 1

i4: sw \$t0, 0x00BB(\$t2)

i5: beq \$t0, \$t2, 0x0089

- No forwarding paths
- RF access R/W optimization
- Control Hazard solved in ID


- 1) Define all conflicts/dependencies. For each of them indicate **whether** it causes an **hazard** and the **theoretical** amount of **stalls**
- 2) Draw the effective **pipeline schema**
- 3) Assuming EX/EX, MEM/EX, and MEM/MEM forwarding paths available + 2)
- 4) Assuming EX/ID + 3)

Exe.1 : Dependencies & Hazards

i1: add \$t1, \$t0, \$t1
i2: add \$t2, \$t1, \$t2
i3: subi \$t0, \$t2, 1
i4: sw \$t0, 0x00BB(\$t2)
i5: beq \$t0, \$t2, 0x0089

Instr. #	Dependency on Instr. #	Register involved	Hazard (yes/no)	# of stalls (theoretical)


Exe.2 : Effective Pipeline Schema



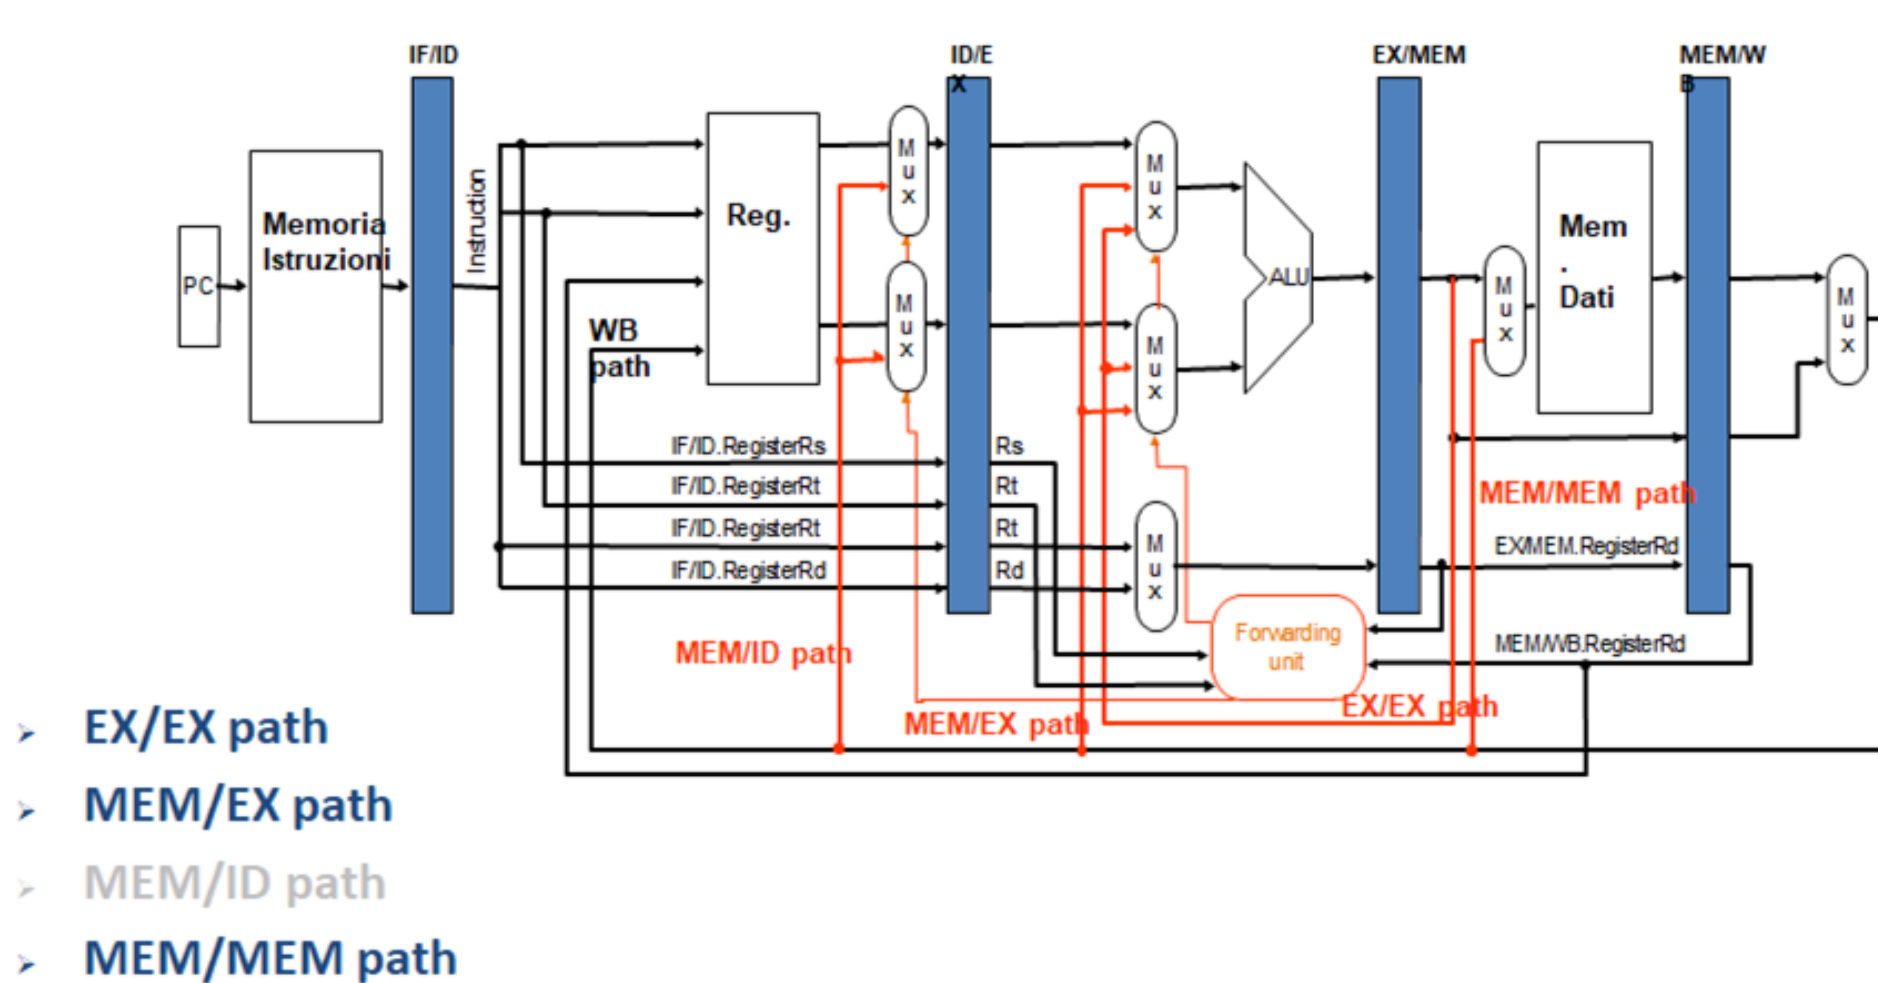
CC 0

[illegible]


Exe.3 : Forwarding Paths

[illegible]

Recall MIPS with Forwarding



Exe.4 : Forwarding Paths + EXE/ID Path

[illegible]

Recall: Three Classes of Hazards

Structural Hazards: Attempt to use the same resource from different instructions simultaneously

Example: Single memory for instructions and data

Data Hazards: Attempt to use a result before it is ready

Example: Instruction depending on a result of a previous instruction still in the pipeline

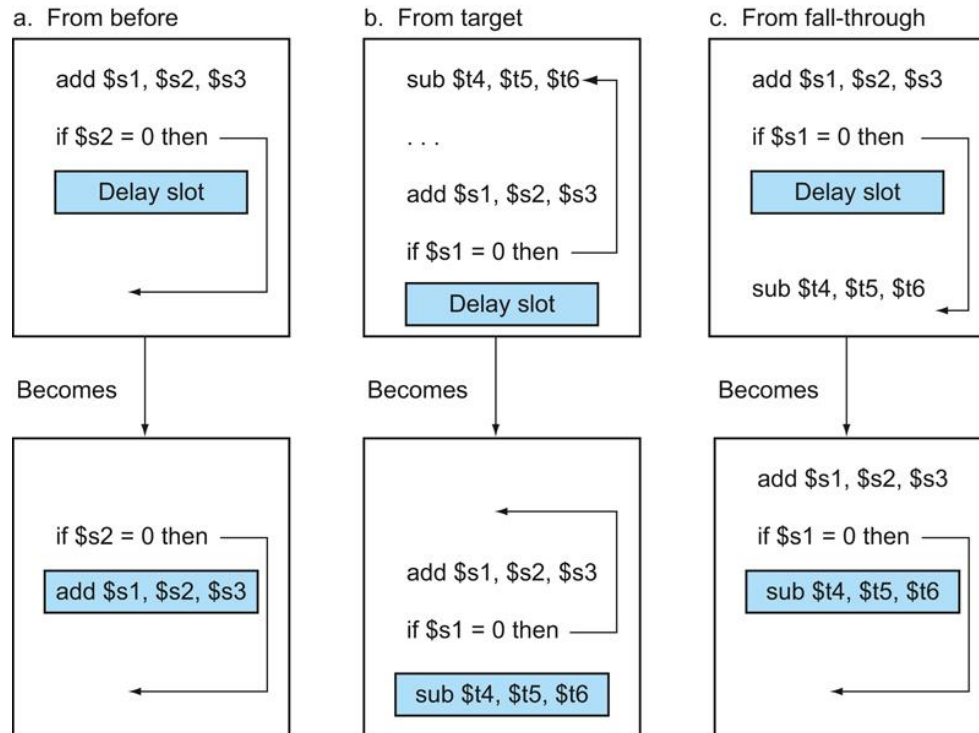
Control Hazards: Attempt to make a decision on the next instruction to execute before the condition is evaluated

Example: Conditional branch execution

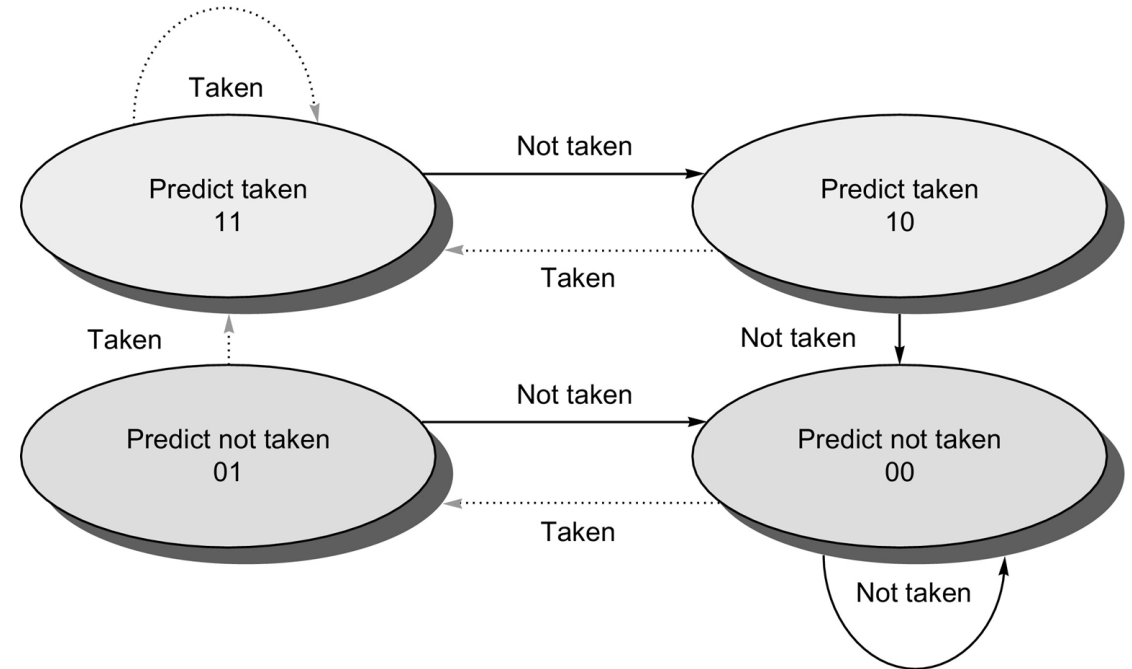
Prediction

Branch vanguard: decomposing branch functionality into prediction and resolution instructions

The IBM z15 High Frequency Mainframe Branch Predictor



Copyright © 2021 Elsevier Inc. All rights reserved.



Recall: Static Branch Prediction Techniques

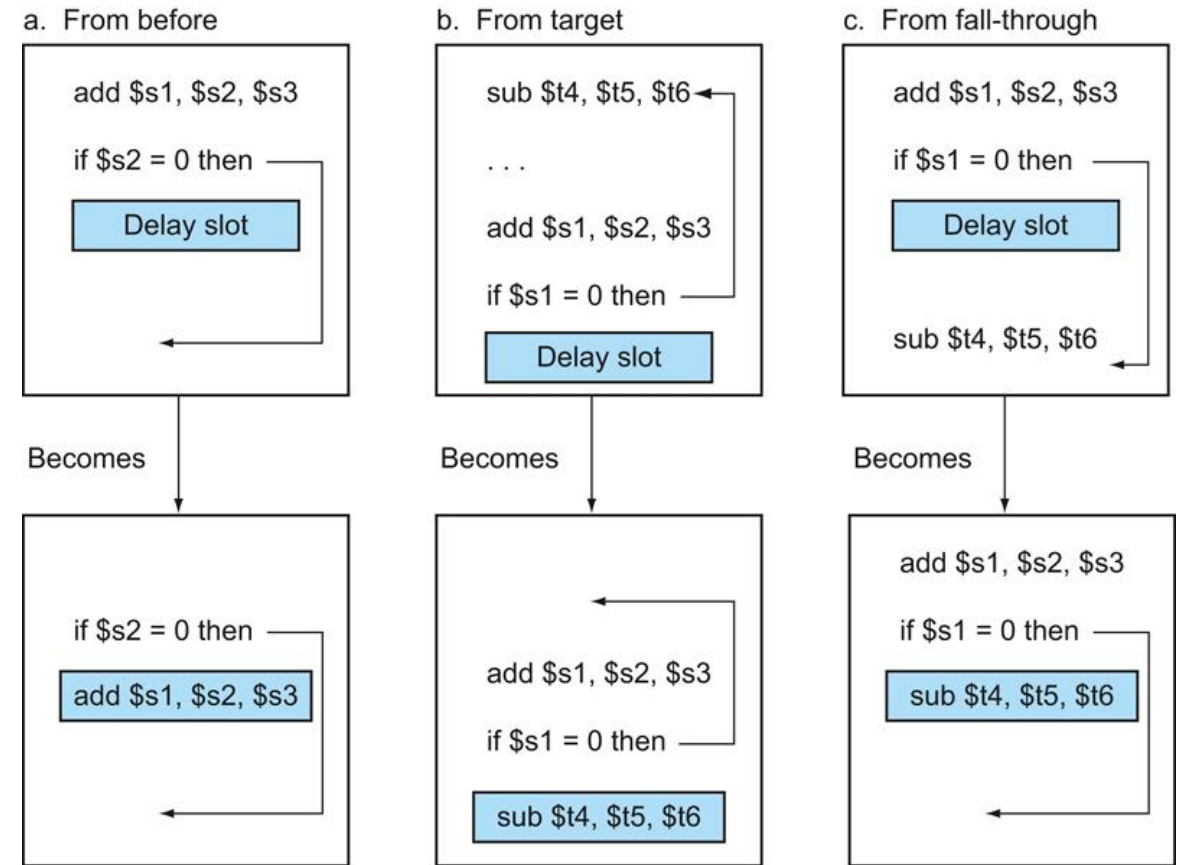
Branch Always Not Taken (Predicted-Not-Taken)

Branch Always Taken (Predicted-Taken)

Backward Taken Forward Not Taken (BTFNT)


Profile-Driven Prediction

Delayed Branch



Copyright © 2021 Elsevier Inc. All rights reserved.

Exe.bonus : Pipeline Schema+Static BP



	Instruction	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C12	C13	C14	C15	C16
1	add \$t1, \$t0, \$t1																
2	add \$t2, \$t1, \$t2																
3	subi \$t0, \$t2, 1																
4	sw \$t0, 0x00BB(\$t2)																
5	beq \$t0, \$t2, 0x0089																
6	NEW INSTRUCTION																



Recall: Three Classes of Hazards

Structural Hazards: Attempt to use the same resource from different instructions simultaneously

Example: Single memory for instructions and data

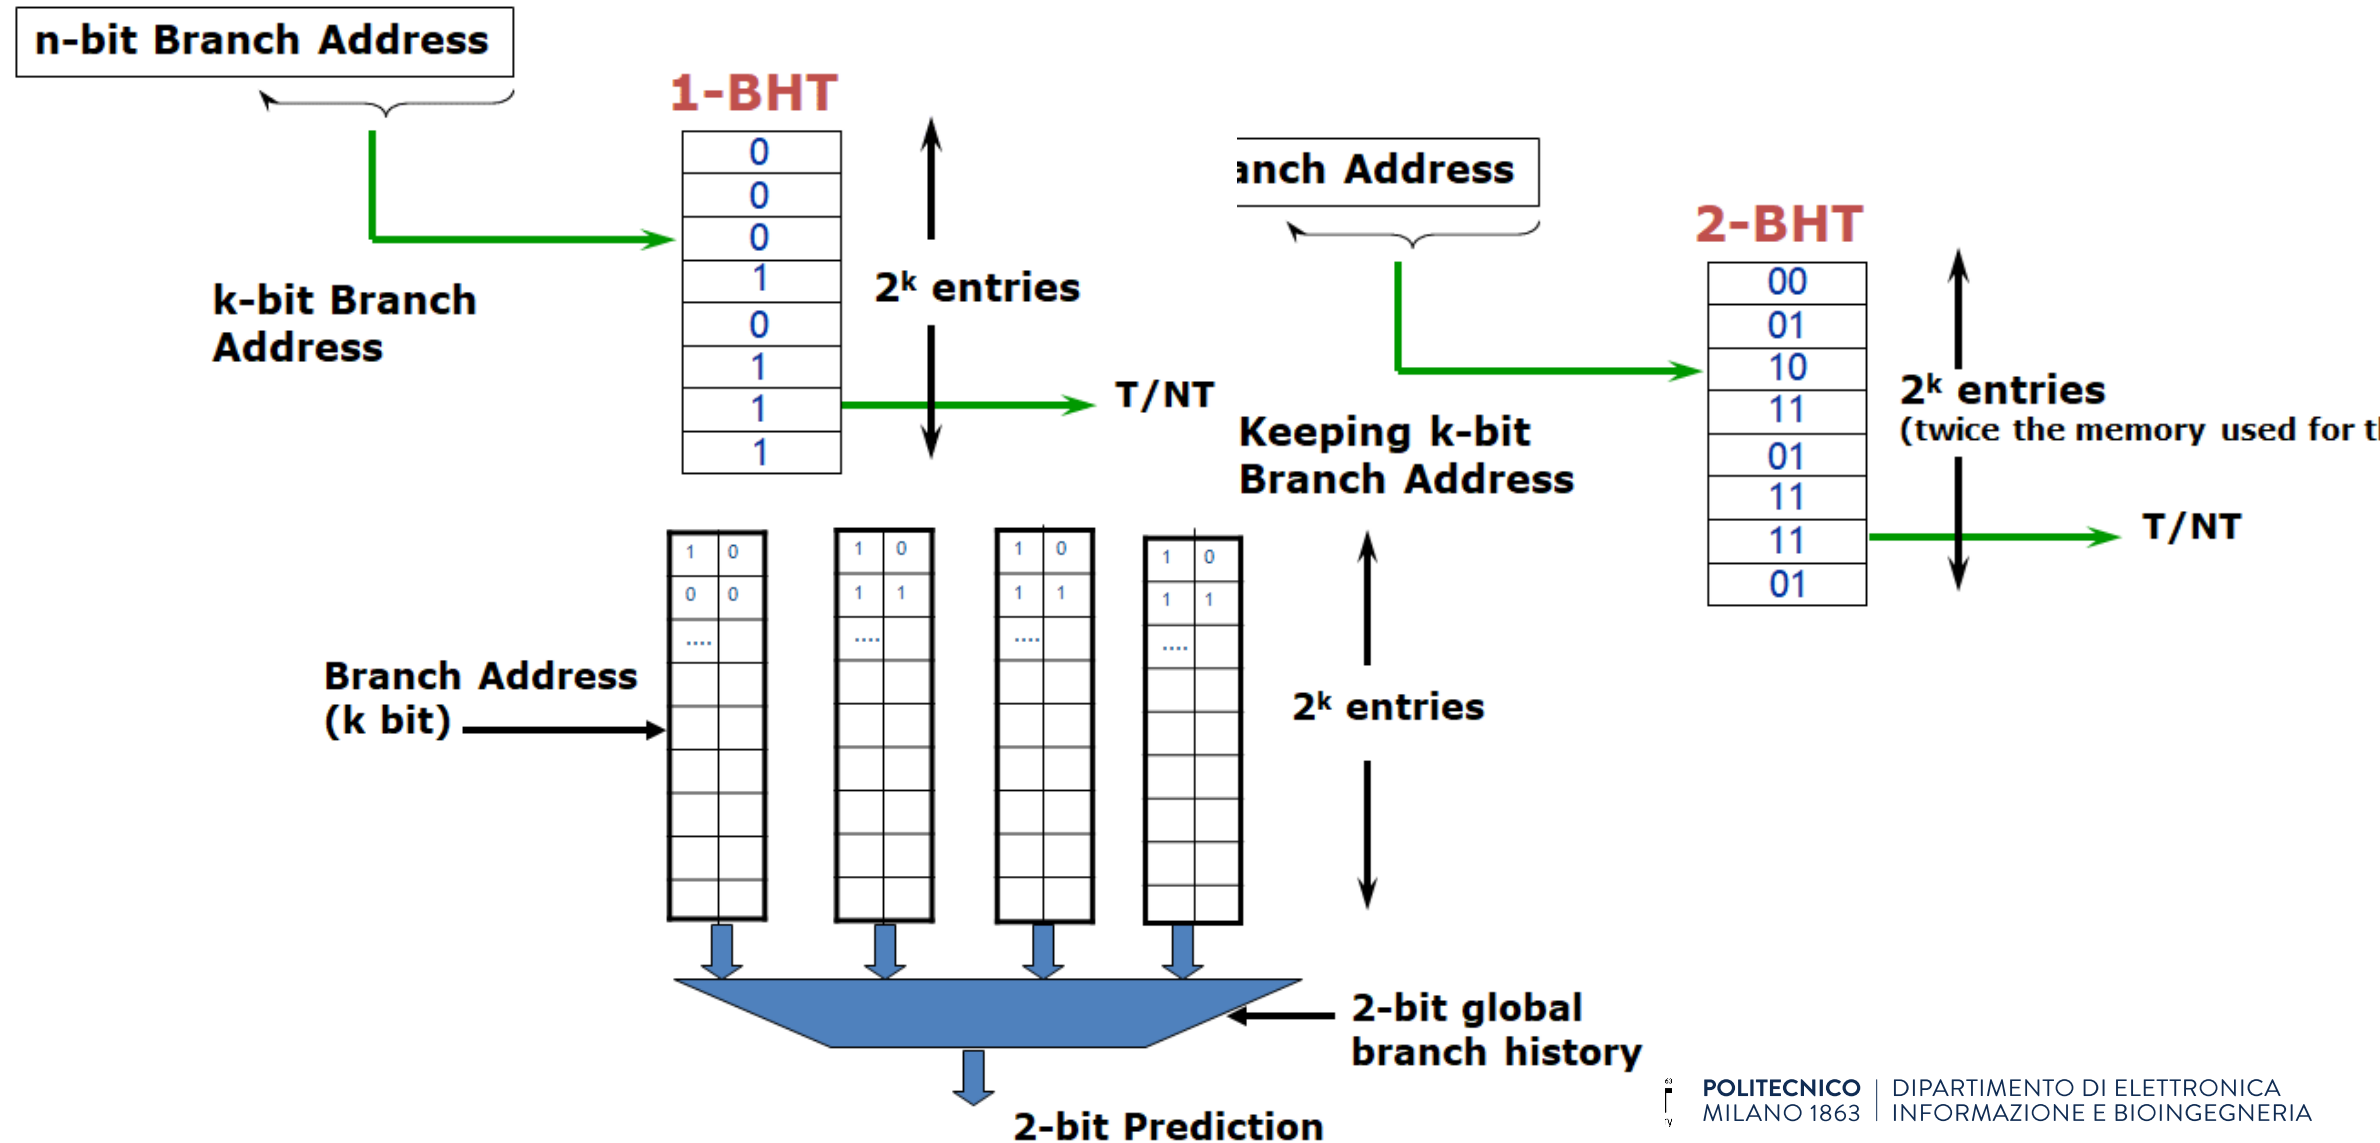
Data Hazards: Attempt to use a result before it is ready

Example: Instruction depending on a result of a previous instruction still in the pipeline

Control Hazards: Attempt to make a decision on the next instruction to execute before the condition is evaluated

Example: Conditional branch execution

Recall: Dynamic Branch Prediction



Dynamic Branch Predictor

- Describe (the answer has to be effectively supported) a 1-BHT and a 2-BHT able to execute the following assembly code (R0 is set to 2000, R1 is set to 0)

```
LOOP:      LD      F1      0      R0
           ADDD    F2      F1      F1
           ADDI    R1      R1      100
LOOP2:     MULTD   F2      F2      F1
           SUBI    R1      R1      1
           BNEZ    R1      LOOP2
           SUBI    R0      R0      2
           BNEZ    R0      LOOP
```

- The obtained result, in terms of mispredictions, is inline with theoretical characteristics of the two predictors? Please effectively support your answer.

A First Consideration

LOOP:	LD	F1	0	R0
	ADDD	F2	F1	F1
	ADDI	R1	R1	100
LOOP2:	MULTD	F2	F2	F1
	SUBI	R1	R1	1
	BNEZ	R1	LOOP2	
	SUBI	R0	R0	2
	BNEZ	R0	LOOP	

How many iterations?


R0 is set to 2000

R1 is set to 0

LOOP:	LD	F1	0	R0
	ADDD	F2	F1	F1
	ADDI	R1	R1	100
LOOP2:	MULTD	F2	F2	F1
	SUBI	R1	R1	1
	BNEZ	R1	LOOP2	
	SUBI	R0	R0	2
	BNEZ	R0	LOOP	

1bit - BHT

LOOP:	LD	F1	0	R0
	ADDD	F2	F1	F1
	ADDI	R1	R1	100
LOOP2:	MULTD	F2	F2	F1
	SUBI	R1	R1	1
	BNEZ	R1	LOOP2	
	SUBI	R0	R0	2
	BNEZ	R0	LOOP	



1bit - BHT


LOOP:	LD	F1	0	R0
	ADDD	F2	F1	F1
	ADDI	R1	R1	100
LOOP2:	MULTD	F2	F2	F1
	SUBI	R1	R1	1
	BNEZ	R1	LOOP2	
	SUBI	R0	R0	2
	BNEZ	R0	LOOP	

n-bit Branch Address

1-BHT



k-bit Branch Address




k-bit Branch Address:
Collide
Not collide

1bit - BHT - Not Collide

```
LOOP:  LD      F1      0      R0
        ADDD    F2      F1     F1
        ADDI    R1      R1     100
LOOP2: MULTD    F2      F2     F1
        SUBI    R1      R1      1
        BNEZ    R1      LOOP2
        SUBI    R0      R0      2
        BNEZ    R0      LOOP
```

R0 is set to 2000

R1 is set to 0



Let us consider that the branch addresses do not collide


	1-BHT	1-BHT	1-BHT	1-BHT
LOOP:				
LOOP2:	T	T	NT	NT
	T	NT	T	NT

1bit - BHT - Not Collide

```

LOOP:  LD      F1      0      R0
        ADDD    F2      F1     F1
        ADDI    R1      R1     100
LOOP2: MULTD    F2      F2     F1
        SUBI    R1      R1      1
        BNEZ    R1      LOOP2
        SUBI    R0      R0      2
        BNEZ    R0      LOOP
    
```

R0 is set to 2000
R1 is set to 0




Let us consider that the branch addresses do not collide

	1-BHT	1-BHT	1-BHT	1-BHT
LOOP:				
LOOP2:	T	T	NT	NT
	T	NT	T	NT

1bit - BHT - Collision

```
LOOP:  LD      F1      0      R0
        ADDD    F2      F1     F1
        ADDI    R1      R1     100
LOOP2: MULTD    F2      F2     F1
        SUBI    R1      R1      1
        BNEZ    R1      LOOP2
        SUBI    R0      R0      2
        BNEZ    R0      LOOP
```

Let us consider that the branch addresses do collide



1-BHT



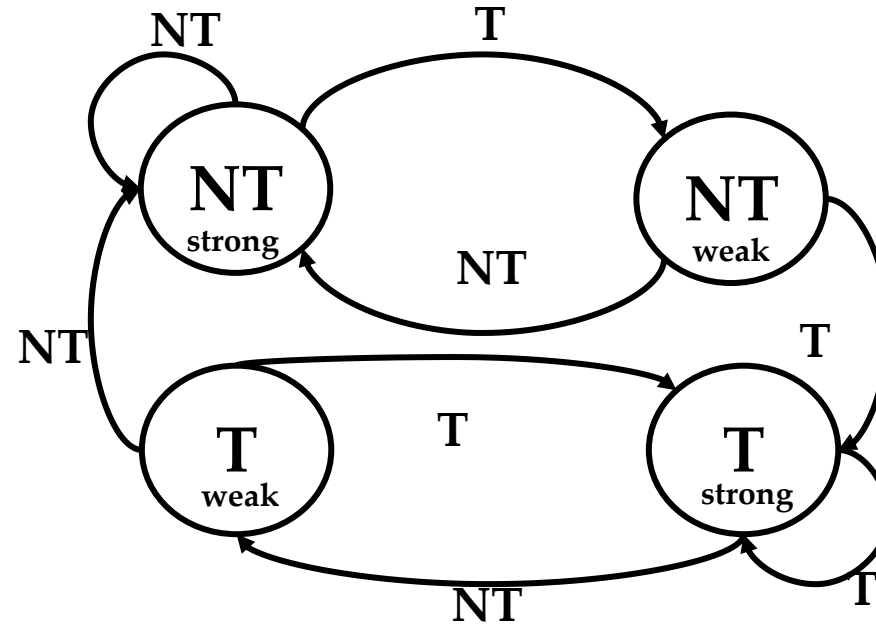
1-BHT



2bit - BHT

LOOP:	LD	F1	0	R0
	ADDD	F2	F1	F1
	ADDI	R1	R1	100
LOOP2:	MULTD	F2	F2	F1
	SUBI	R1	R1	1
	BNEZ	R1	LOOP2	
	SUBI	R0	R0	2
	BNEZ	R0	LOOP	

R0 is set to 2000
R1 is set to 0




2bit - BHT


```

LOOP:  LD    F1    0    R0
       ADDD  F2    F1   F1
       ADDI  R1    R1   100
LOOP2: MULTD F2    F2   F1
       SUBI  R1    R1   1
       BNEZ  R1    LOOP2
       SUBI  R0    R0   2
       BNEZ  R0    LOOP
    
```


Let us consider that the branch addresses do collide




2-BHT




2-BHT



2-BHT





2-BHT





SUMMARY

Assumption: NO collision

WORST CASES



BEST CASES







Thanks for your attention

Davide Conficconi <davide.conficconi@polimi.it>

Acknowledgements

E. Del Sozzo, Marco D. Santambrogio, D. Sciuto

Part of this material comes from:

- “Computer Organization and Design” and “Computer Architecture A Quantitative Approach” Patterson and Hennessy books
- News and paper cited throughout the lecture

and are **properties of their respective owners**